

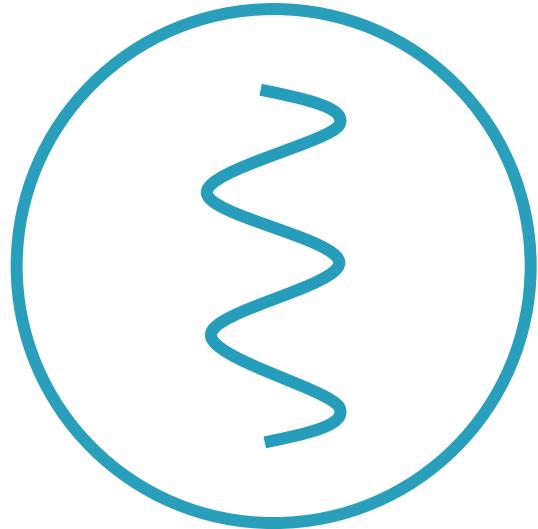
# Concurrency Factors

**Many commonplace tasks can be parallelized**

**The existence of multi-core machines**

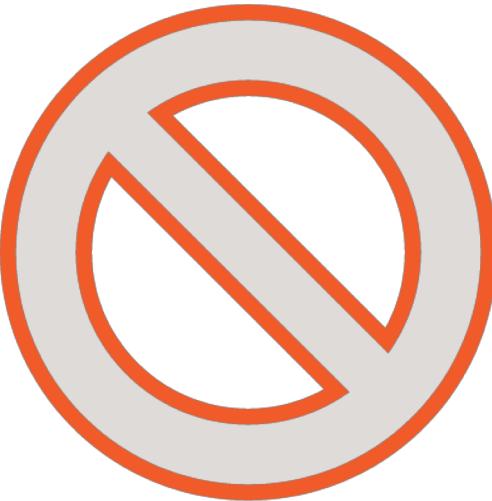


# Java Concurrency Support



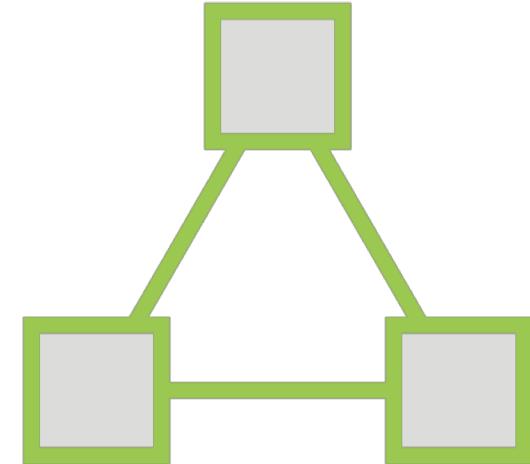
## Thread

Initiating and managing  
threaded execution



## Synchronized

Obtaining and releasing  
locks

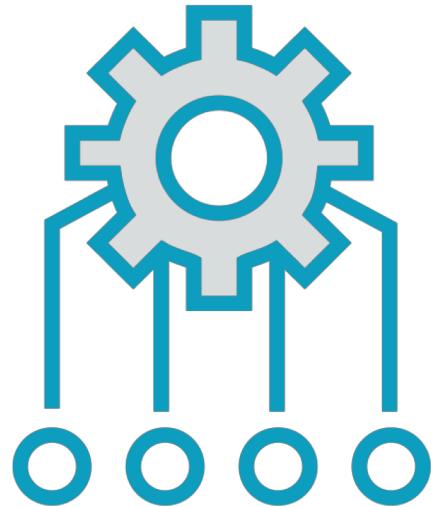


## Wait/Notify/NotifyAll

Thread coordination

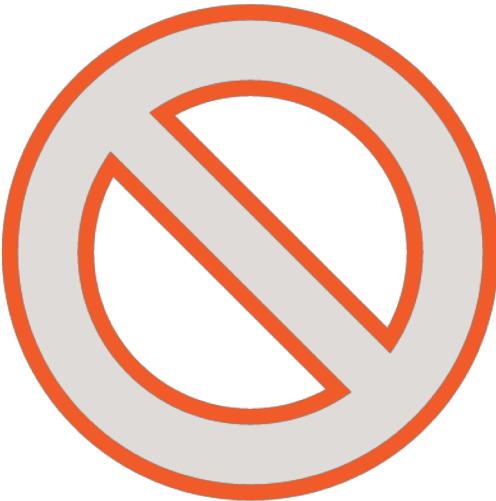


# Java Concurrency Support



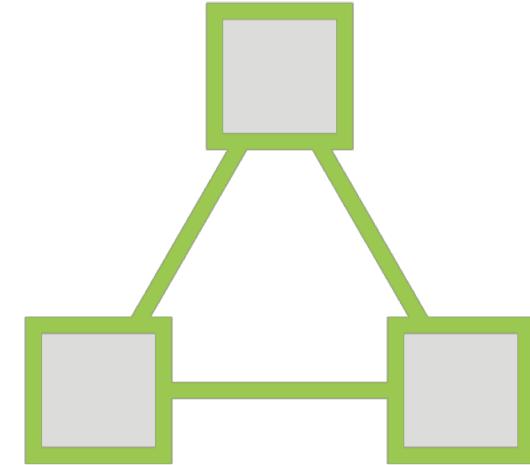
## Executor Service

Initiating and managing  
threaded execution



## Synchronized

Obtaining and releasing  
locks

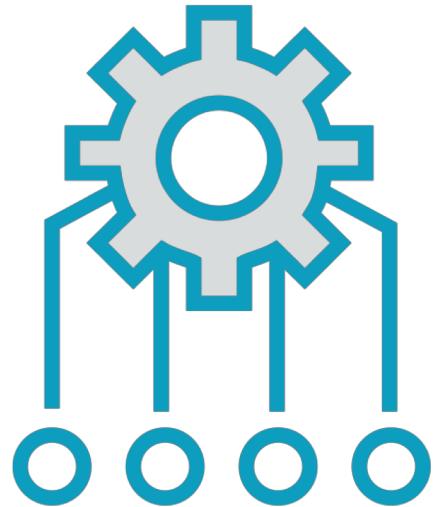


## Wait/Notify/NotifyAll

Thread coordination



# Java Concurrency Support



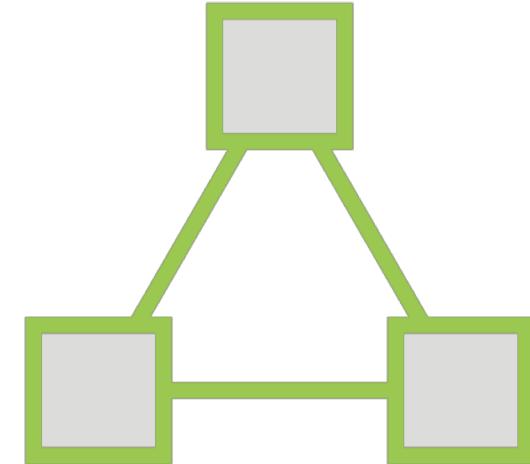
## Executor Service

Initiating and managing  
threaded execution



## Lock

Obtaining and releasing  
locks

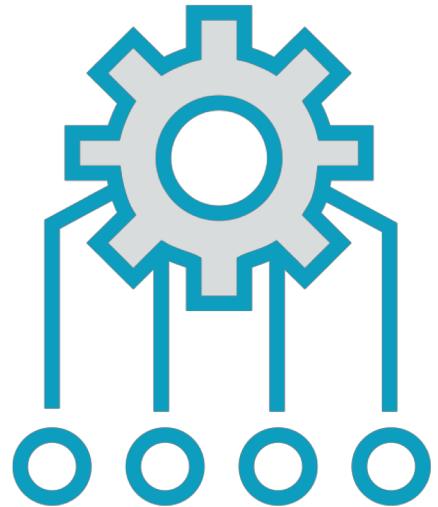


## Wait/Notify/NotifyAll

Thread coordination



# Java Concurrency Support



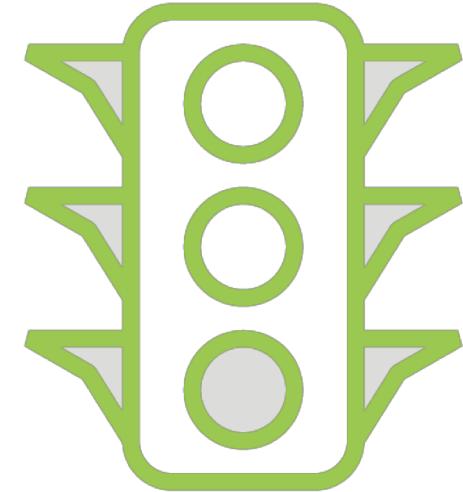
## Executor Service

Initiating and managing  
threaded execution



## Lock

Obtaining and releasing  
locks



## Condition/Semaphore /CountDownLatch/ CyclicBarrier

Thread coordination



# Concurrency Challenges

**Even with the new tools, we can still write inefficient concurrent code**

## **Challenges:**

- Selecting the correct concurrency tool
- Knowing how to use and configure the tools
- Knowing how to avoid certain pitfalls



# Overview



**ThreadPoolExecutor Optimization**

**The ForkJoinPool**

**Reducing Lock Contention**

**Atomic Variables and Concurrent Collections**

**Avoiding Synchronization**



# ThreadPoolExecutor Optimization

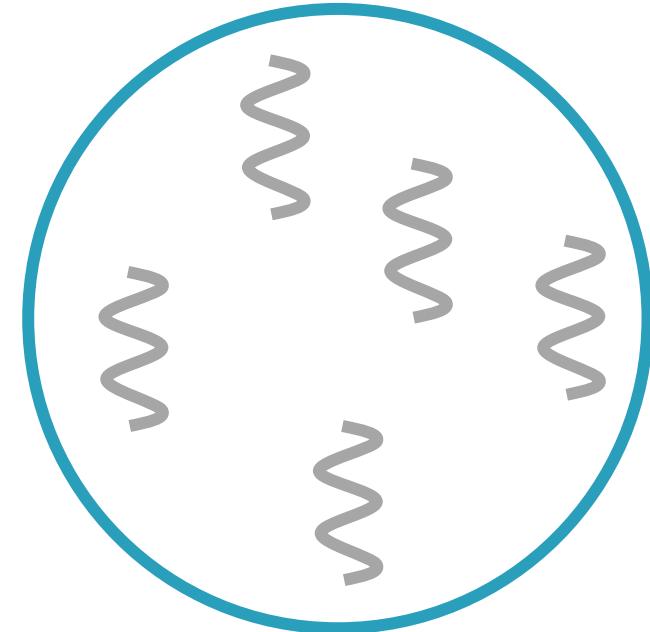
---



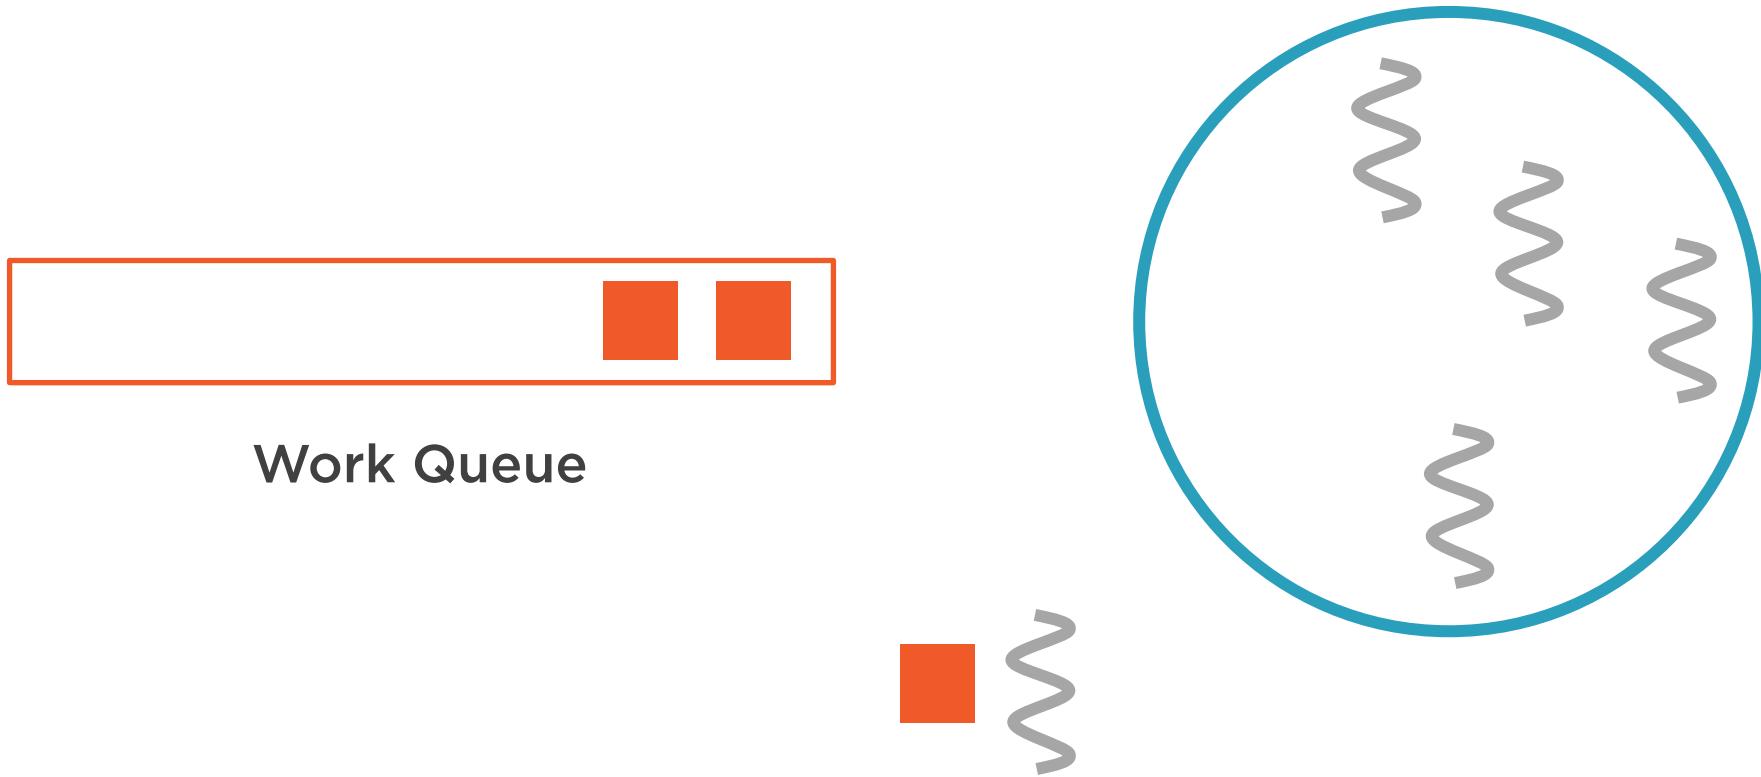
# Thread Pools



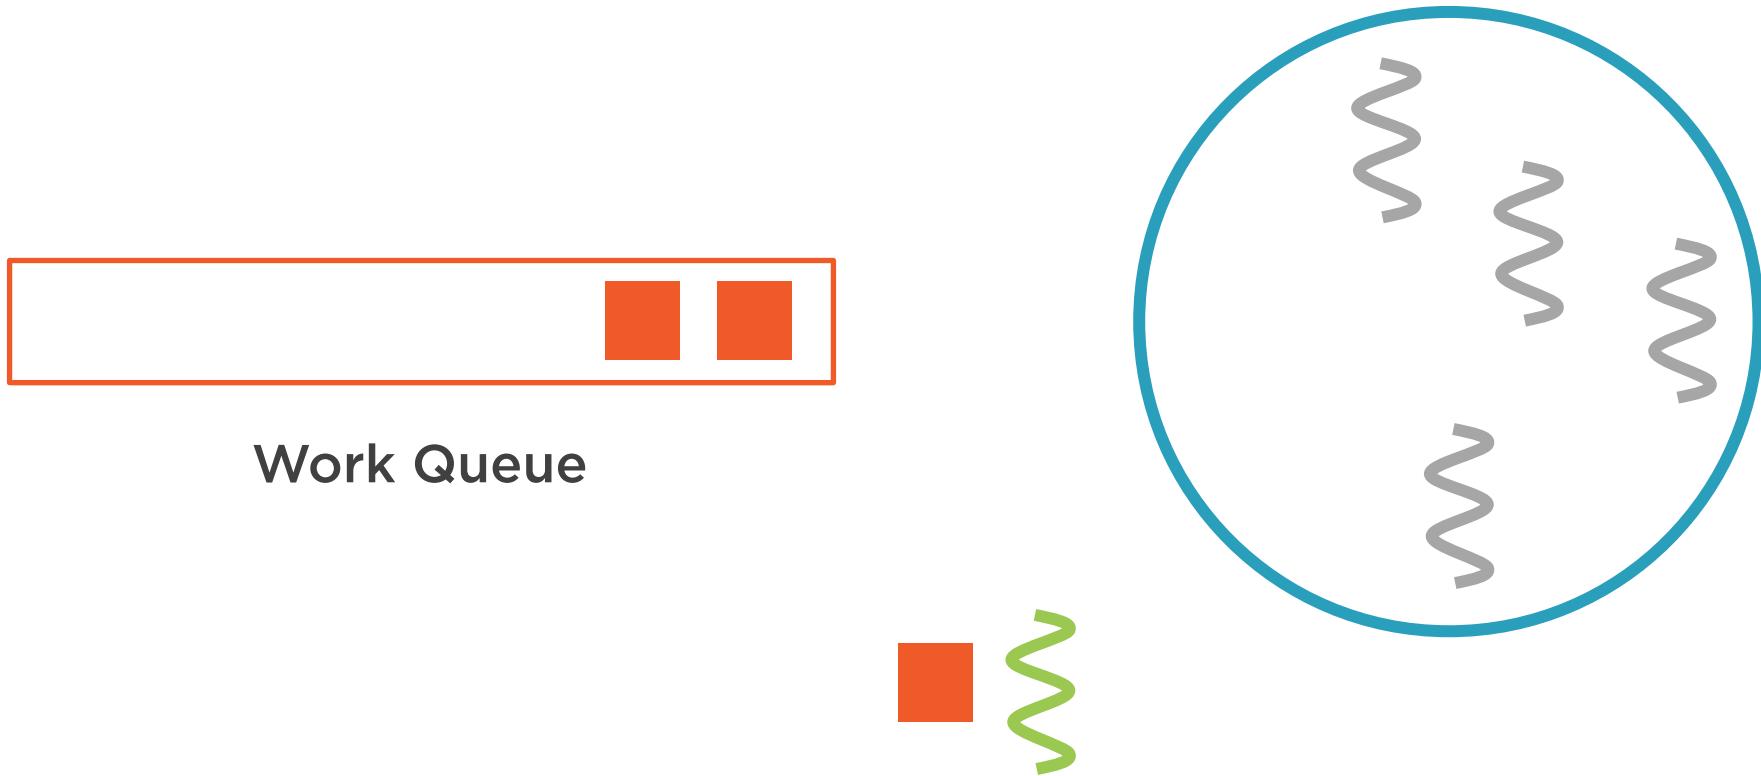
Work Queue



# Thread Pools



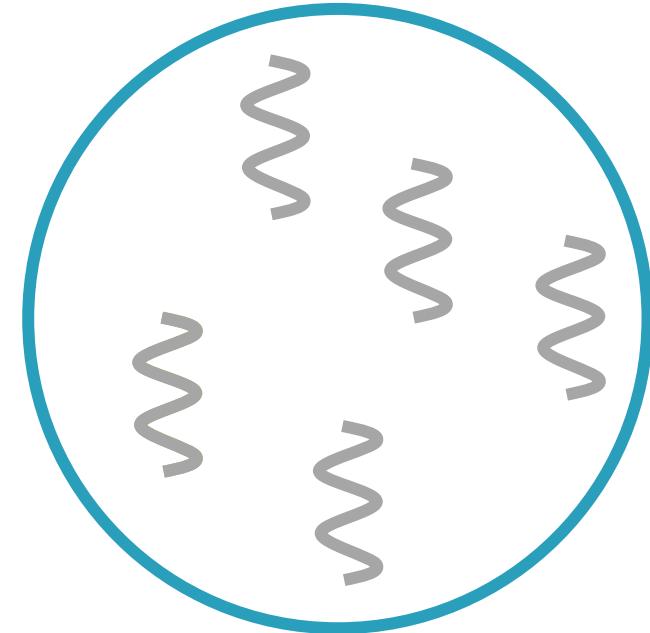
# Thread Pools



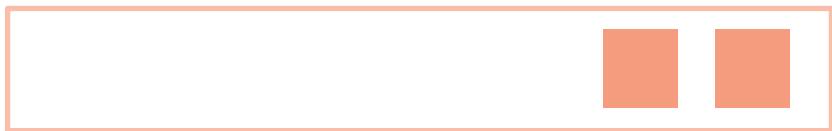
# Thread Pools



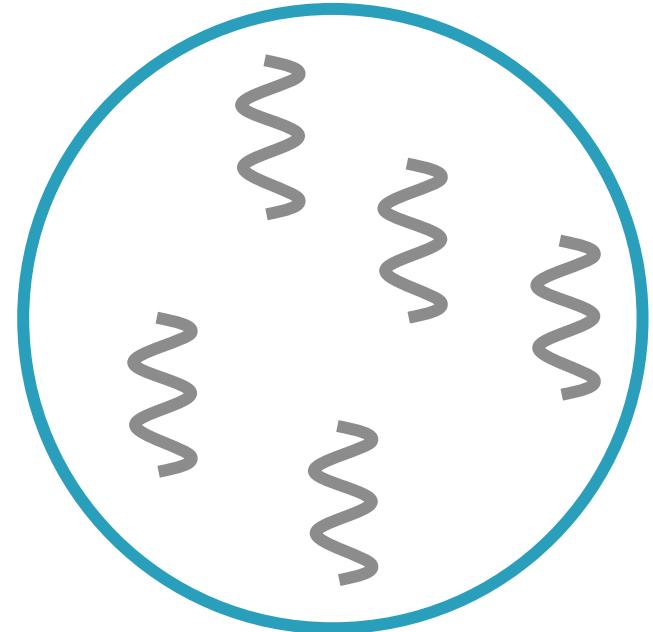
Work Queue



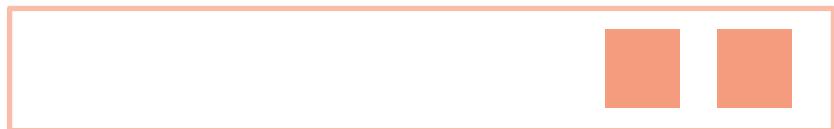
# Thread Pools



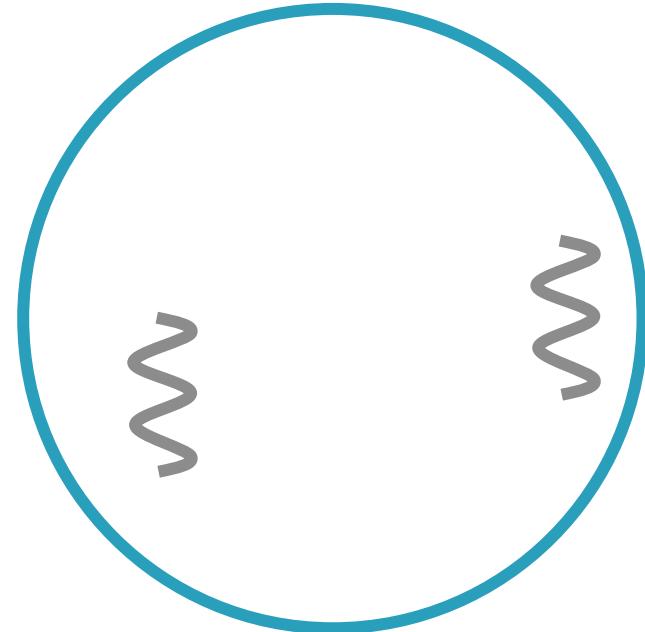
Work Queue



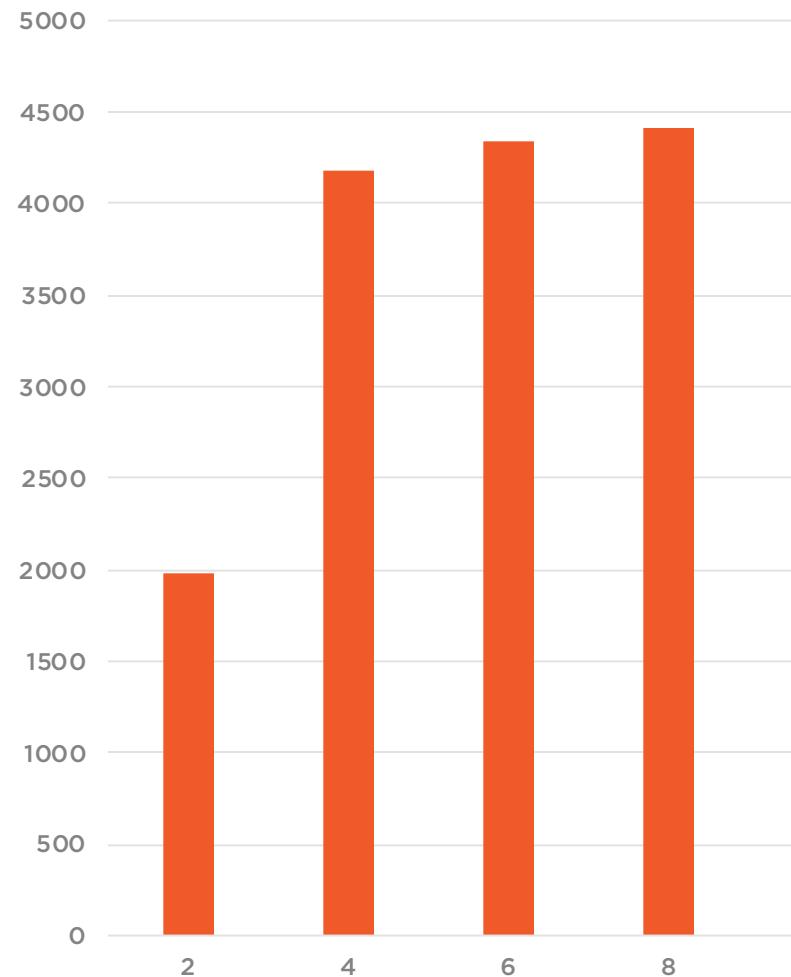
# Thread Pools



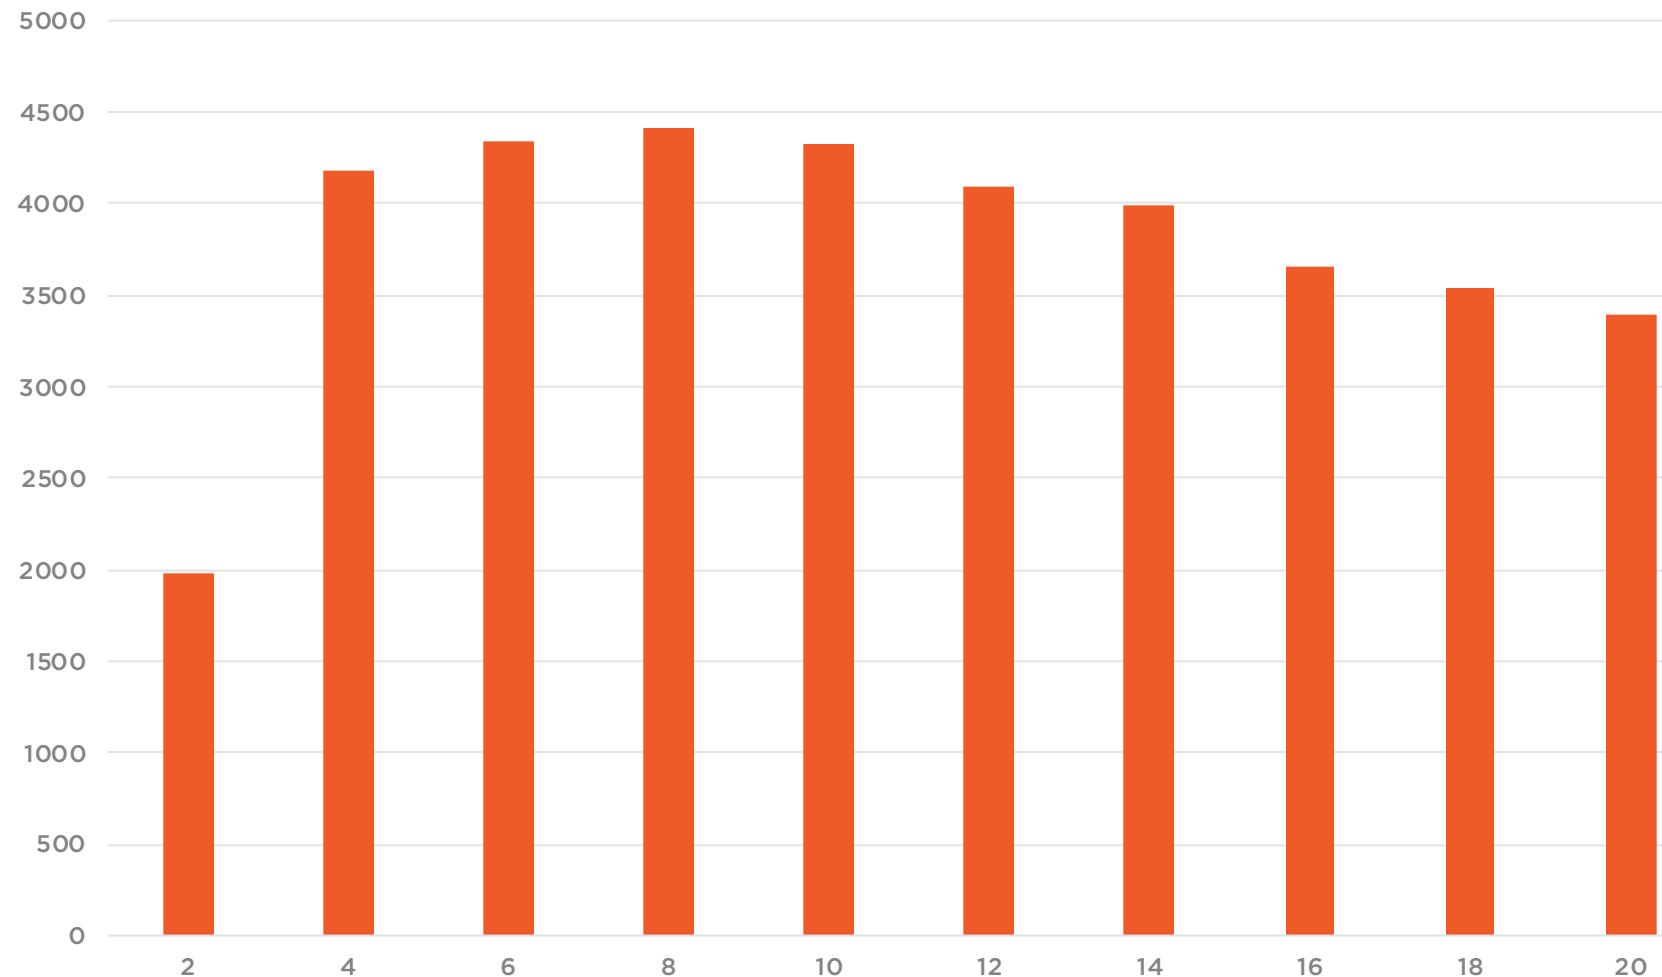
Work Queue



# Thread Pool Size



# Thread Pool Size



# Drawbacks of Having Too Many Threads

The CPU has to dedicate more time to scheduling threads

Context switches require expensive context unloading and loading operations

Caches may be invalidated and data locality may be lost



Tuning the size of the thread pool  
is not an exact science.



# Thread Pool Sizing Approaches



Model-based approach



Experimental approach



# Thread Pool Sizing Approaches



Model-based approach



Experimental approach



$$N_{\text{threads}} = N_{\text{cpu}} * U_{\text{cpu}} * (1 + W/C)$$

$N_{\text{cpu}}$  = number of CPU cores

$U_{\text{cpu}}$  = desired CPU utilization,  $0 < U_{\text{cpu}} \leq 1$

$W/C$  = ratio of wait time to compute time



$$N_{\text{threads}} = N_{\text{cpu}} * U_{\text{cpu}} * (1 + 0)$$

$N_{\text{cpu}}$  = number of CPU cores

$U_{\text{cpu}}$  = desired CPU utilization,  $0 < U_{\text{cpu}} \leq 1$

W/C = ratio of wait time to compute time



$$N_{\text{threads}} = N_{\text{cpu}} * 1 * (1 + O)$$

$N_{\text{cpu}}$  = number of CPU cores

$U_{\text{cpu}}$  = desired CPU utilization,  $0 < U_{\text{cpu}} \leq 1$

W/C = ratio of wait time to compute time



$$N_{\text{threads}} = N_{\text{cpu}} * 1 * 1$$

$N_{\text{cpu}}$  = number of CPU cores

$U_{\text{cpu}}$  = desired CPU utilization,  $0 < U_{\text{cpu}} \leq 1$

W/C = ratio of wait time to compute time



$$N_{\text{threads}} = N_{\text{cpu}} * 0.6 * 1$$

$N_{\text{cpu}}$  = number of CPU cores

$U_{\text{cpu}}$  = desired CPU utilization,  $0 < U_{\text{cpu}} \leq 1$

W/C = ratio of wait time to compute time



$$N_{\text{threads}} = N_{\text{cpu}} * 1 * (1 + 75/25)$$

$N_{\text{cpu}}$  = number of CPU cores

$U_{\text{cpu}}$  = desired CPU utilization,  $0 < U_{\text{cpu}} \leq 1$

W/C = ratio of wait time to compute time



$$N_{\text{threads}} = N_{\text{cpu}} * 1 * (1 + 3)$$

$N_{\text{cpu}}$  = number of CPU cores

$U_{\text{cpu}}$  = desired CPU utilization,  $0 < U_{\text{cpu}} \leq 1$

W/C = ratio of wait time to compute time



# Variations of the Model

**Start with the model, then tweak the results**

**Tweak the formula itself and use different coefficients**

**Use a different model**



# Other Factors Affecting Thread Pool Sizing

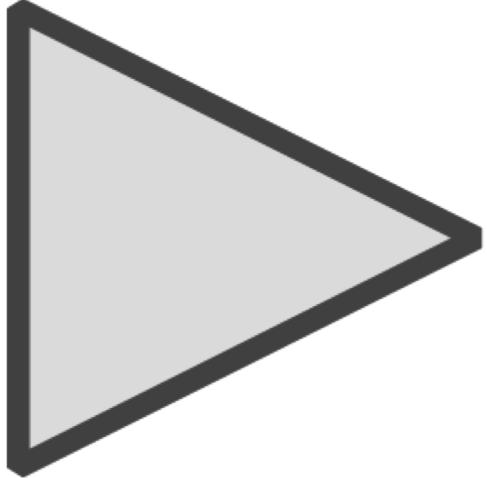
**Memory**

**File / socket handles**

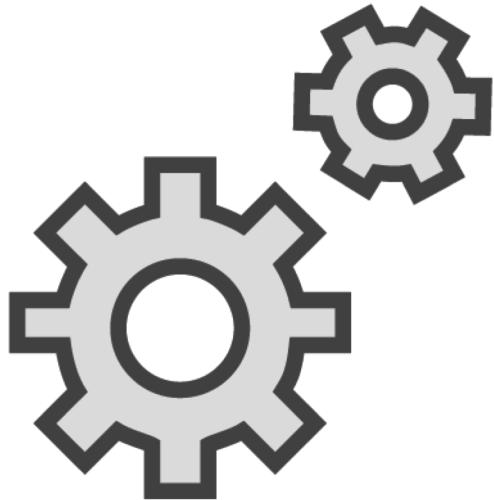
**Database connections**



# Experimental Approach to Thread Pool Sizing



**Send traffic to the application**



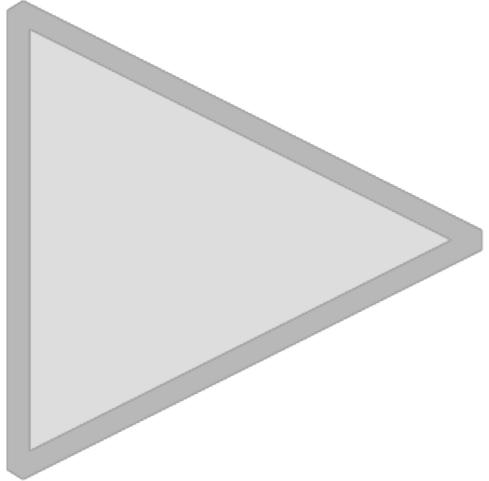
**Vary the thread pool sizes**



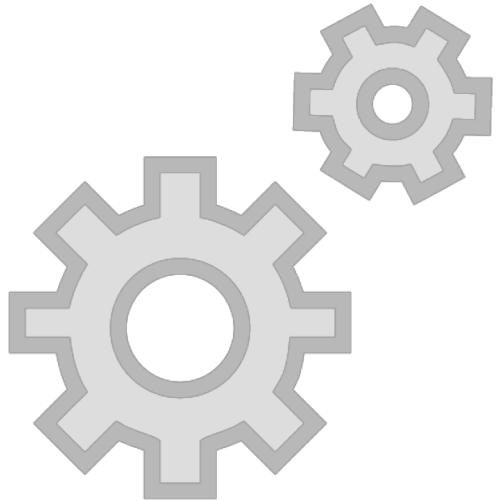
**Monitor system performance and thread characteristics**



# Experimental Approach to Thread Pool Sizing



Send traffic to the application



Vary the thread pool sizes



**Monitor system performance and thread characteristics**



# Thread Pool Sizing Best Practices

Instead of hard-coding thread pool sizes,  
make it configurable or calculate it  
dynamically

If plausible, split CPU and I/O work into  
separate thread pools that can be  
configured individually

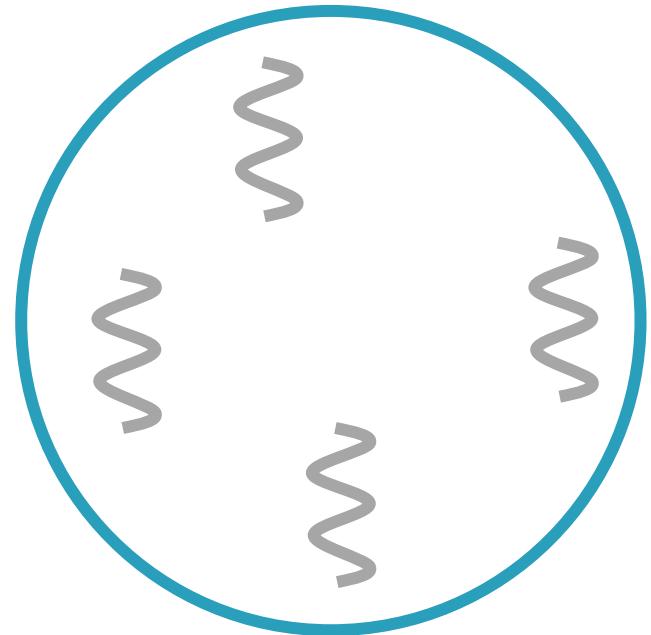
Focus first on tuning core thread pools  
before auxiliary ones



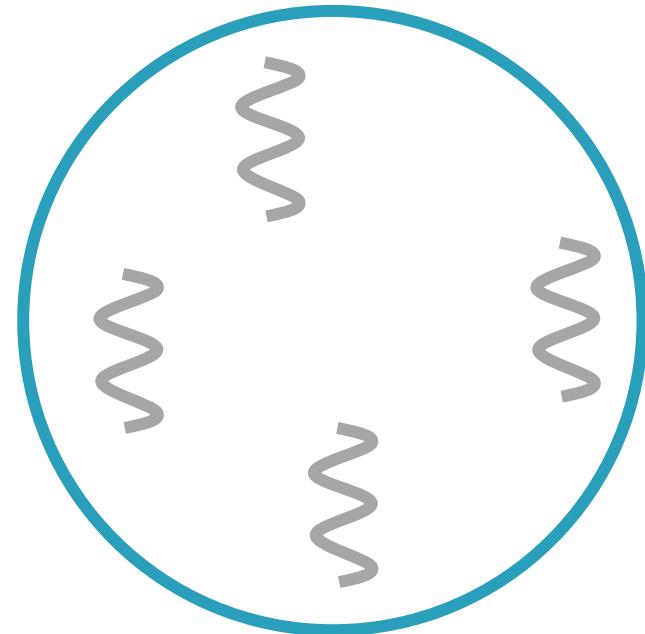
```
public ThreadPoolExecutor(int corePoolSize,  
                         int maximumPoolSize,  
                         long keepAliveTime,  
                         TimeUnit unit,  
                         BlockingQueue<Runnable> workQueue)
```



# ThreadPoolExecutor



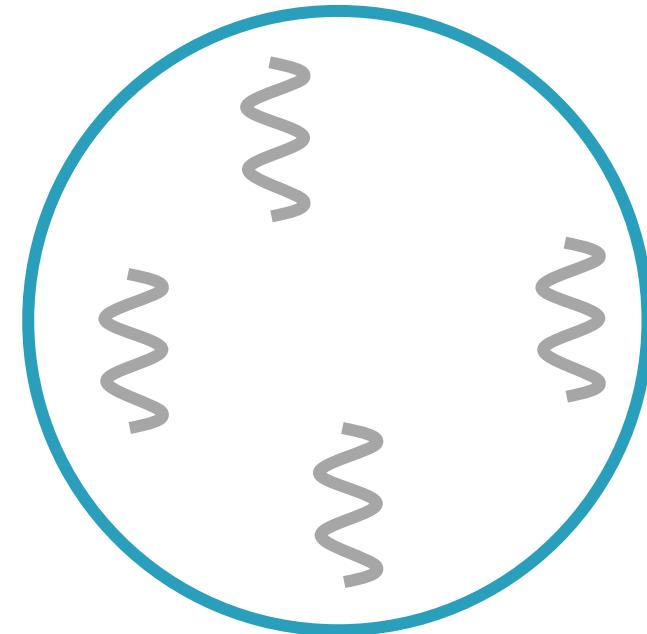
# ThreadPoolExecutor



# ThreadPoolExecutor



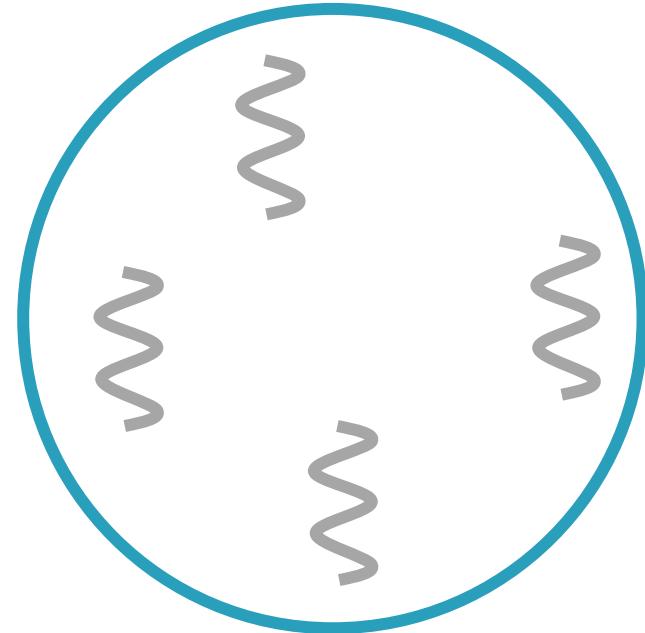
Unbounded / Bounded / Synchronous



# ThreadPoolExecutor



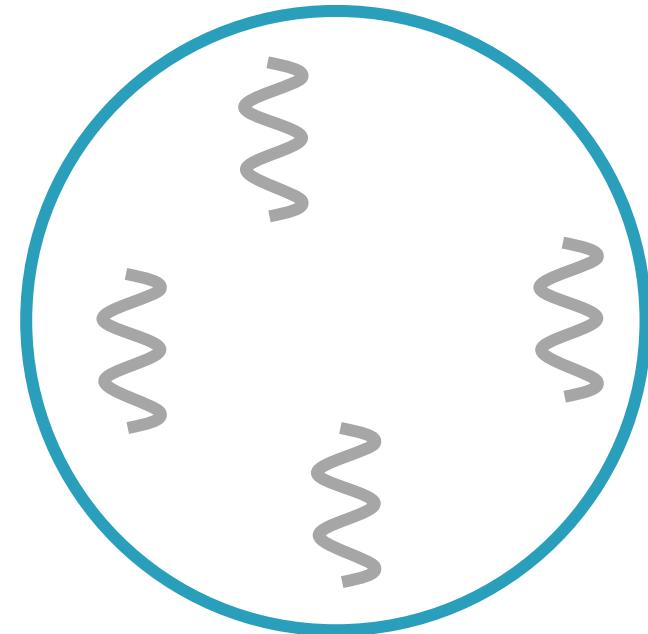
Unbounded



# ThreadPoolExecutor



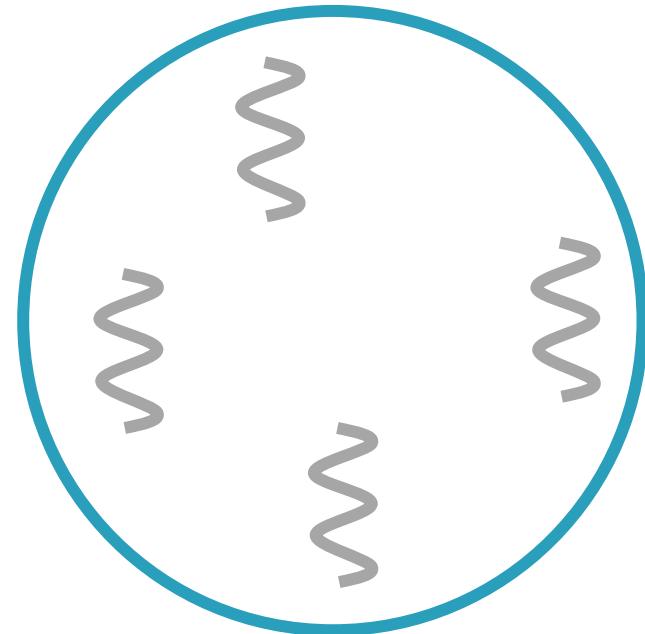
Unbounded



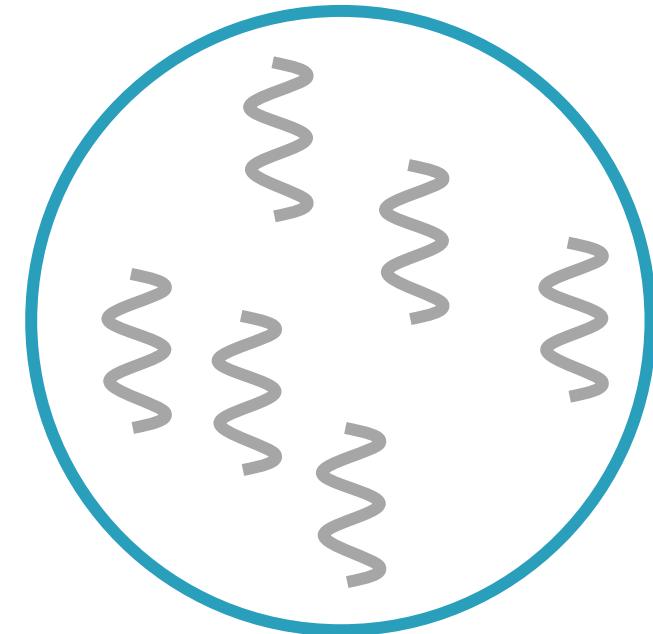
# ThreadPoolExecutor



Bounded



# ThreadPoolExecutor



# Server Application Best Practices

**Determine optimal pool size**

**Use a bounded queue with a limit just large enough to handle a burst of requests**

**Set:**

**corePoolSize = maxPoolSize = optimalSize**

**Or:**

**corePoolSize = optimalSizeMin**

**maxPoolSize = optimalSizeMax**



# Work Queue Implementation

## **LinkedBlockingQueue, ArrayBlockingQueue**

Tasks are started in the order they  
were received

## **PriorityBlockingQueue**

Tasks are started in order of  
priority



# The ForkJoinPool

---



# Recursive Model



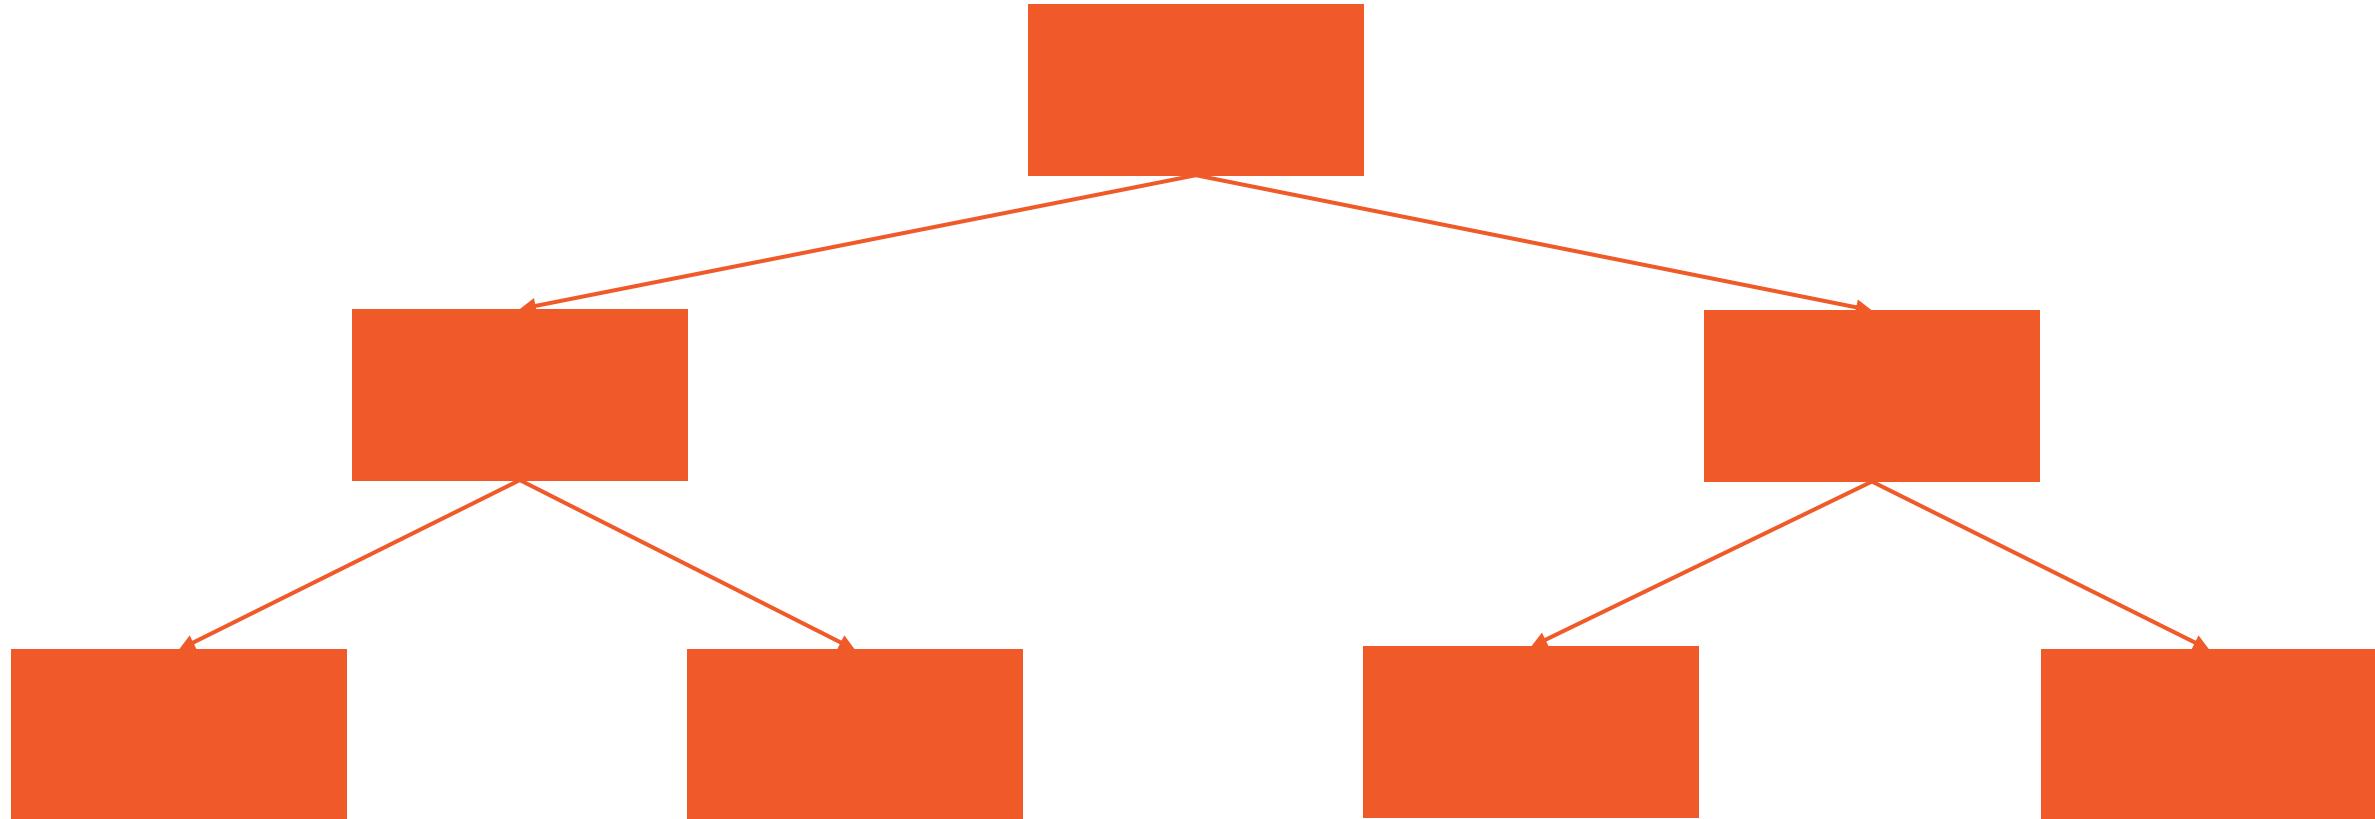
# Recursive Model



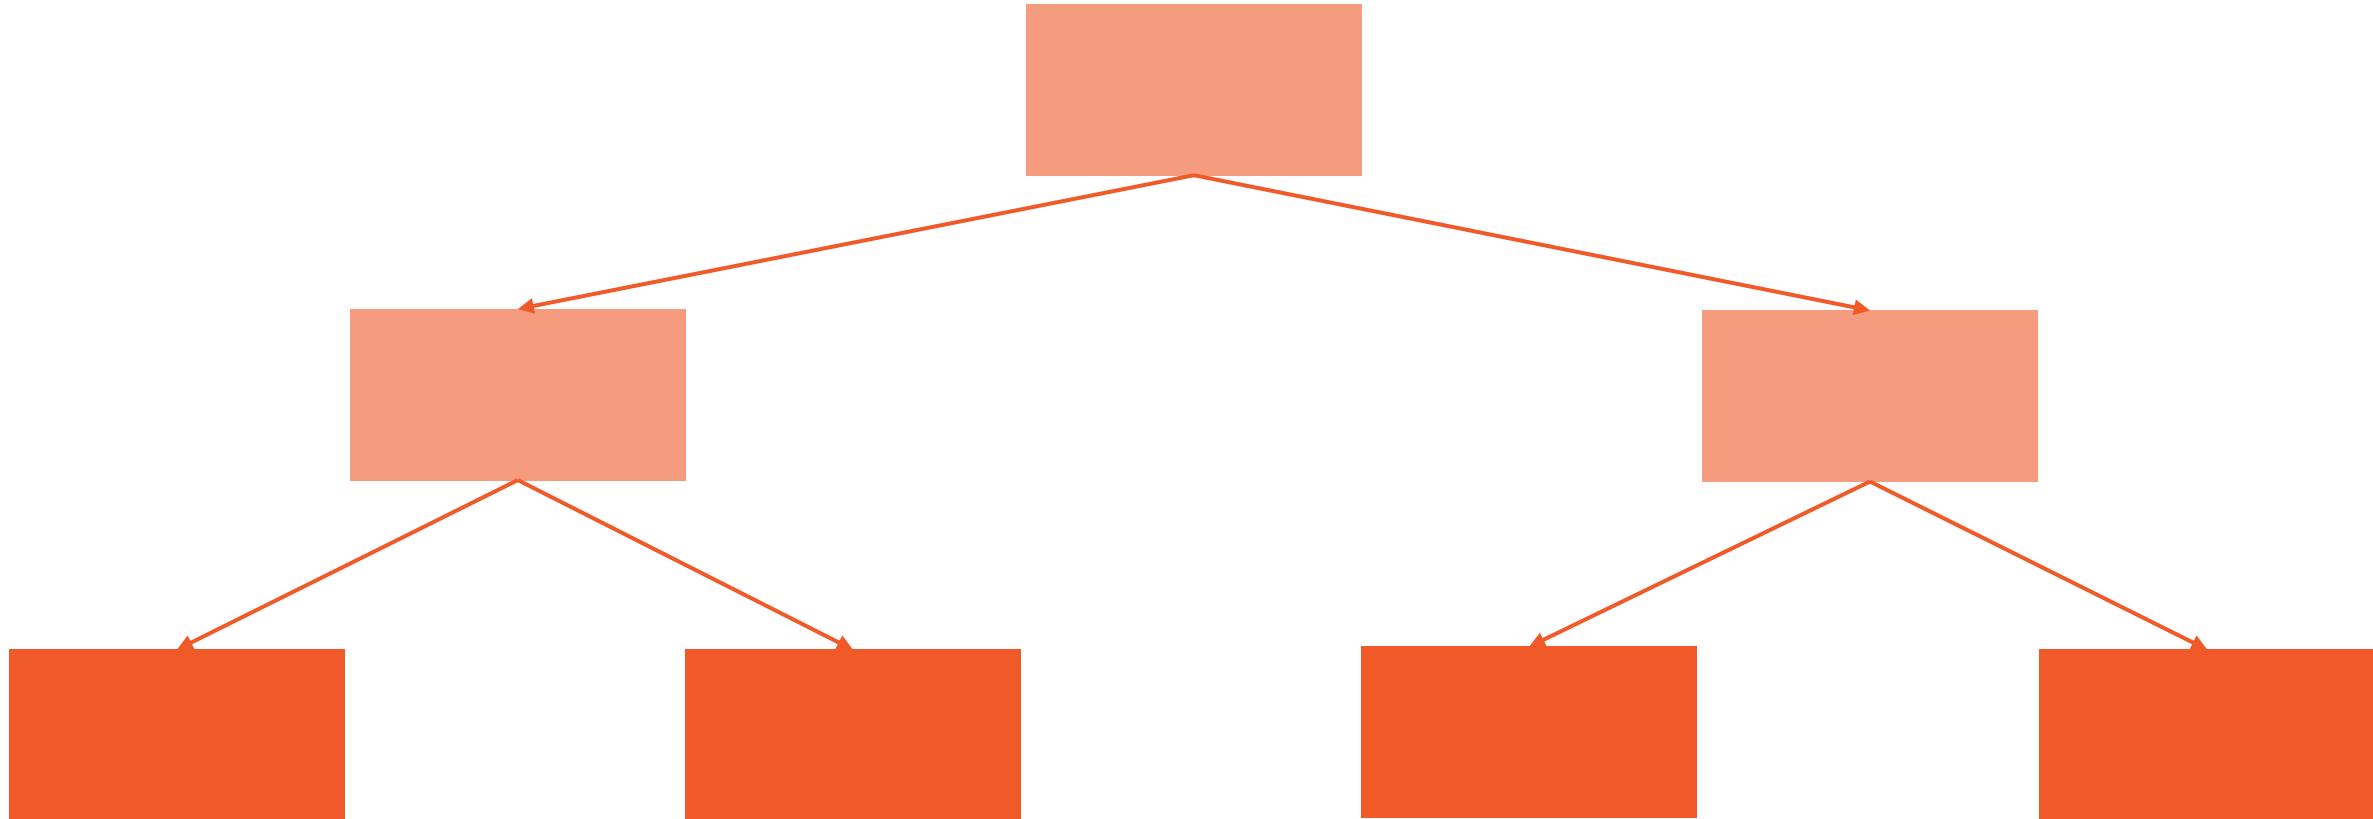
# Recursive Model



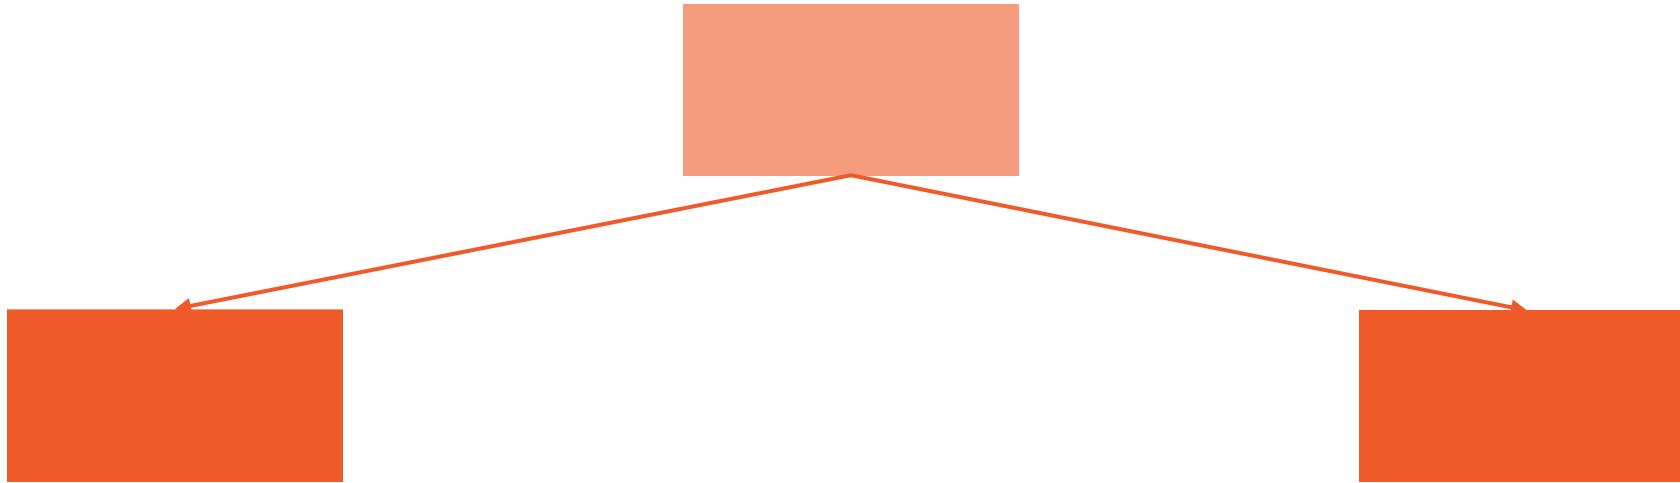
# Recursive Model



# Recursive Model



# Recursive Model



# Recursive Model



```
public class MyRecursiveTask extends RecursiveTask<Long> {  
    private Task task;  
  
    public MyRecursiveTask(Task task) {  
        this.task = task;  
    }  
  
    protected Long compute() {  
        if (this.task.size() > threshold) {  
            Pair<RecursiveTask> subTasks = splitTask();  
            subTasks.getLeft().fork();  
            subTasks.getRight().fork();  
            ...  
        } else { ... }  
    }  
}
```



```
public class MyRecursiveTask extends RecursiveTask<Long> {  
    private Task task;  
  
    public MyRecursiveTask(Task task) {  
        this.task = task;  
    }  
  
    protected Long compute() {  
        if (this.task.size() > threshold) {  
            Pair<RecursiveTask> subTasks = splitTask();  
            subTasks.getLeft().fork();  
            subTasks.getRight().fork();  
            ...  
        } else { ... }  
    }  
}
```



```
public class MyRecursiveTask extends RecursiveTask<Long> {  
    ...  
  
    protected Long compute() {  
        if (this.task.size() > threshold) {  
            Pair<RecursiveTask> subTasks = splitTask();  
            subTasks.getLeft().fork();  
            subTasks.getRight().fork();  
            long result = 0;  
            result += subTasks.getLeft().join();  
            result += subTasks.getRight().join();  
            return result;  
        } else { ... }  
    }  
}
```



```
public class MyRecursiveTask extends RecursiveTask<Long> {  
    ...  
    protected Long compute() {  
        if (this.task.size() > threshold) {  
            Pair<RecursiveTask> subTasks = splitTask();  
            ...  
        } else {  
            System.out.println("Executing task myself");  
            return executeTask(task);  
        }  
    }  
}
```



```
public class MyRecursiveTask extends RecursiveTask<Long> {  
    ...  
    protected void compute() {  
        if (this.task.size() > threshold) {  
            Pair<RecursiveTask> subTasks = splitTask();  
            subTasks.getLeft().fork();  
            subTasks.getRight().fork();  
            long result = 0;  
            result += subTasks.getLeft().join();  
            result += subTasks.getRight().join();  
            return result;  
        } else { ... }  
    }  
}
```



```
public class MyRecursiveTask extends RecursiveAction {  
    ...  
    protected void compute() {  
        if (this.task.size() > threshold) {  
            Pair<RecursiveTask> subTasks = splitTask();  
            subTasks.getLeft().fork();  
            subTasks.getRight().fork();  
            long result = 0;  
            result += subTasks.getLeft().join();  
            result += subTasks.getRight().join();  
            return result;  
        } else { ... }  
    }  
}
```



```
public class MyRecursiveTask extends RecursiveAction {  
    ...  
    protected void compute() {  
        if (this.task.size() > threshold) {  
            Pair<RecursiveTask> subTasks = splitTask();  
            subTasks.getLeft().fork();  
            subTasks.getRight().fork();  
        } else { ... }  
    }  
}
```



```
MyRecursiveTask myRecurTask = new MyRecursiveTask(largeTask);  
ForkJoinPool forkJoinPool = new ForkJoinPool(8);  
forkJoinPool.invoke(myRecurTask);
```



```
MyRecursiveTask myRecurTask = new MyRecursiveTask(largeTask);  
ForkJoinPool forkJoinPool = new ForkJoinPool(8);  
forkJoinPool.invoke(myRecurTask);
```



```
MyRecursiveTask myRecurTask = new MyRecursiveTask(largeTask);  
ForkJoinPool forkJoinPool = new ForkJoinPool(8);  
forkJoinPool.invoke(myRecurTask);
```



```
public class UserRequest extends RecursiveTask<Result> {  
    protected Result compute() {  
        if (this.numItems < 100) {  
            // directly execute UserRequest  
        } else {  
            // split into smaller UserRequest subtasks  
        }  
    }  
}
```



```
public class UserRequest extends RecursiveTask<Result> {  
    protected Result compute() {  
        if (this.numItems < 100) {  
            // directly execute UserRequest  
        } else {  
            // split into smaller UserRequest subtasks  
        }  
    }  
}
```



# ForkJoinPool

Can handle a large amount of tasks with very few threads

Used by the Java 8 Parallel Streams feature under the hood

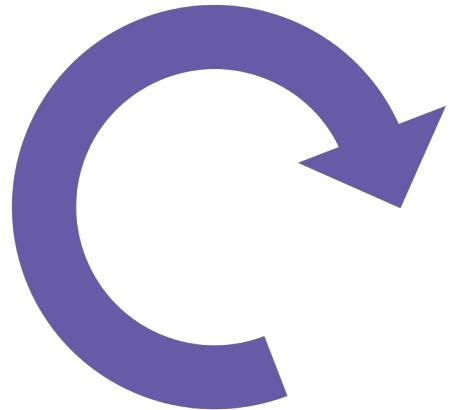


# Reducing Lock Contention

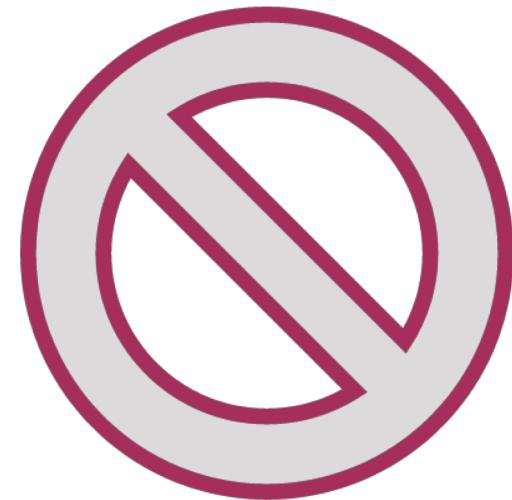
---



# Blocking Mechanisms



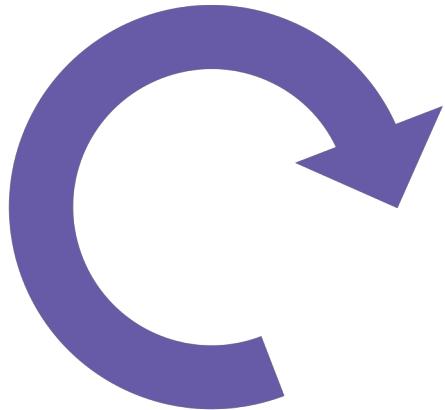
Spin-Wait



OS Suspension

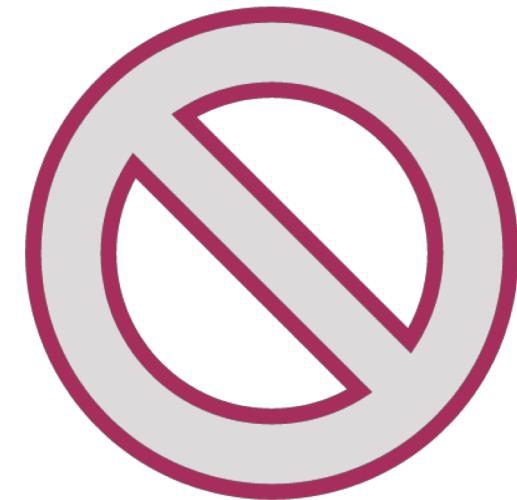


# Blocking Mechanisms



**Spin-Wait**

More efficient for short waits



**OS Suspension**

Preferable for longer waits



# JVM Blocking Implementations

Spin-wait + suspend after X number of spins

Spin-wait or suspend, depending on profiling data

Default to suspension



Locks are not bad, lock contention  
is.



# Performance Impacts of Lock Contention

**Results in serial execution and reduces the scalability of the application**

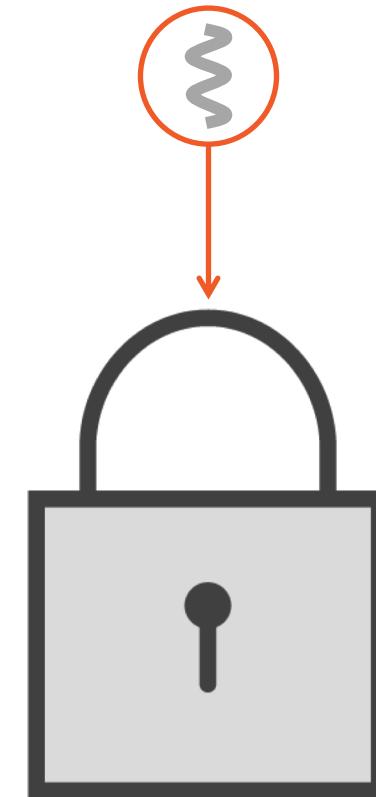
**Leads to context switches which hurt performance**



# Lock Classification



High Contention



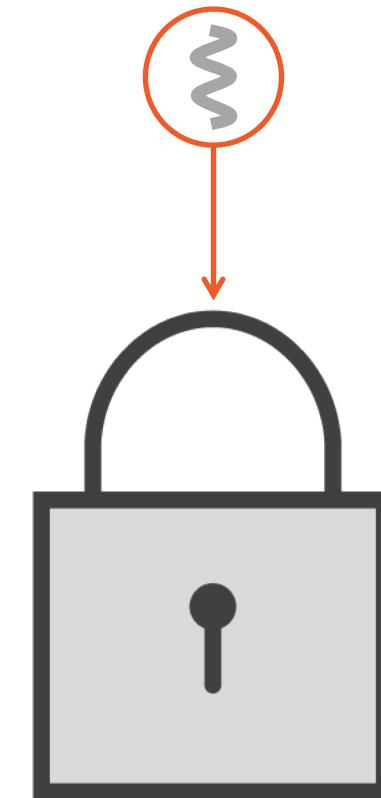
Low Contention



# Lock Classification



High Contention



Low Contention



First make it right, then make it fast, if it's not already fast enough



# Reducing Lock Contention

**Reduce the duration for which the lock is held**

**Reduce the frequency at which the lock is requested**

**Replace locks with alternative synchronization methods that allow for greater concurrency**



# Reducing the Duration for Which a Lock Is Held

**Reduce the size of the critical block to the point where only the shared data access is protected**



```
public class AggregateScore {  
    private static final double MAX_SCORE_PER_SCORER = 100;  
    private double score;  
  
    public void synchronized updateScore(Scorer scorer) {  
        if (scorer.isApplicable(this)) {  
            double scorerScore = scorer.calculateScore();  
            this.score += scorerScore < MAX_SCORE_PER_SCORER ?  
                scorerScore : MAX_SCORE_PER_SCORER;  
        }  
    }  
}
```



```
public class AggregateScore {  
    private static final double MAX_SCORE_PER_SCORER = 100;  
    private double score;  
  
    public void synchronized updateScore(Scorer scorer) {  
        if (scorer.isApplicable(this)) {  
            double scorerScore = scorer.calculateScore();  
            this.score += scorerScore < MAX_SCORE_PER_SCORER ?  
                scorerScore : MAX_SCORE_PER_SCORER;  
        }  
    }  
}
```



```
public class AggregateScore {  
    private static final double MAX_SCORE_PER_SCORER = 100;  
    private double score;  
  
    public void synchronized updateScore(Scorer scorer) {  
        if (scorer.isApplicable(this)) {  
            double scorerScore = scorer.calculateScore();  
            this.score += scorerScore < MAX_SCORE_PER_SCORER ?  
                scorerScore : MAX_SCORE_PER_SCORER;  
        }  
    }  
}
```



```
public class AggregateScore {  
    private static final double MAX_SCORE_PER_SCORER = 100;  
    private double score;  
  
    public void updateScore(Scorer scorer) {  
        if (scorer.isApplicable(this)) {  
            double scorerScore = scorer.calculateScore();  
            synchronized(this) {  
                this.score += scorerScore < MAX_SCORE_PER_SCORER ?  
                    scorerScore : MAX_SCORE_PER_SCORER; }  
            }  
    }  
}
```



# Reducing the Demand for a Lock

## Lock splitting:

- Takes a single lock that is guarding different independent shared state variables and splits it into multiple locks for each variable



```
public class ServerStatistics {  
    private long numActiveUsers;  
    private long numDBQueries;  
  
    public void synchronized incrementNumActiveUsers(int increment) {  
        numActiveUsers += increment;  
    }  
    public void synchronized incrementNumDBQueries(int increment) {  
        numDBQueries += increment;  
    }  
    public long synchronized getNumActiveUsers { ... }  
    public long synchronized getNumDBQueries { ... }  
}
```



```
public class ServerStatistics {  
    private long numActiveUsers;  
    private long numDBQueries;  
  
    public void synchronized incrementNumActiveUsers(int increment) {  
        numActiveUsers += increment;  
    }  
    public void synchronized incrementNumDBQueries(int increment) {  
        numDBQueries += increment;  
    }  
    public long synchronized getNumActiveUsers { ... }  
    public long synchronized getNumDBQueries { ... }  
}
```



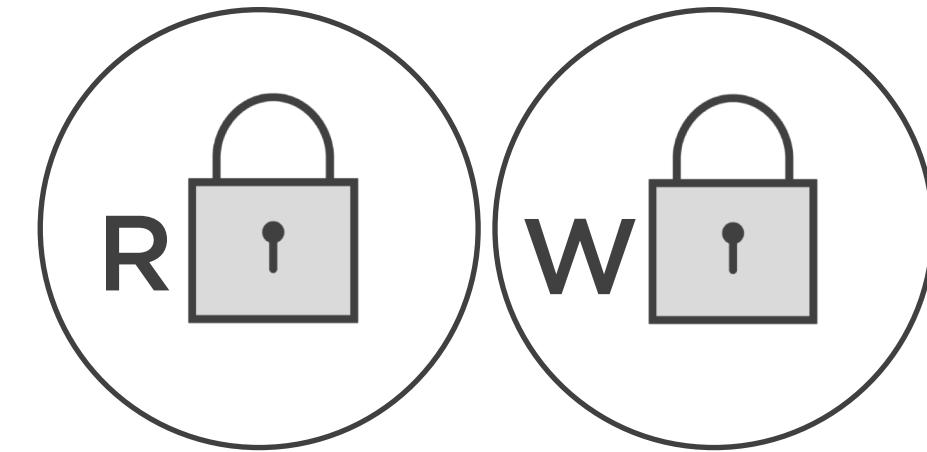
```
public class ServerStatistics {  
    private long numActiveUsers;  
    private long numDBQueries;  
    private Object activeUsersLock = new Object();  
    private Object dbQueriesLock = new Object();  
  
    public void incrementNumActiveUsers(int increment) {  
        synchronized (activeUsersLock) { numActiveUsers += increment; }  
    }  
  
    public void incrementNumDBQueries(int increment) {  
        synchronized (dbQueriesLock) { numDBQueries += increment; }  
    }  
    ...  
}
```

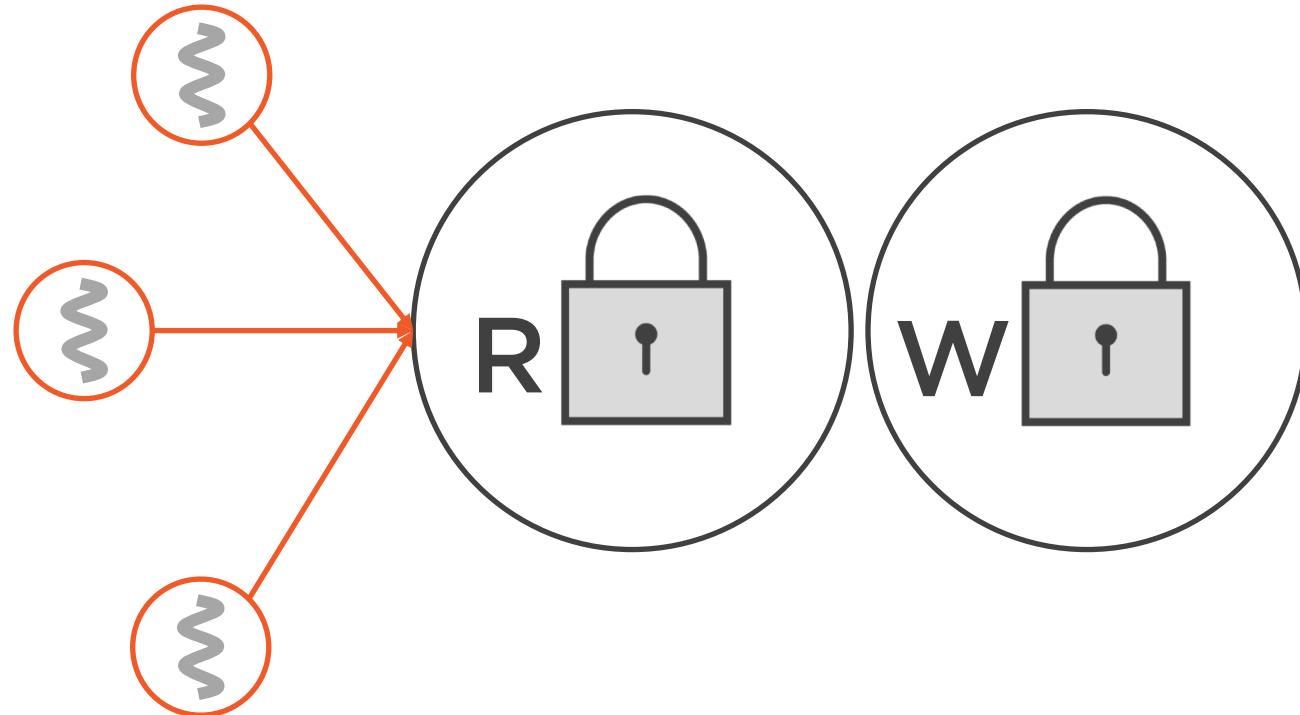


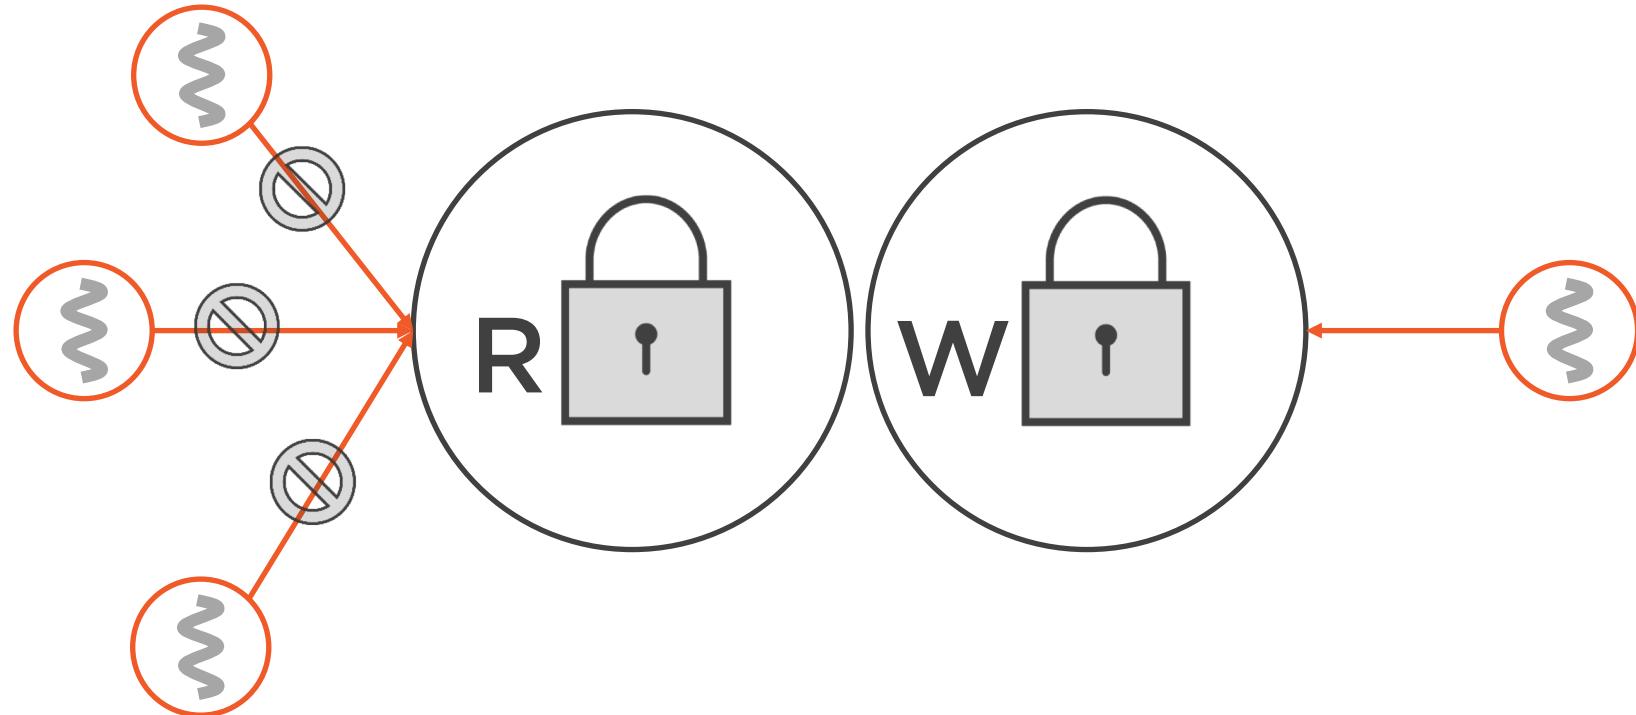
# ReadWriteLock

Ideal when; a collection or data structure is written to infrequently but reads happen at a high frequency









# Lock Striping

**Reduces the granularity at which a lock is applied**

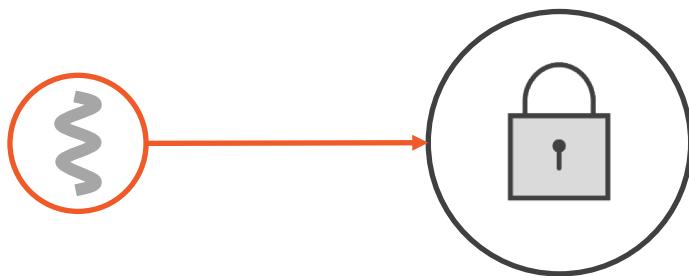
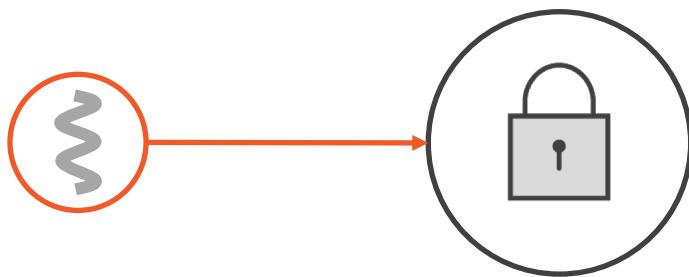
**Applicable when you have a variable sized collection of independent objects**











```
public class Users {  
    private Object[] locks = new Object[16]; ...  
  
    private Object getLock(long id) {  
        return locks[id % locks.length];  
    }  
  
    public void updateUserLocation(long userId, float lat, float lon) {  
        synchronized (getLock(userId)) {  
            User user = db.getUser(userId);  
            user.setLat(lat);  
            user.setLon(lon);  
        }  
    }  
}
```



```
public class Users {  
    private Object[] locks = new Object[16]; ...  
private Object getLock(long id) {  
    return locks[id % locks.length];  
}  
  
public void updateUserLocation(long userId, float lat, float lon) {  
    synchronized (getLock(userId)) {  
        User user = db.getUser(userId);  
        user.setLat(lat);  
        user.setLon(lon);  
    }  
}  
}
```



```
public class Users {  
    private Object[] locks = new Object[16]; ...  
  
    private Object getLock(long id) {  
        return locks[id % locks.length];  
    }  
  
    public void updateUserLocation(long userId, float lat, float lon) {  
        synchronized (getLock(userId)) {  
            User user = db.getUser(userId);  
            user.setLat(lat);  
            user.setLon(lon);  
        }  
    }  
}
```



# Lock Class Features

**Ability to interrupt a thread that's waiting for a lock**

**Ability to set a timeout on the amount of time to wait for a lock**

**More scalable under high contention than the synchronized implementation**



# Atomic Variables and Concurrent Collections

---



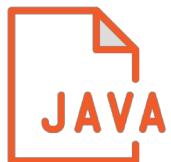
# Java Atomic Classes



**AtomicBoolean**



**AtomicInteger, AtomicIntegerArray**



**AtomicLong, AtomicLongArray**



**AtomicReference, AtomicReferenceArray**



# Atomic Classes

**Enables atomic updates**

**Enables atomic arithmetic operations in the case of AtomicInteger and AtomicLong objects**

**Atomicity prevents data corruption**

**Uses the Compare-And-Swap instruction that is provided by modern CPUs**



```
public class AggregateScore {  
    private static final double MAX_SCORE_PER_SCORER = 100;  
    private double score;  
  
    public void updateScore(Scorer scorer) {  
        if (scorer.isApplicable(this)) {  
            double scorerScore = scorer.calculateScore();  
            synchronized(this) {  
                this.score += scorerScore < MAX_SCORE_PER_SCORER ?  
                    scorerScore : MAX_SCORE_PER_SCORER; }  
            }  
    }  
}
```



```
public class AggregateScore {  
    private static final double MAX_SCORE_PER_SCORER = 100;  
private AtomicLong score = new AtomicLong();  
  
    public void updateScore(Scorer scorer) {  
        if (scorer.isApplicable(this)) {  
            double scorerScore = scorer.calculateScore();  
            synchronized(this) {  
                this.score += scorerScore < MAX_SCORE_PER_SCORER ?  
                    scorerScore : MAX_SCORE_PER_SCORER; }  
            }  
    }  
}
```



```
public class AggregateScore {  
    private static final double MAX_SCORE_PER_SCORER = 100;  
    private AtomicLong score = new AtomicLong();  
  
    public void updateScore(Scorer scorer) {  
        if (scorer.isApplicable(this)) {  
            double scorerScore = scorer.calculateScore();  
            cappedScore = scorerScore < MAX_SCORE_PER_SCORER ?  
                scorerScore : MAX_SCORE_PER_SCORER;  
            this.score.getAndAdd(Double.doubleToLongBits(cappedScore));  
        }  
    }  
}
```



# Even Faster Atomics

When there are frequent simultaneous writes by multiple threads

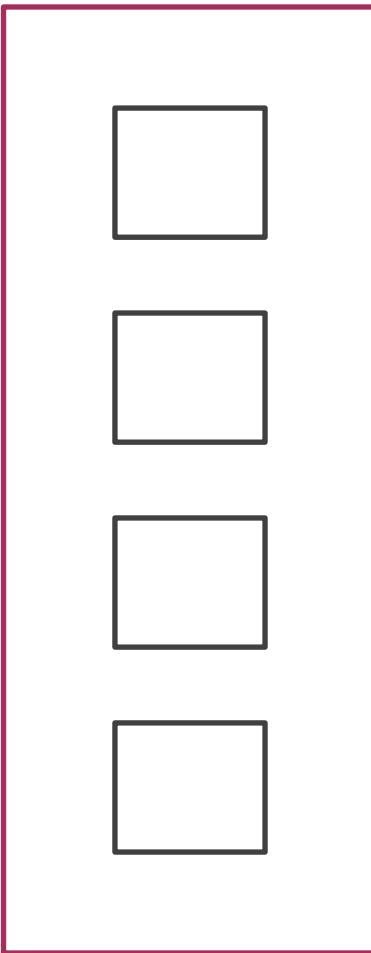
**LongAdder**

**LongAccumulator**

**Double Adder**

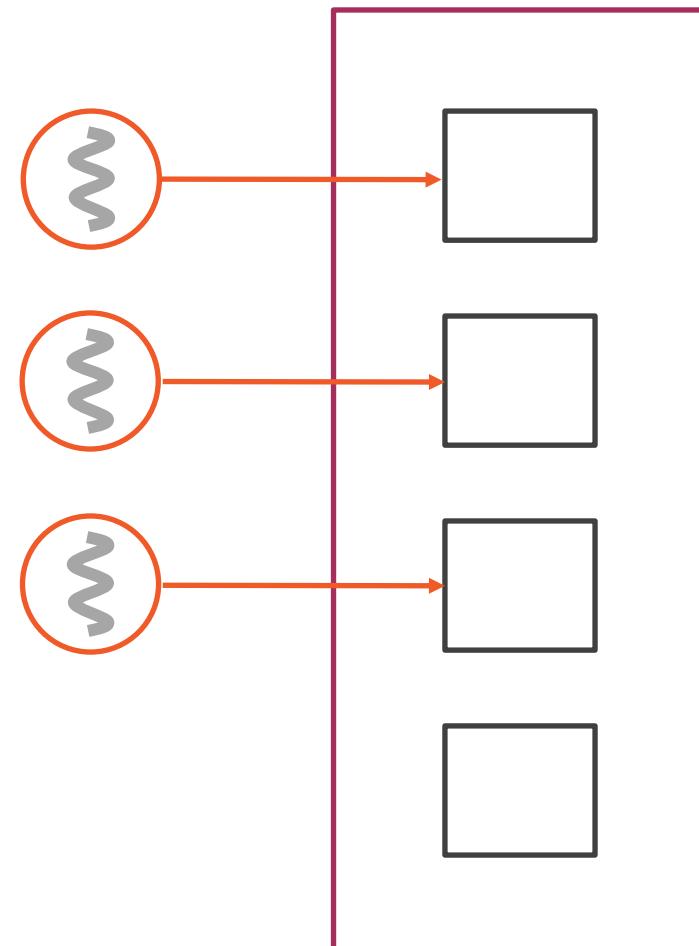
**DoubleAccumulator**





LongAdder





LongAdder



```
public class ServerStatistics {  
    private long numActiveUsers;  
    private long numDBQueries;  
    private Object activeUsersLock = new Object();  
    private Object dbQueriesLock = new Object();  
  
    public void incrementNumActiveUsers(int increment) {  
        synchronized (activeUsersLock) { numActiveUsers += increment; }  
    }  
  
    public void incrementNumDBQueries(int increment) {  
        synchronized (dbQueriesLock) { numDBQueries += increment; }  
    }  
    ...  
}
```



```
public class ServerStatistics {  
    private LongAdder numActiveUsers = new LongAdder();  
    private LongAdder numDBQueries = new LongAdder();  
  
    public void incrementNumActiveUsers(int increment) {  
        numActiveUsers.add(increment);  
    }  
  
    public void incrementNumDBQueries(int increment) {  
        numDBQueries.add(increment);  
    }  
    ...  
}
```



# Concurrent Collections

**ConcurrentHashMap**

**CopyOnWriteArrayList**



# ConcurrentHash Map

Supports highly concurrent updates and fully concurrent retrievals

Earlier implementations used lock striping to enable high concurrency

Current implementations use a combination of Compare-And-Swap operations and volatile puts and gets

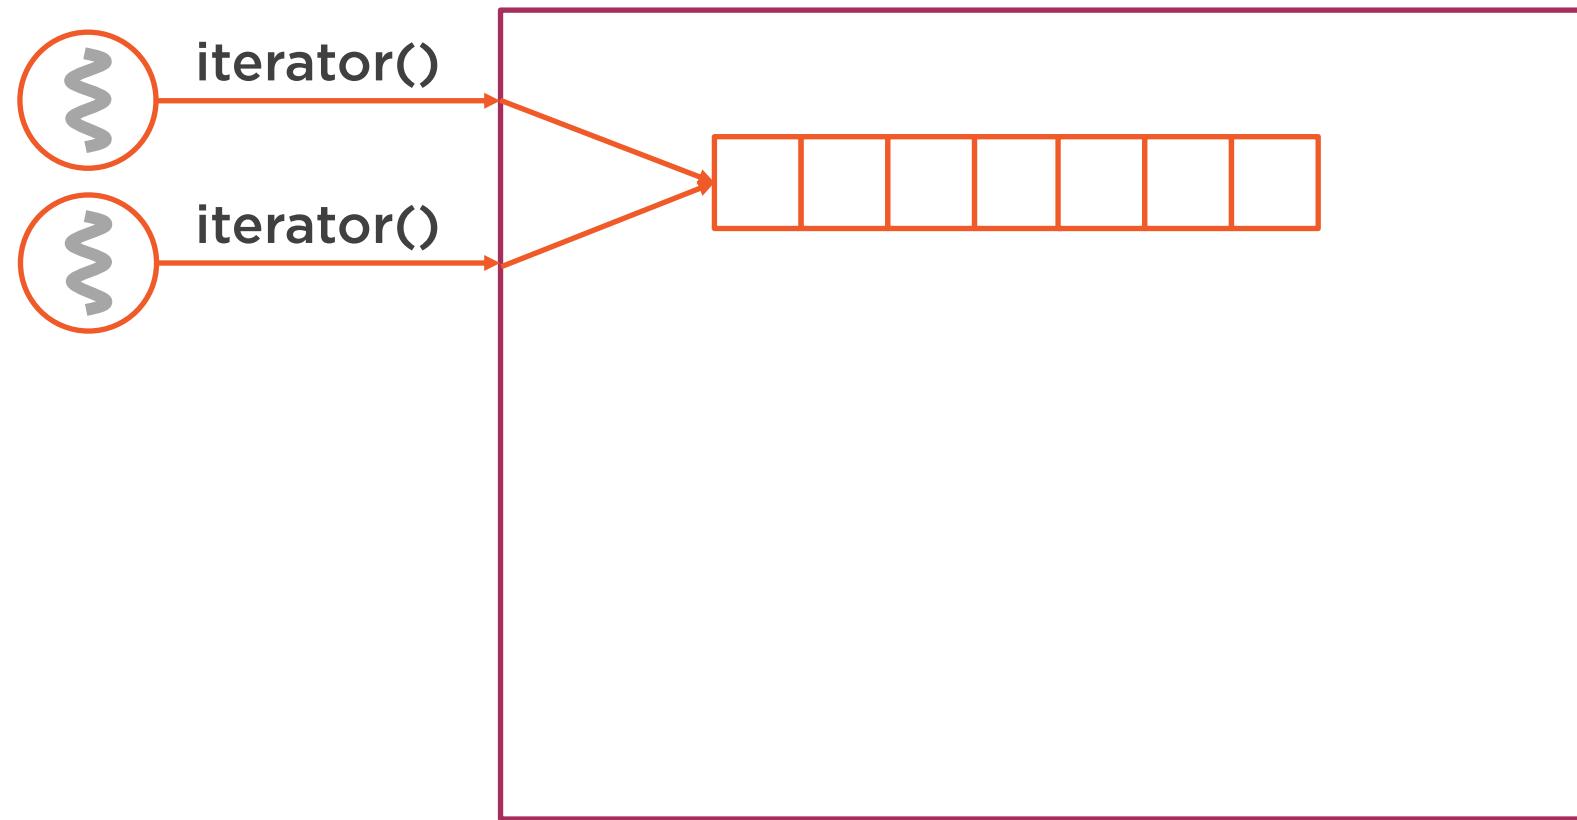
Drop-in replacement for synchronized maps



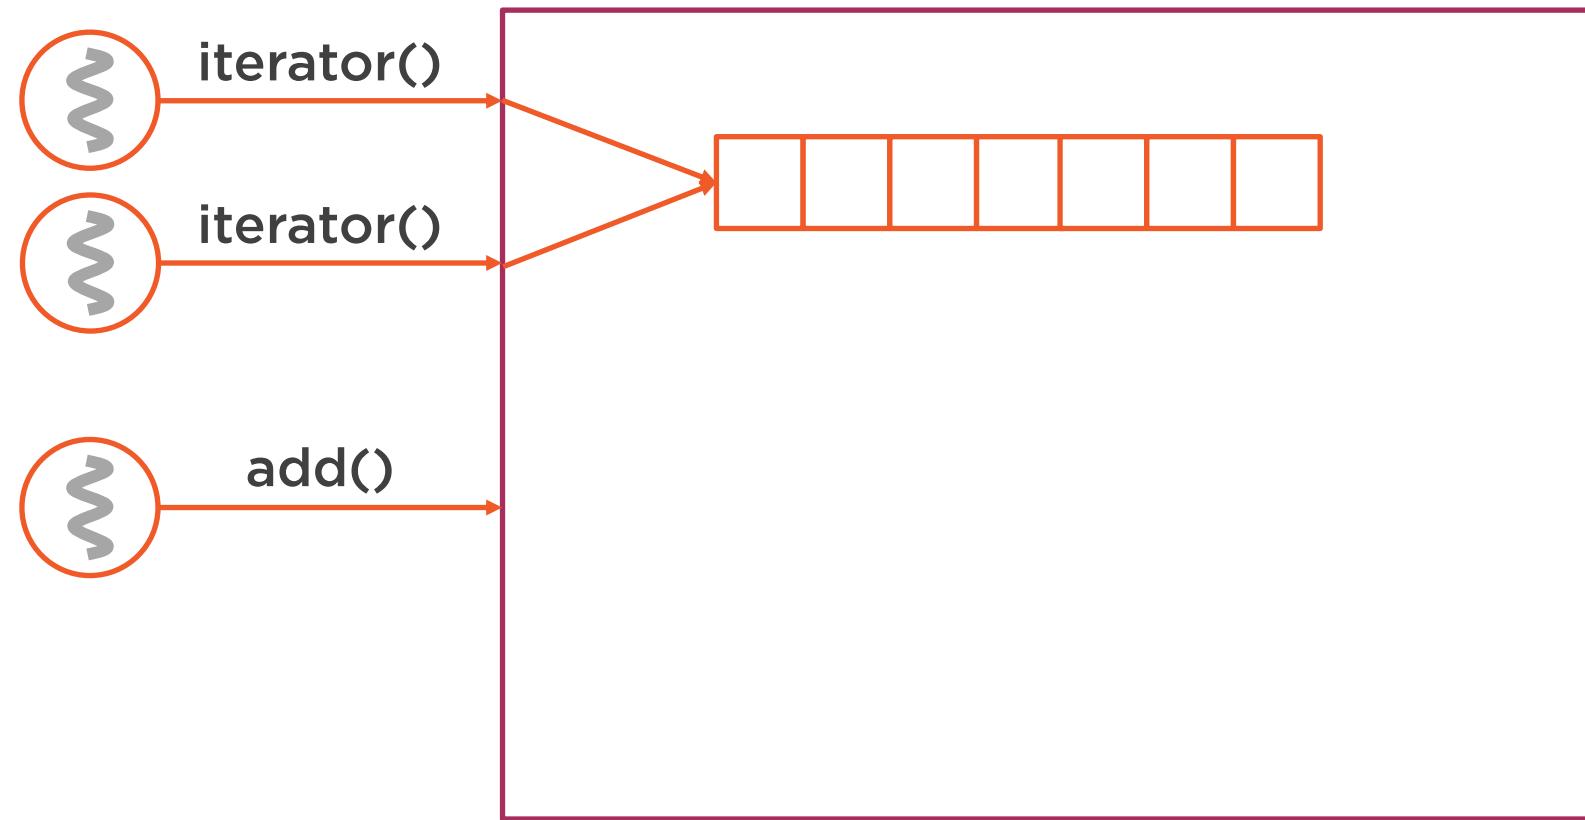
# CopyOnWriteArrayList



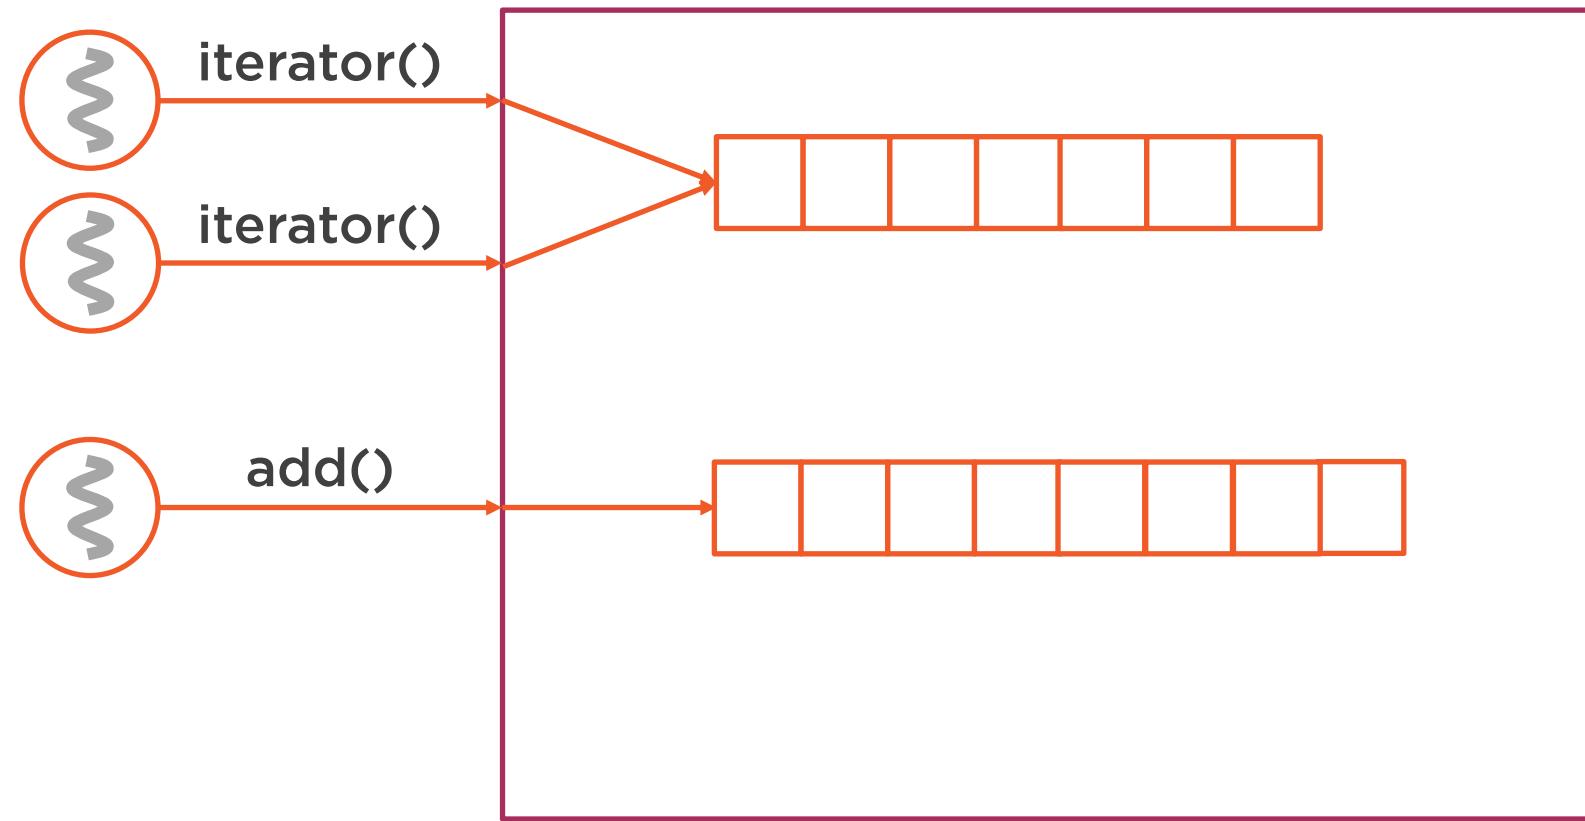
# CopyOnWriteArrayList



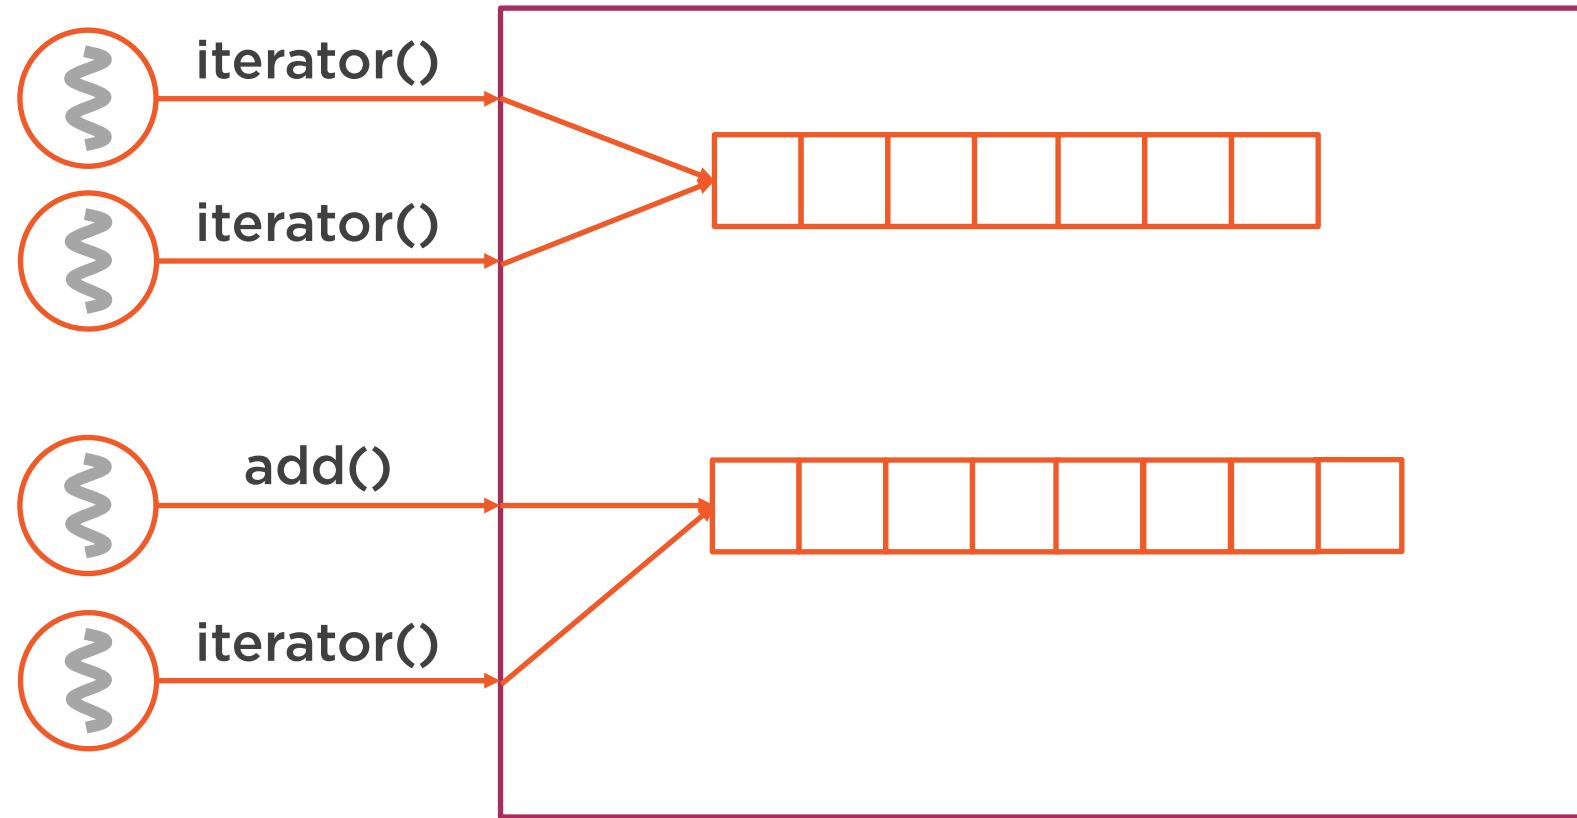
# CopyOnWriteArrayList



# CopyOnWriteArrayList



# CopyOnWriteArrayList



# CopyOnWriteArrayList Alternative

**CopyOnWriteArrayList doesn't fit your use case**

**Would like to avoid using a single lock over the entire list**

**Don't need index based access or ability to insert in the middle of the list**

**ConcurrentLinkedQueue**



```
ConcurrentLinkedQueue<String> list = new  
ConcurrentLinkedQueue<>();  
  
// adding to the queue  
list.add("Some String")  
  
// iterating the queue  
for (String str : list) {  
    System.out.println(str);  
}  
}
```



# Avoiding Synchronization

---



# The volatile Keyword

Let's the compiler and Java runtime know that this variable is shared

**Compiler:**

- Will not re-order operations surrounding the variables

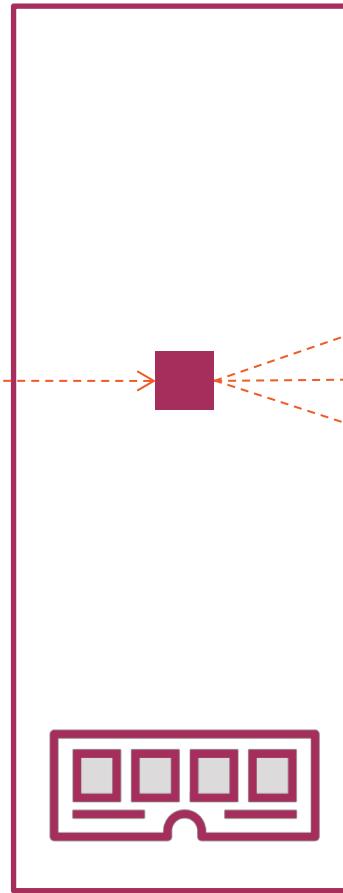
**Runtime:**

- Volatiles are directly read from and written to memory and are not cached
- When a thread writes to a volatile variable, all other variables visible to it are flushed to memory





**direct write to memory**



# Volatile Variables

**Most commonly used as a completion, interruption or status flag**

**Can be used when:**

- **Updates to a single variable are only being done by one thread**
- **You have multiple writer threads but writes to the variable do not depend on its current value**



# Immutable Objects

**Do not need synchronized access**

**Once constructed, their state cannot be changed**

**The more you use immutable classes when writing concurrent code, the less synchronization and locking that's needed**



# ThreadLocal

Instead of sharing, each thread has their own copy of an object



```
public class MySerializer {  
    Kryo kryoSerializer = new Kryo(); // shared unsafe object  
    kryoSerializer.register(SomeClass.class);  
  
    public synchronized byte[] serializer(SomeClass object) {  
        Output output =  
            new Output(new ByteArrayOutputStream(), 100);  
        kryoSerializer.writeObject(output, object);  
        return output.toBytes();  
    }  
}
```

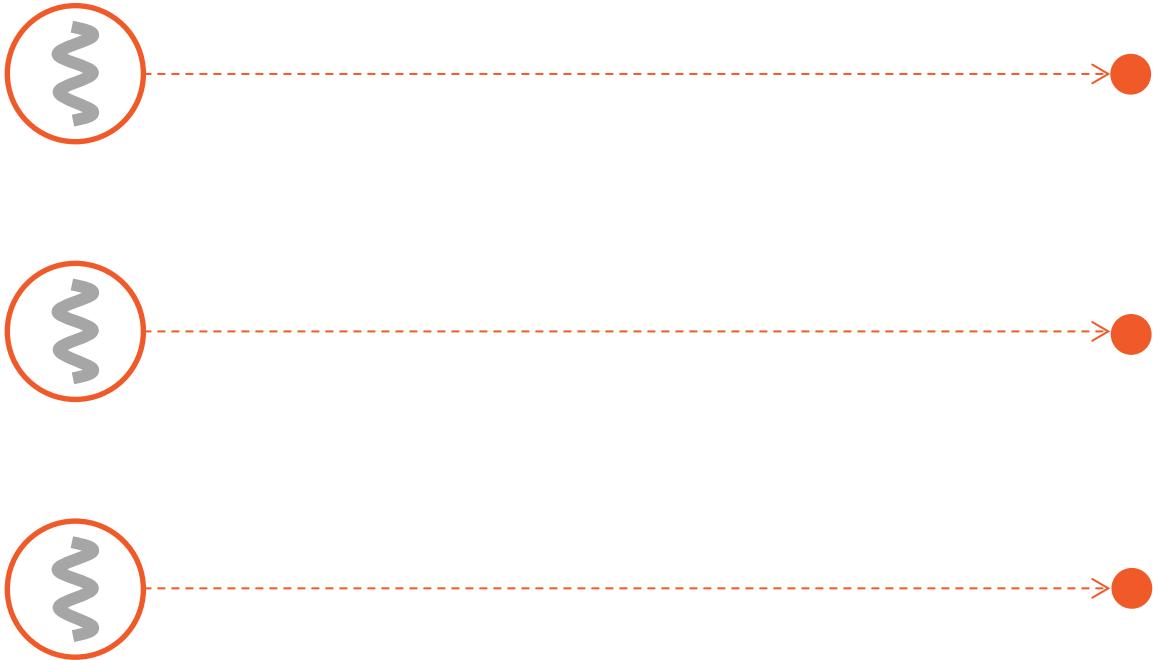


```
public class MySerializer {  
    ThreadLocal<Kryo> kryoSerializers = new ThreadLocal<Kryo>() {  
        protected Kryo initialValue() {  
            Kryo kryoSerializer = new Kryo();  
            kryoSerializer.register(SomeClass.class);  
            return kryo; }  
    };  
    public byte[] serializer(SomeClass object) {  
        Output output = new Output(new ByteArrayOutputStream(), 100);  
        kryoSerializers.get().writeObject(output, object);  
        return output.toBytes();  
    }  
}
```



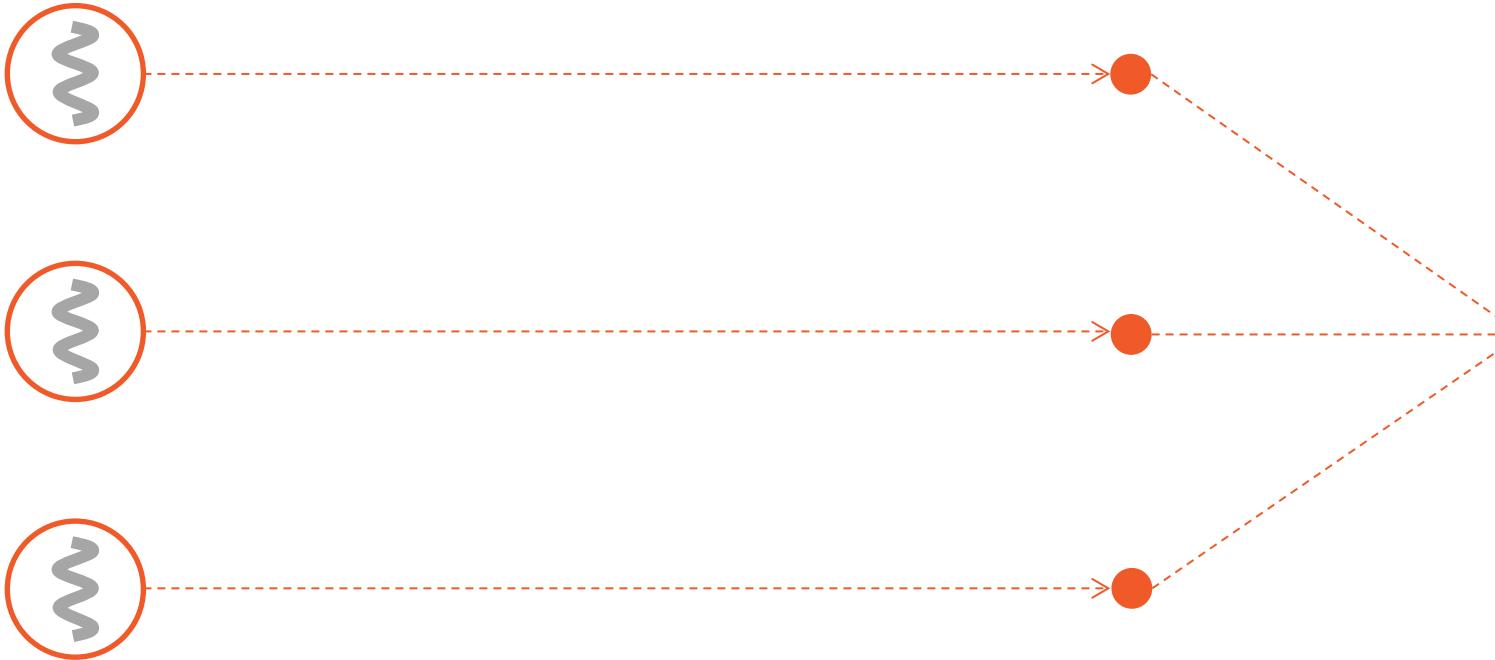
```
public class MySerializer {  
    ThreadLocal<Kryo> kryoSerializers = new ThreadLocal<Kryo>() {  
        protected Kryo initialValue() {  
            Kryo kryoSerializer = new Kryo();  
            kryoSerializer.register(SomeClass.class);  
            return kryo; }  
    };  
    public byte[] serializer(SomeClass object) {  
        Output output = new Output(new ByteArrayOutputStream(), 100);  
        kryoSerializers.get().writeObject(output, object);  
        return output.toBytes();  
    }  
}
```





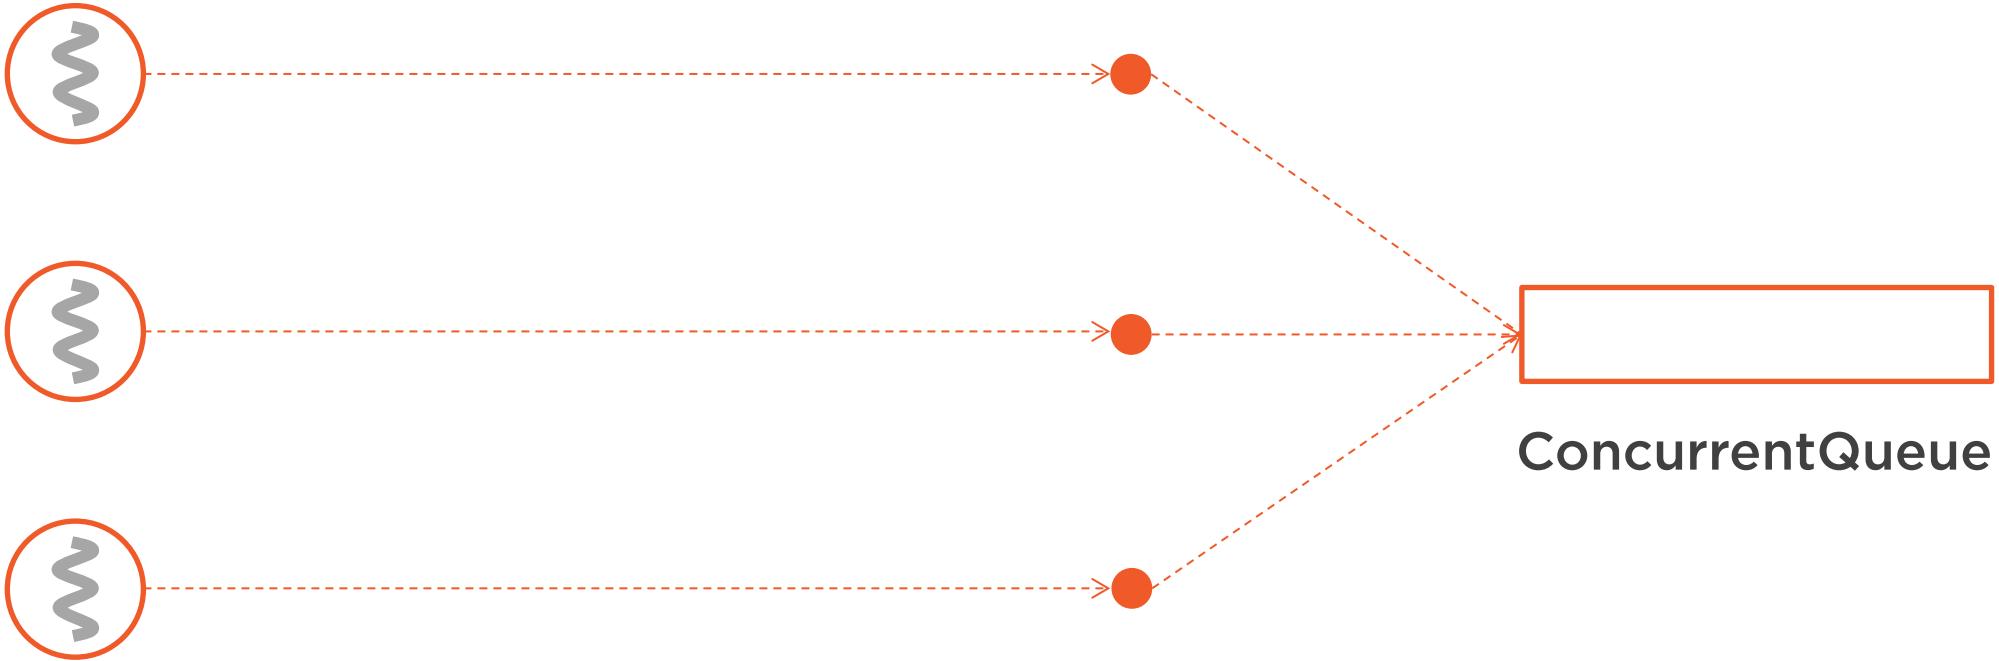
Parallel shared-nothing execution





Parallel shared-nothing execution





Parallel shared-nothing execution



# Concurrent Queues

## Non-Blocking Queues

`ConcurrentLinkedQueue`

`ConcurrentLinkedDeque`

## Blocking Queues

`ArrayBlockingQueue`

`LinkedBlockingQueue`

`LinkedBlockingDeque`



# Concurrent Queues

## Non-Blocking Queues

`ConcurrentLinkedQueue`

`ConcurrentLinkedDeque`

## Blocking Queues

`ArrayBlockingQueue`

`LinkedBlockingQueue`

`LinkedBlockingDeque`

Use when:



# Concurrent Queues

## Non-Blocking Queues

ConcurrentLinkedQueue

ConcurrentLinkedDeque

## Blocking Queues

ArrayBlockingQueue

LinkedBlockingQueue

LinkedBlockingDeque

### Use when:

- You want the consuming thread to wait for a period of time for a message to arrive



# Concurrent Queues

## Non-Blocking Queues

ConcurrentLinkedQueue

ConcurrentLinkedDeque

## Blocking Queues

ArrayBlockingQueue

LinkedBlockingQueue

LinkedBlockingDeque

### Use when:

- You want the consuming thread to wait for a period of time for a message to arrive
- You want to control the growth rate of the queue



# Resources

**Online articles and blog posts**

**Java Concurrency in Practice**

**Java Performance: The Definitive Guide**

