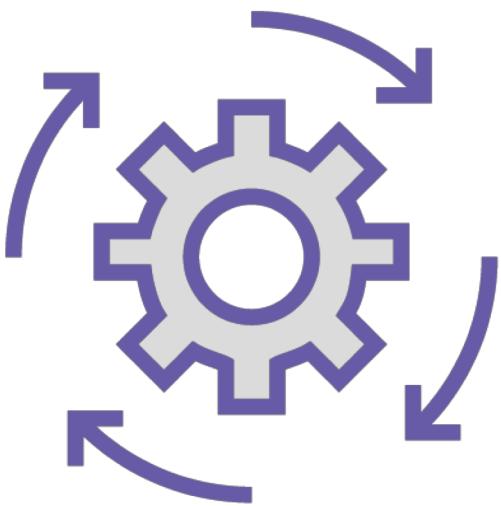


Language Types

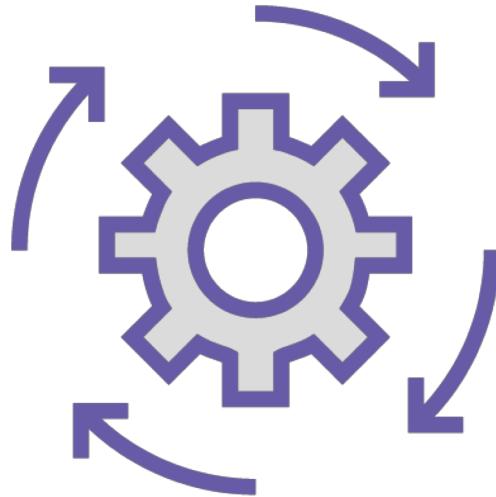


Compiled



Interpreted



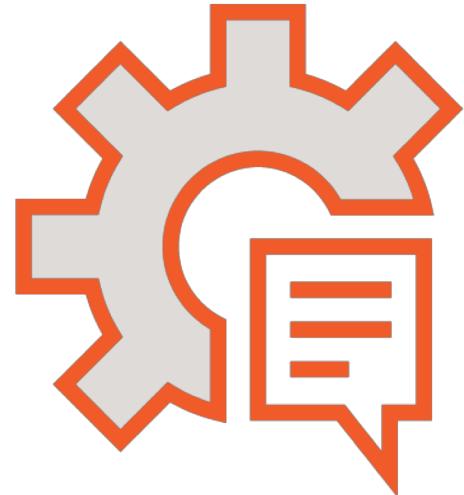


Compiled

Transforms program code into binary instructions

Compiled binaries target specific CPU architectures





Interpreted

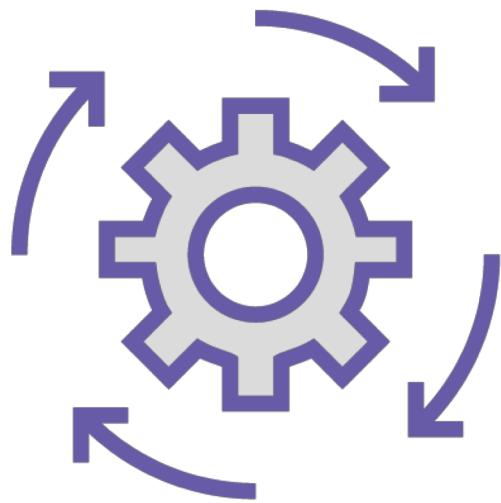
Translates each line of code into binary instructions as the line is executed

Is portable

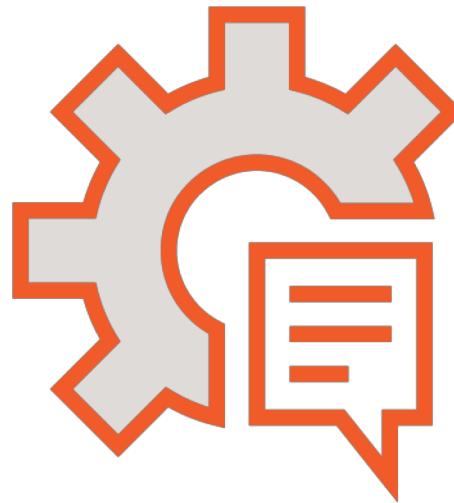
Is not as fast as compiled languages



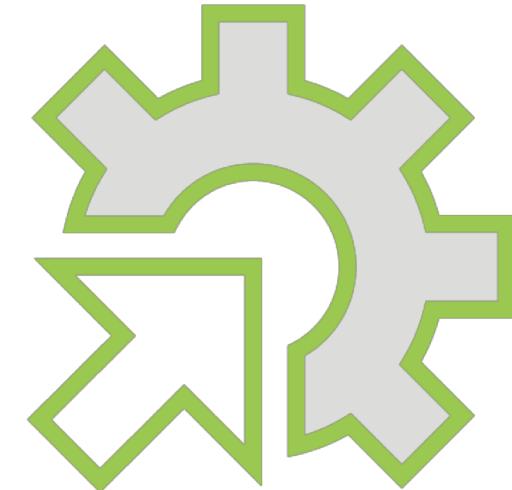
Language Types



Compiled



Interpreted



Intermediate



Intermediate Language Benefits

Compiler can perform type checking and optimizations

Bytecode is compiled only once, and is close to machine code

Maintains portability



Just-In-Time Compiler

Introduced in 1999

Identifies hot methods (“hotspots”) and
compiles them



Overview



JIT Compilation Modes

JIT Tuning Flags

Garbage Collection Intro

Measuring Garbage Collection Performance

Garbage Collection Tuning



JIT Compilation Modes



JIT Compiler Types



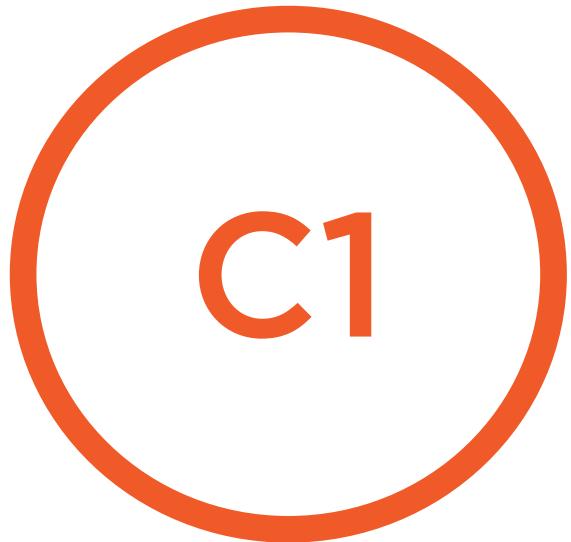
Compiler 1 / Client Compiler



Compiler 2 / Server Compiler



JIT Compiler Types

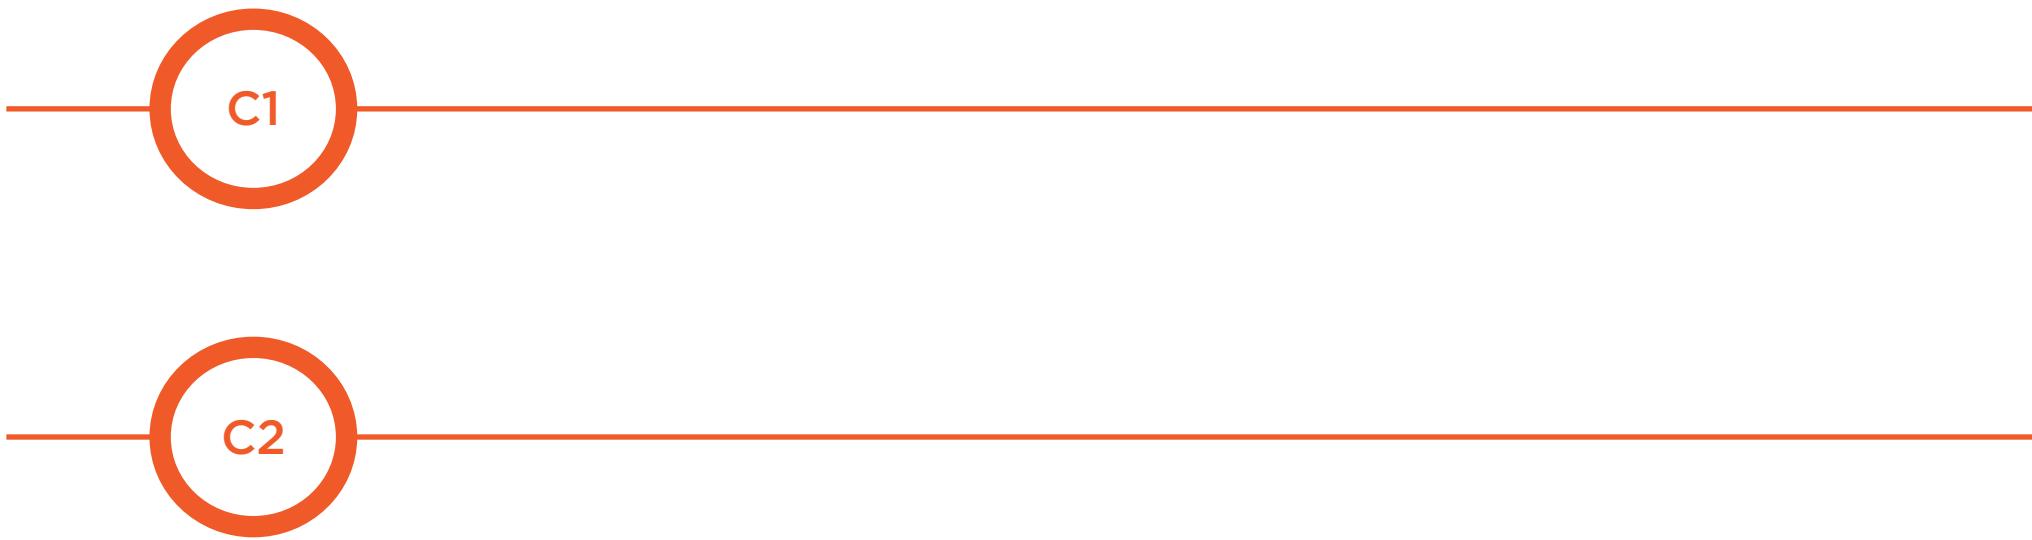


Compiler 1 / Client Compiler
Optimized for startup performance

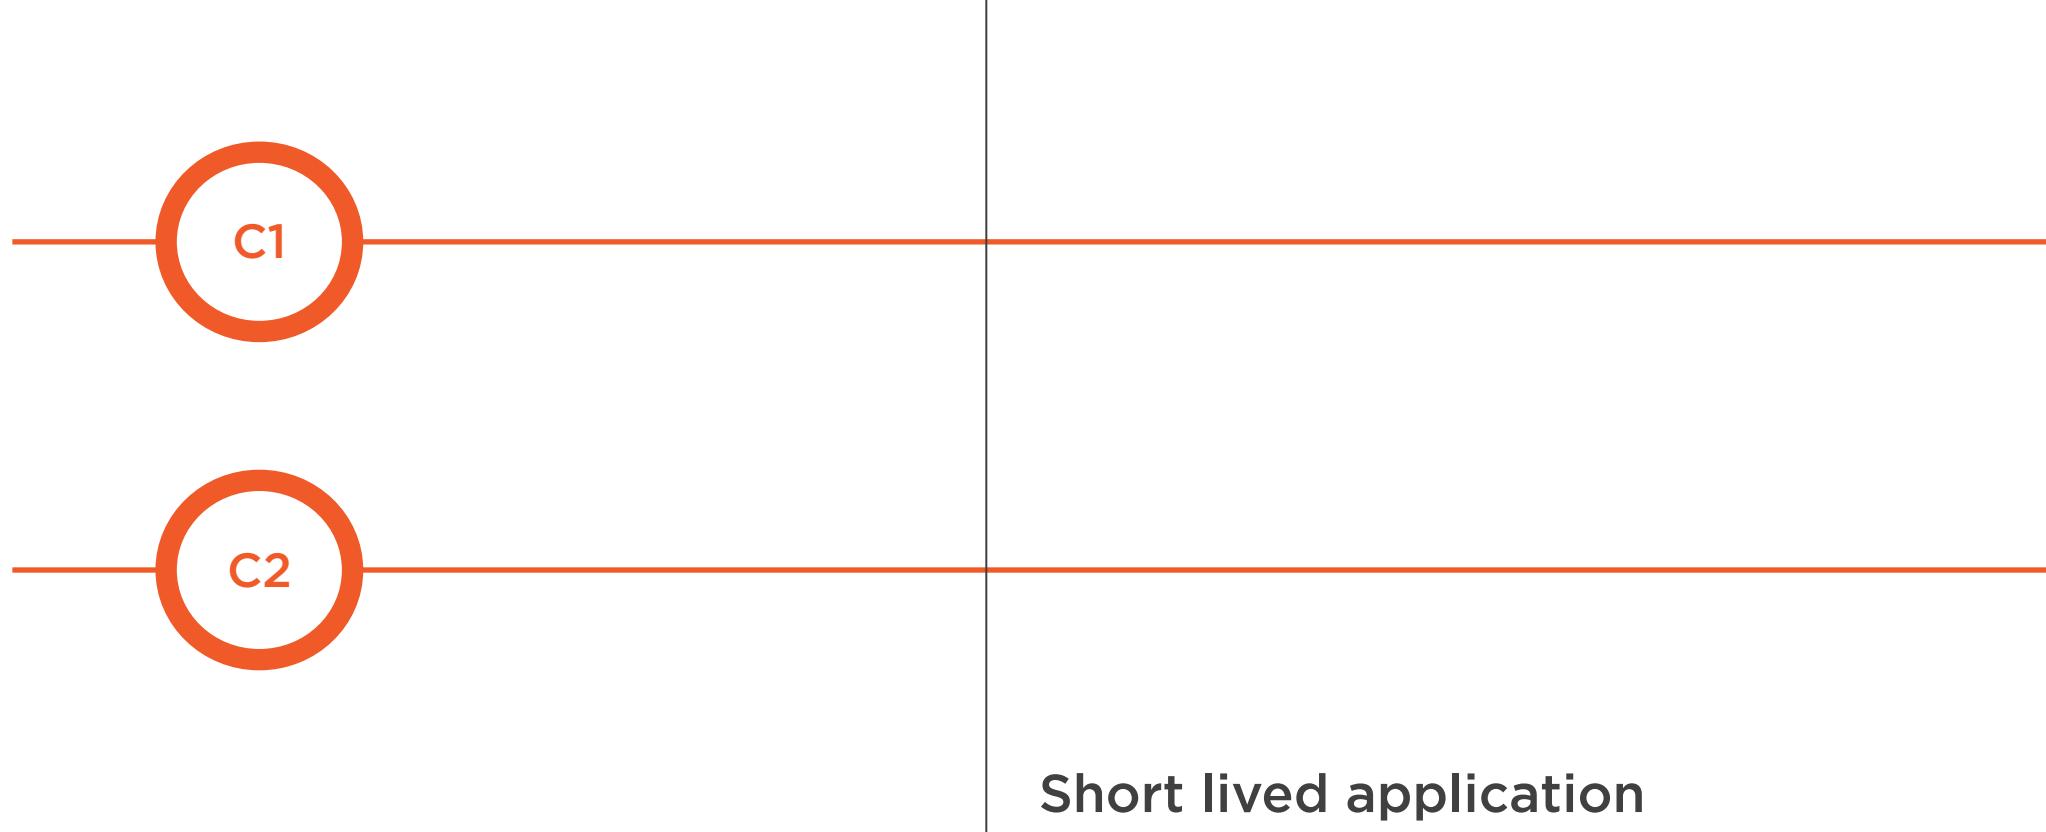


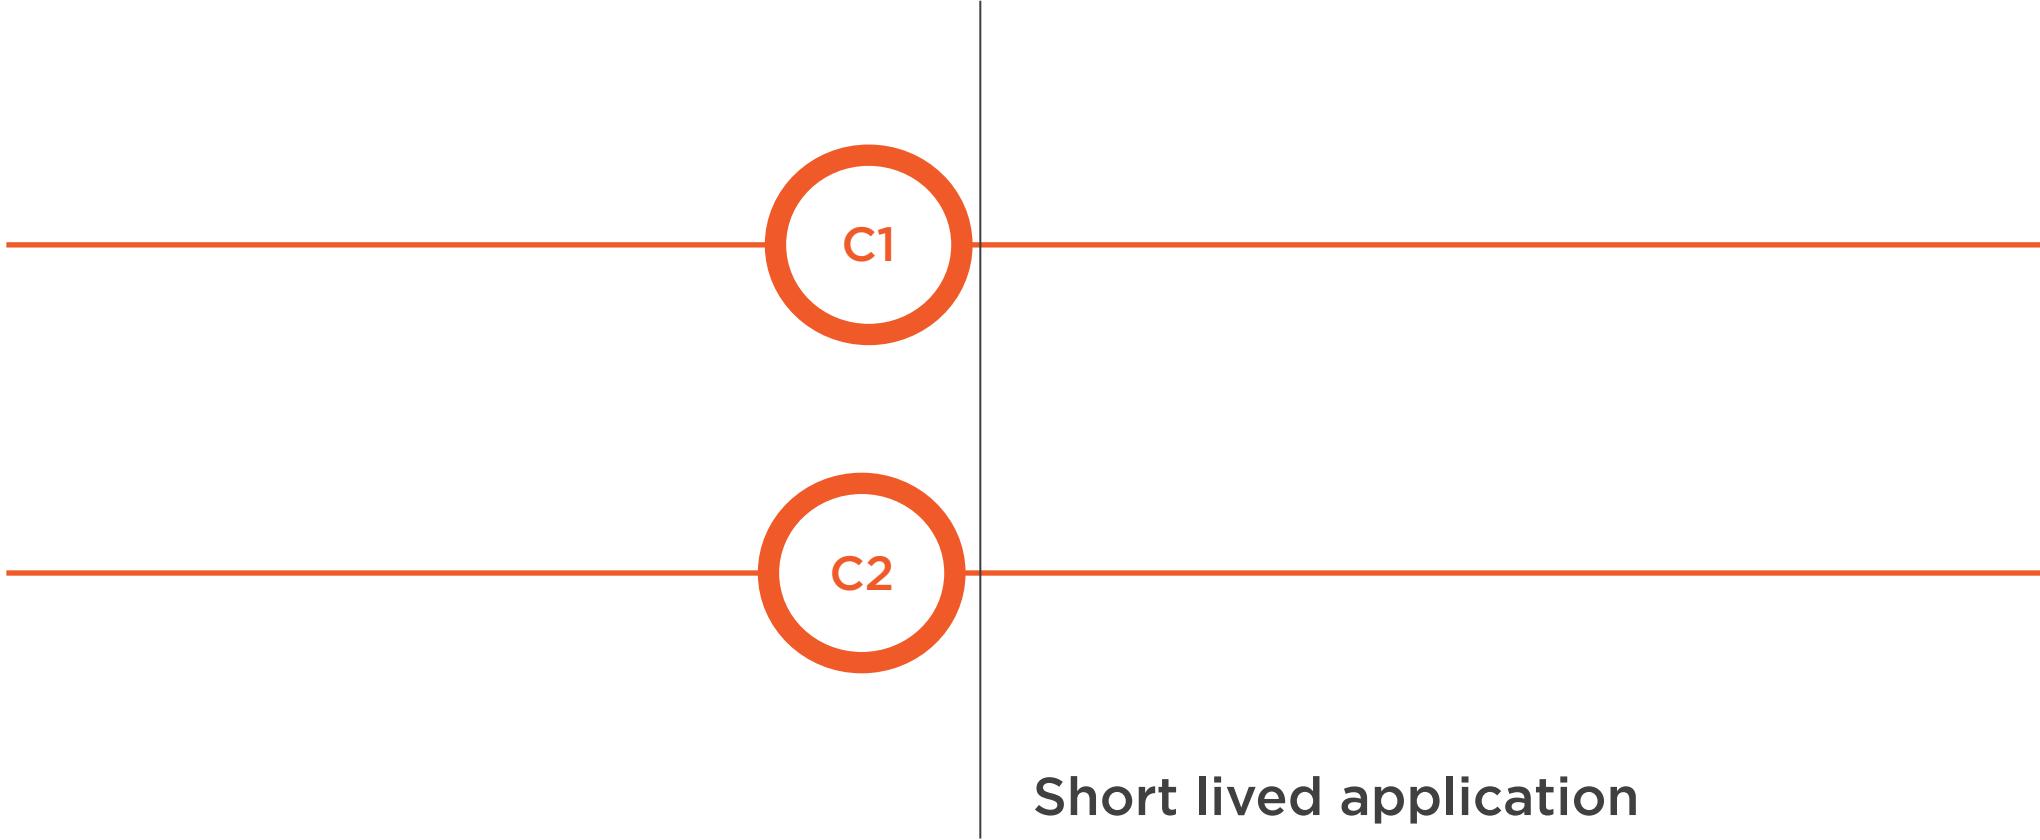
Compiler 2 / Server Compiler
Waits and profiles so that it can
apply more aggressive optimizations



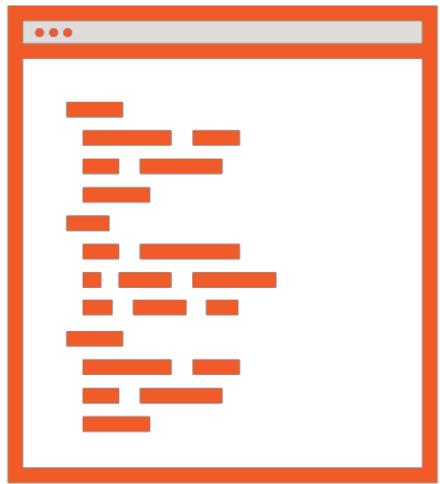








Tiered Compilation



Just-In-Time Compiler

Introduced in Java 7

**Hot methods are first compiled with C1,
and then C2 as they get hotter**

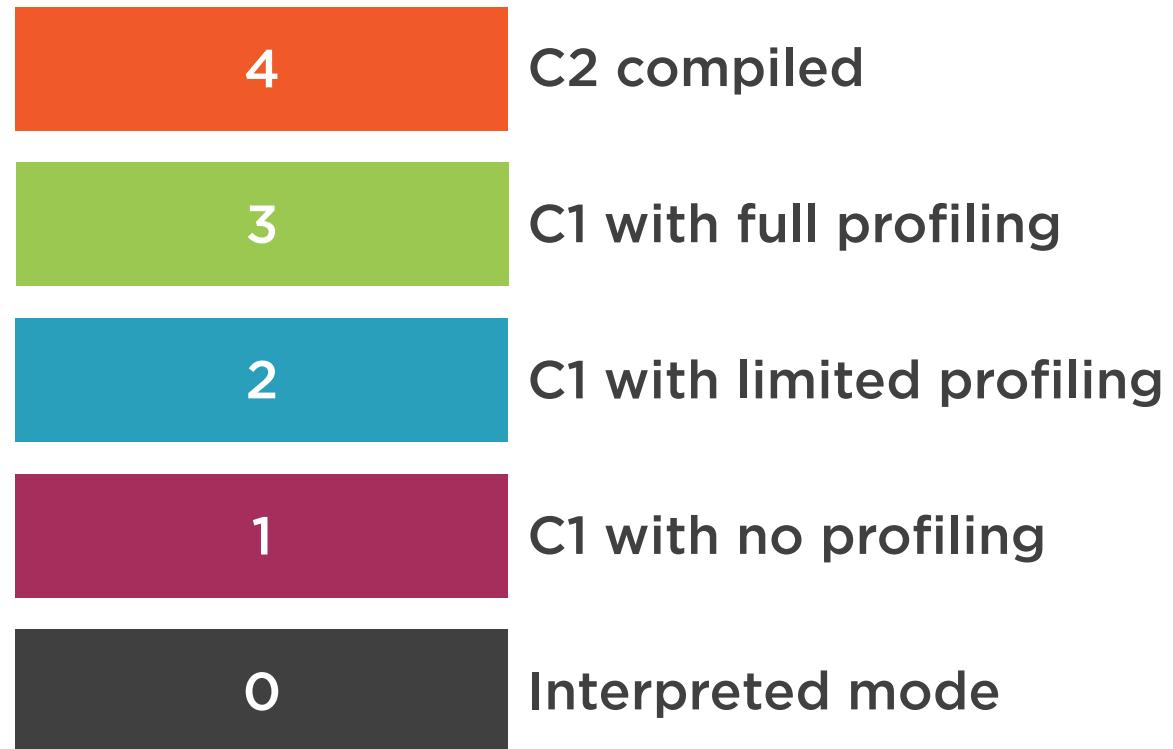
Set as default in Java 8



Tiered Compilation Levels



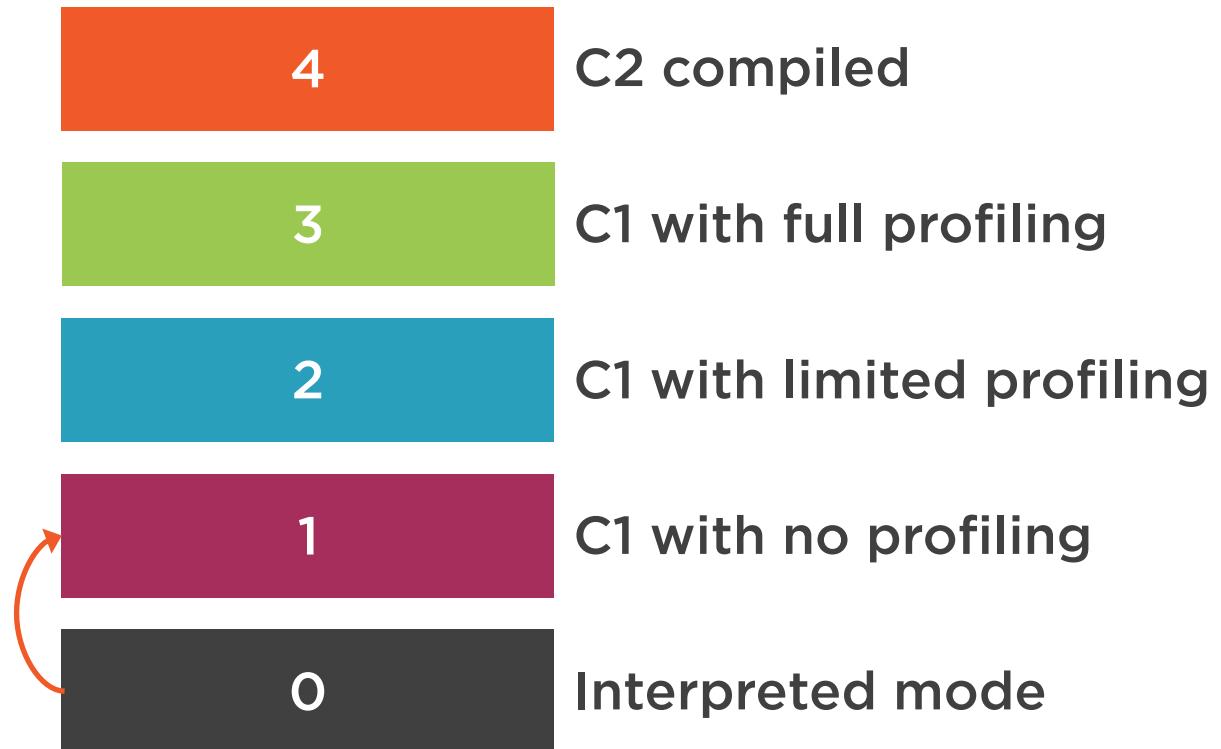
Tiered Compilation Levels



Tiered Compilation Levels



Tiered Compilation Levels



Tiered Compilation Levels



Choosing Your Compilation Mode

Start with **Tiered
Compilation**

Test with **C1** or

Test with **C2**

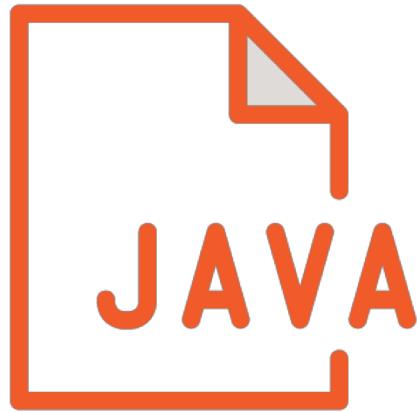


Choosing Your Compilation Mode

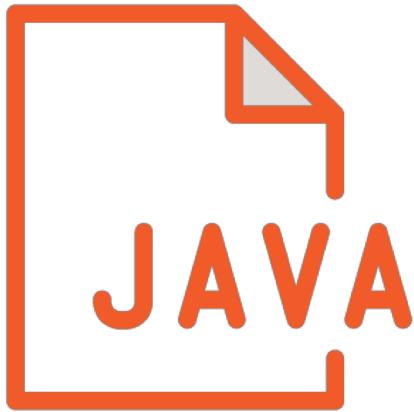
Start with **Tiered
Compilation**

Test with **C2**



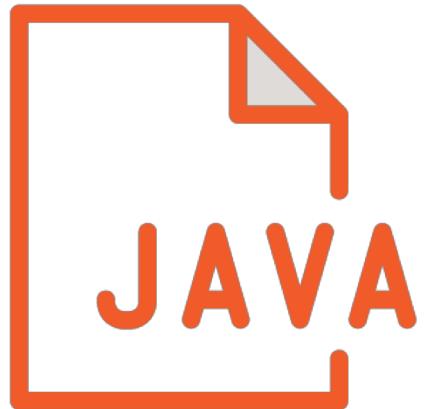


Java 7

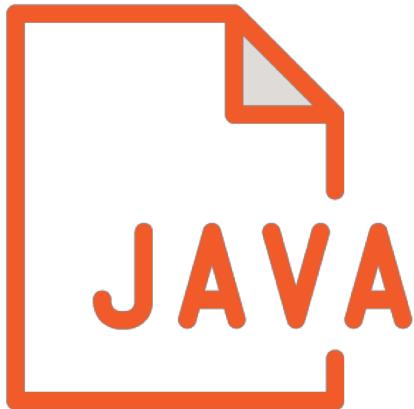


Java 8





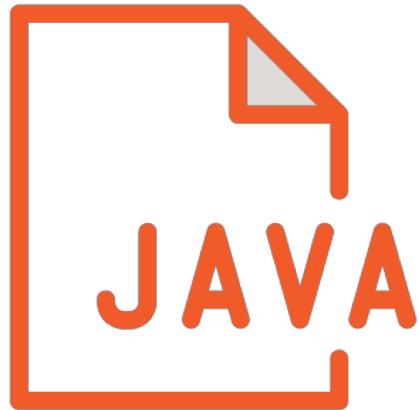
Java 7



Java 8

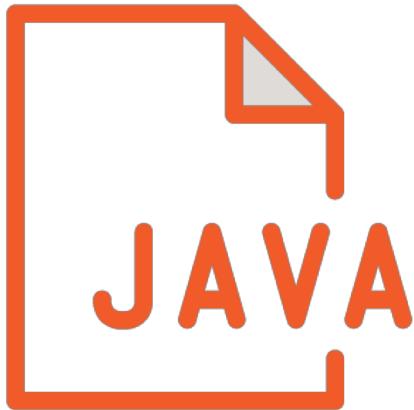
Default: server mode





Java 7

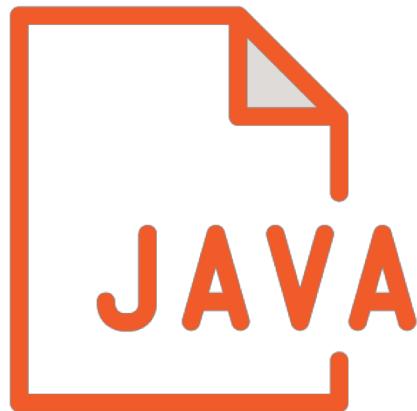
Default: server mode



Java 8

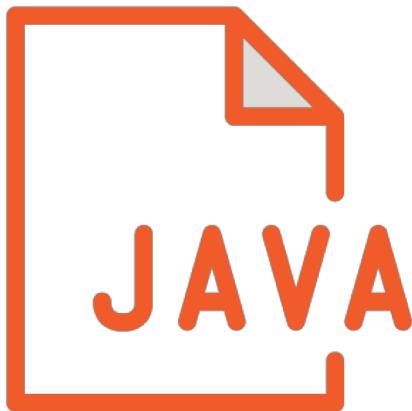
Default: tiered mode





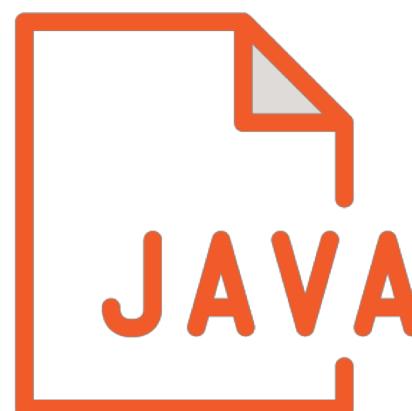
Java 7

Default: server mode



Java 8

Default: tiered mode



Java 8

Set: server mode



JIT Tuning Flags



```
java -server
```

Setting the Compilation Mode in Java 7



```
java -client
```

Setting the Compilation Mode in Java 7



```
java -XX:TieredStopAtLevel=1
```

Setting the Client Compilation Mode in Java 8

Available levels

- **Level 0: Interpreted mode**
- **Level 1: C1 compiled with no profiling**



```
java -XX:-TieredCompilation
```

Setting the Server Compilation Mode in Java 8



Hot Method Definition

i: invocation counter

The number of times a
method has been called

b: backedge counter

The number of times any
loops in a method have
branched back



Hot Method Definition – Tiered Compilation

```
i > Tier3InvocationThreshold ||  
(i > Tier3MinInvocationThreshold && i + b >  
Tier3CompileThreshold)
```



Hot Method Definition – Tiered Compilation

```
i > 200 ||  
(i > 100 && i + b > 2000)
```



Hot Method Definition – Tiered Compilation

```
i > 5000 ||  
(i > 600 && i + b > 15000)
```



```
java -XX:Tier4InvocationThreshold=4000 -  
XX:Tier4CompileThreshold=10000
```

Lowering the Compilation Thresholds in Tiered Mode



Hot Method Definition – Server Compilation

$i + b > \text{CompileThreshold}$



Hot Method Definition – Server Compilation

$i + b > 10000$



```
java -XX:-TieredCompilation -XX:CompileThreshold=8000
```

Lowering the Compilation Thresholds in Server Mode

- Methods get compiled earlier
- “Warm” methods get compiled



Code Cache

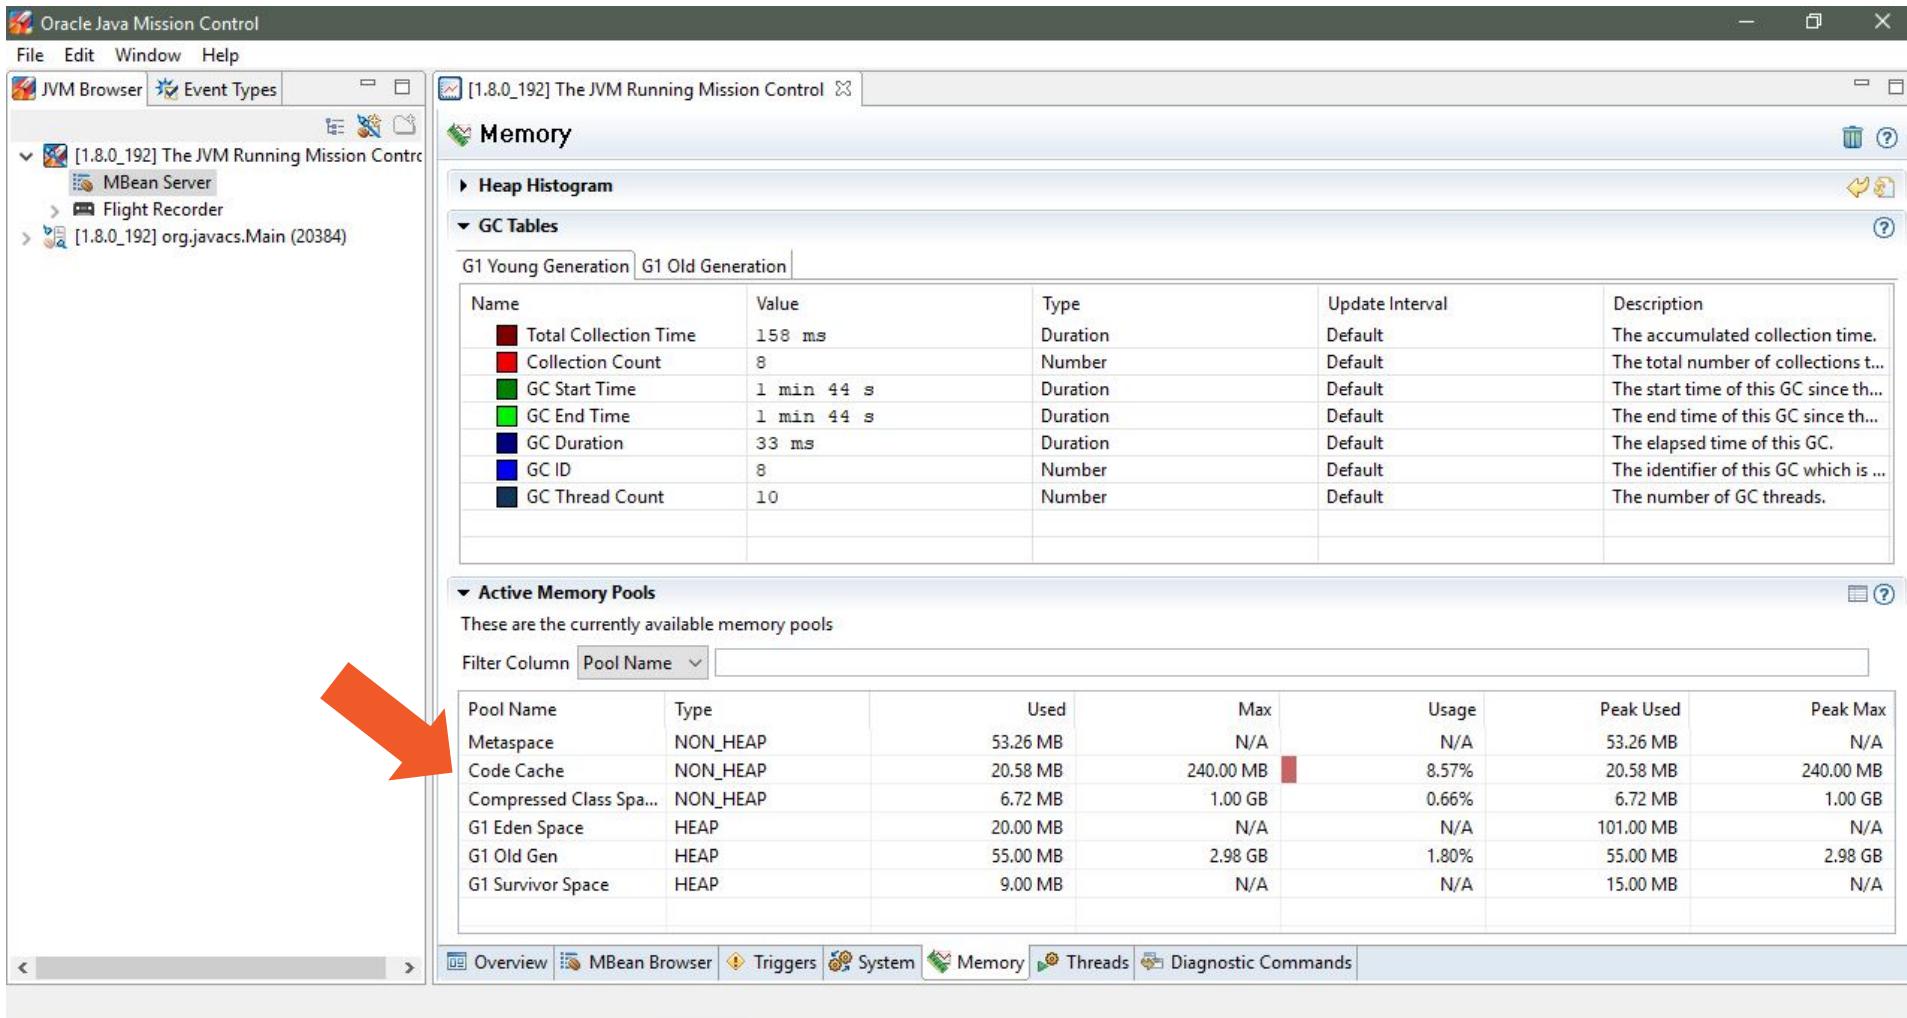
An area of native memory where compiled code is stored for future executions

If the code cache is filled up no more JIT compilation is done!

In Java 8 and up the reserved code cache size is usually sufficient by default



Code Cache Monitoring



The screenshot shows the Oracle Java Mission Control interface, specifically the Memory tab. The left sidebar shows a tree view with nodes like 'MBean Server', 'Flight Recorder', and '[1.8.0_192] org.javacs.Main (20384)'. The main panel displays two tables: 'GC Tables' and 'Active Memory Pools'. A red arrow points to the 'Code Cache' row in the 'Active Memory Pools' table.

GC Tables

Name	Value	Type	Update Interval	Description
Total Collection Time	158 ms	Duration	Default	The accumulated collection time.
Collection Count	8	Number	Default	The total number of collections ...
GC Start Time	1 min 44 s	Duration	Default	The start time of this GC since th...
GC End Time	1 min 44 s	Duration	Default	The end time of this GC since th...
GC Duration	33 ms	Duration	Default	The elapsed time of this GC.
GC ID	8	Number	Default	The identifier of this GC which is ...
GC Thread Count	10	Number	Default	The number of GC threads.

Active Memory Pools

Pool Name	Type	Used	Max	Usage	Peak Used	Peak Max
Metaspace	NON_HEAP	53.26 MB	N/A	N/A	53.26 MB	N/A
Code Cache	NON_HEAP	20.58 MB	240.00 MB	8.57%	20.58 MB	240.00 MB
Compressed Class Spa...	NON_HEAP	6.72 MB	1.00 GB	0.66%	6.72 MB	1.00 GB
G1 Eden Space	HEAP	20.00 MB	N/A	N/A	101.00 MB	N/A
G1 Old Gen	HEAP	55.00 MB	2.98 GB	1.80%	55.00 MB	2.98 GB
G1 Survivor Space	HEAP	9.00 MB	N/A	N/A	15.00 MB	N/A



```
java -XX:ReservedCodeCacheSize=<N>
```

Setting the Code Cache Size



Java 9 Code Cache Changes

JVM internal
code

Profiled (lightly
optimized) code

Non-profiled
(fully
optimized) code



Java 9 Code Cache Changes

JVM internal
code

Profiled (lightly
optimized) code

Non-profiled
(fully
optimized) code



Java 9 Code Cache Changes

JVM internal
code

Profiled (lightly
optimized) code

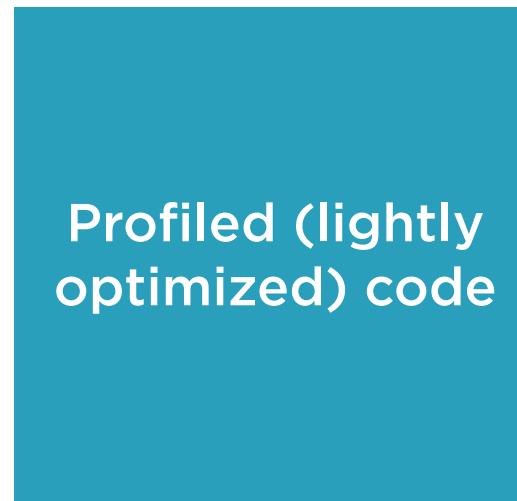
Non-profiled
(fully
optimized) code



Java 9 Code Cache Changes



NonNMethodCodeHeapSize



ProfiledCodeHeapSize



NonProfiledCodeHeapSize



```
java -XX:ReservedCodeCacheSize=<N>
```

Setting the Code Cache Size in Java 9

**Will turn off the segmented code cache feature in Java 9
and above**



Introduction to Garbage Collection



Garbage Collection

Mechanism by which the JVM reclaims memory on the behalf of the application when it's no longer needed

Consists of:

- Finding no longer reachable objects
- Deallocating the memory
- Compacting the heap



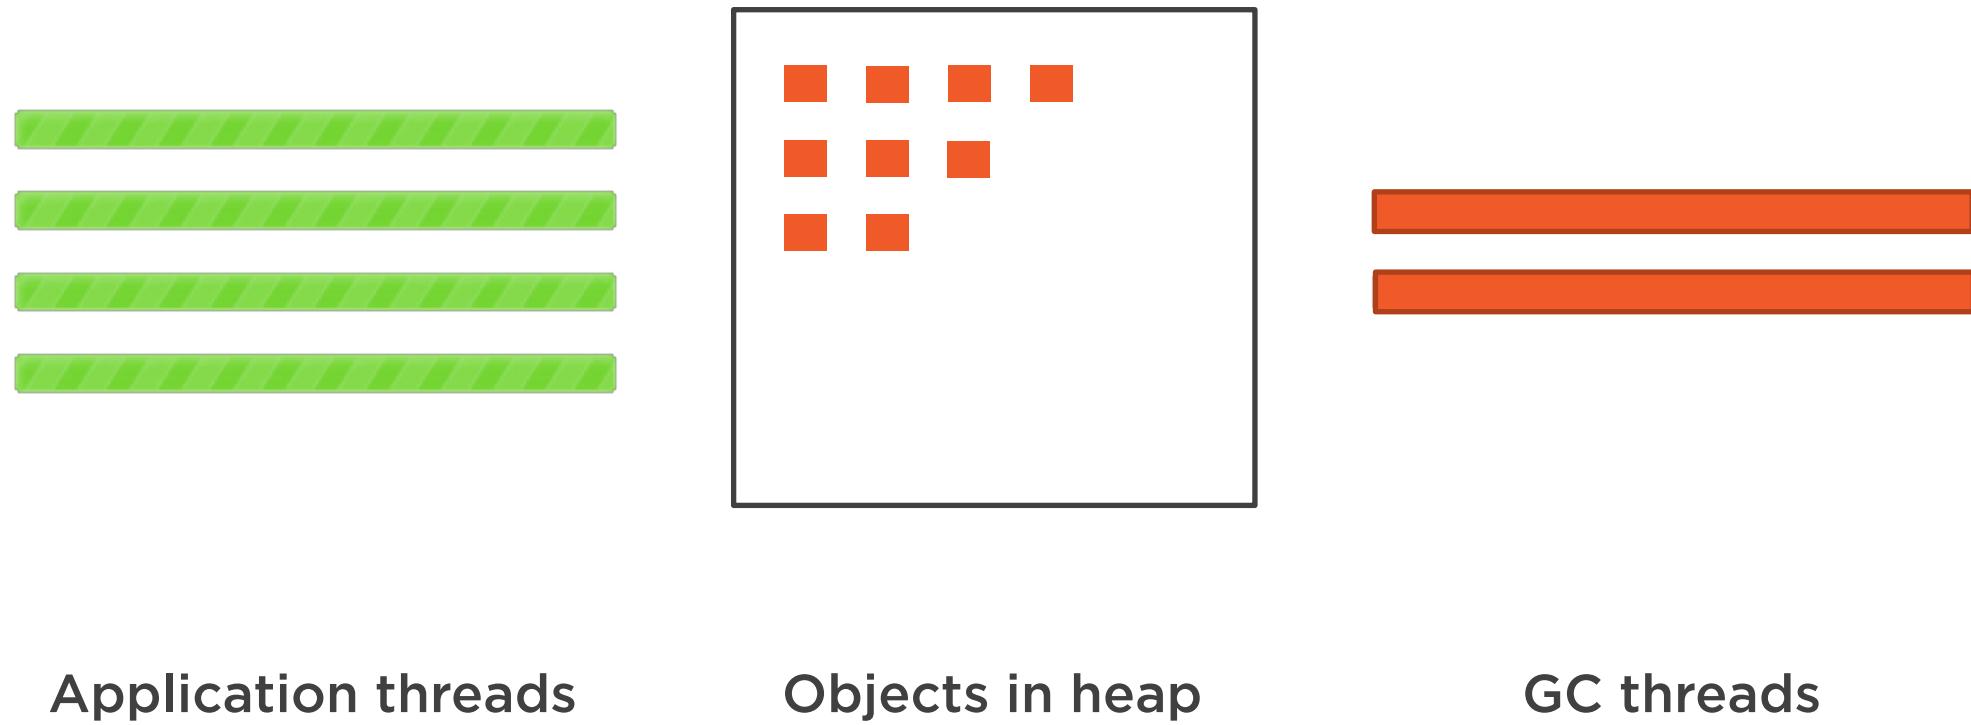
GC Pauses

GC threads are used to collect garbage and compact the heap

During this process application threads are typically stopped so that object references aren't lost



Stop-The-World Pauses



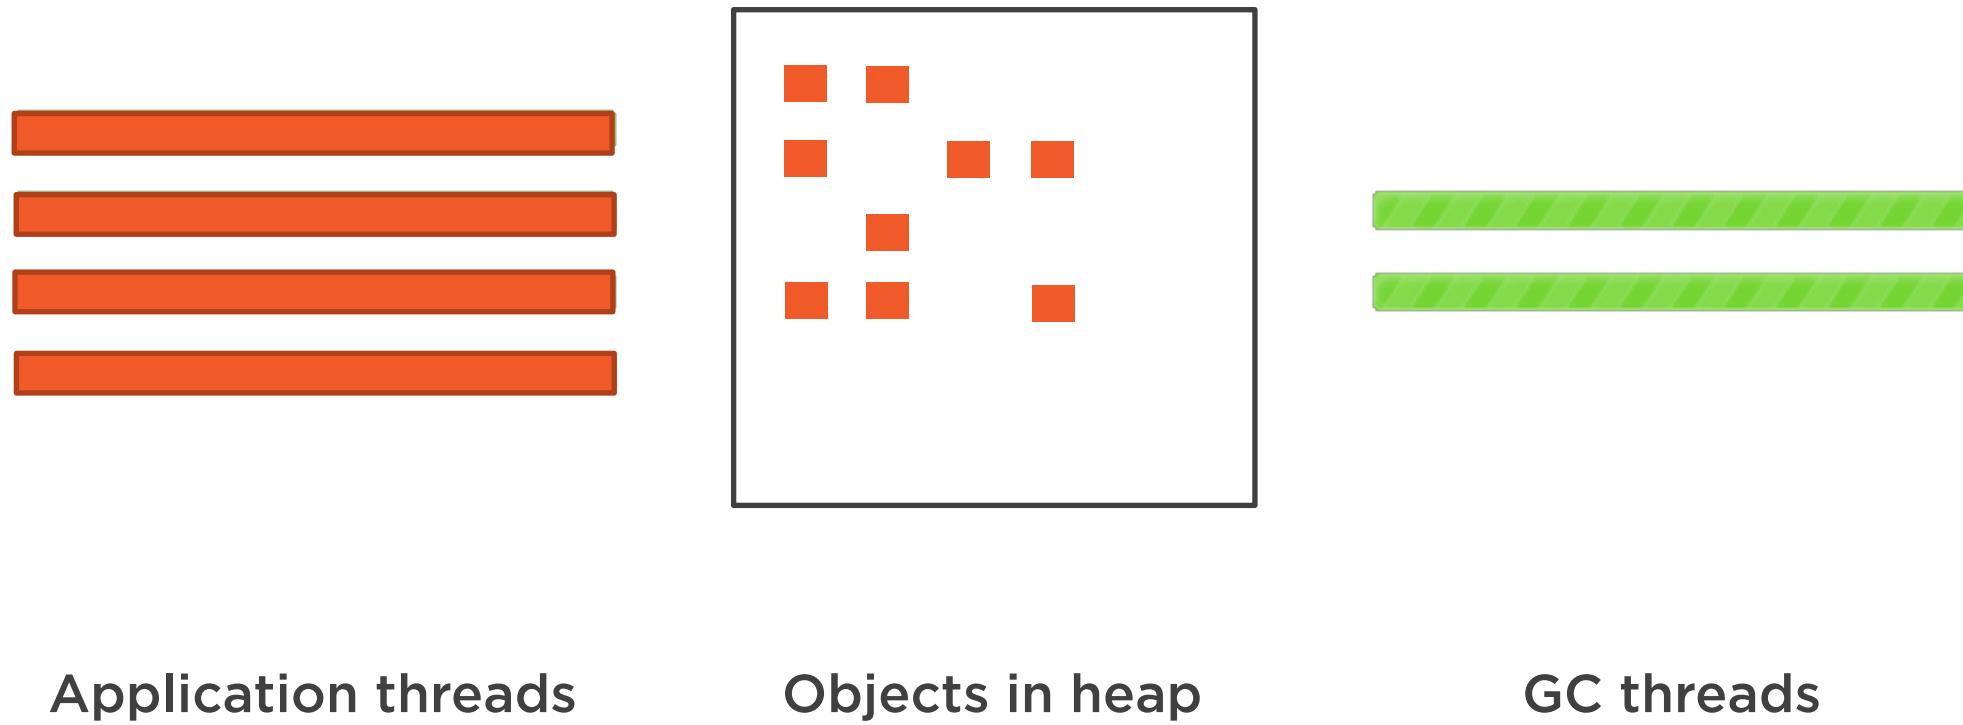
Application threads

Objects in heap

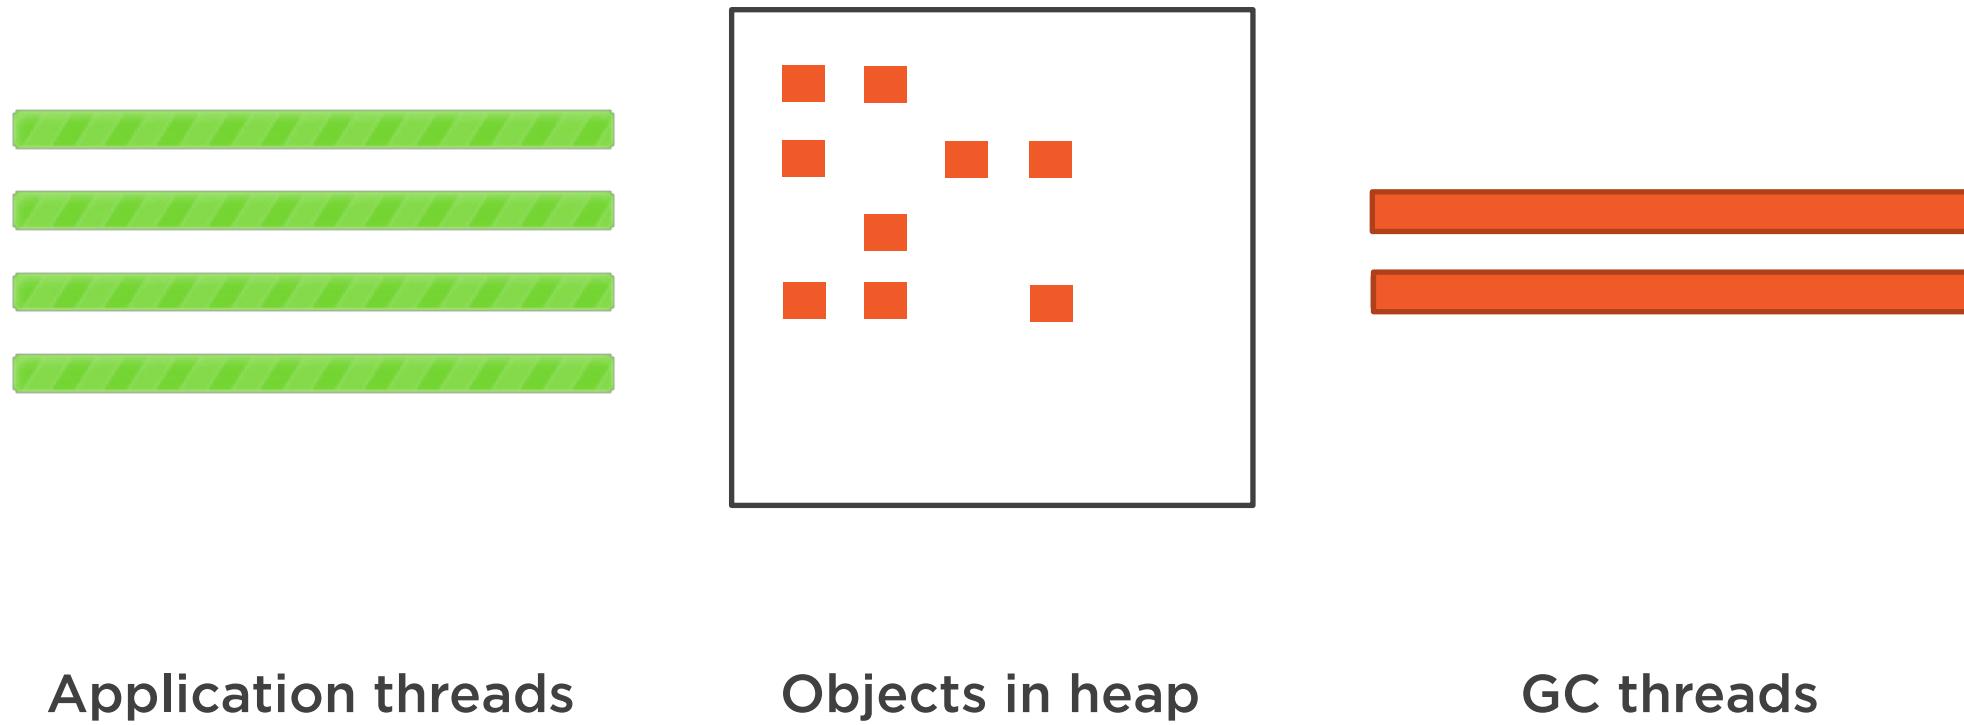
GC threads



Stop-The-World Pauses



Stop-The-World Pauses



OpenJDK Garbage Collectors

Serial

Parallel

Concurrent Mark Sweep

Garbage First



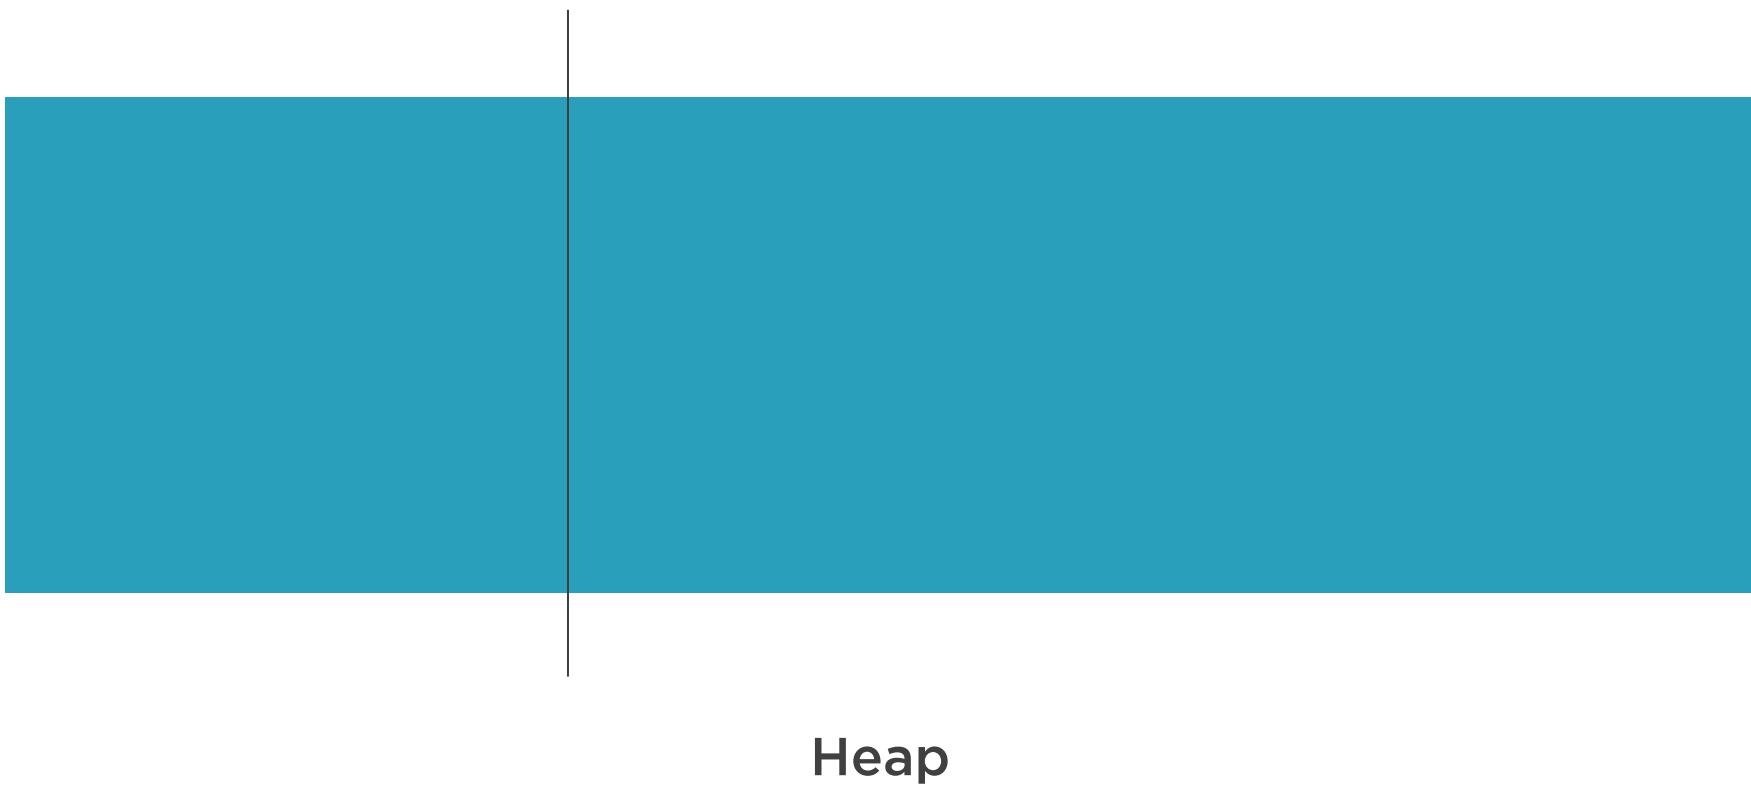
Generational Garbage Collection



Heap



Generational Garbage Collection



Generational Garbage Collection

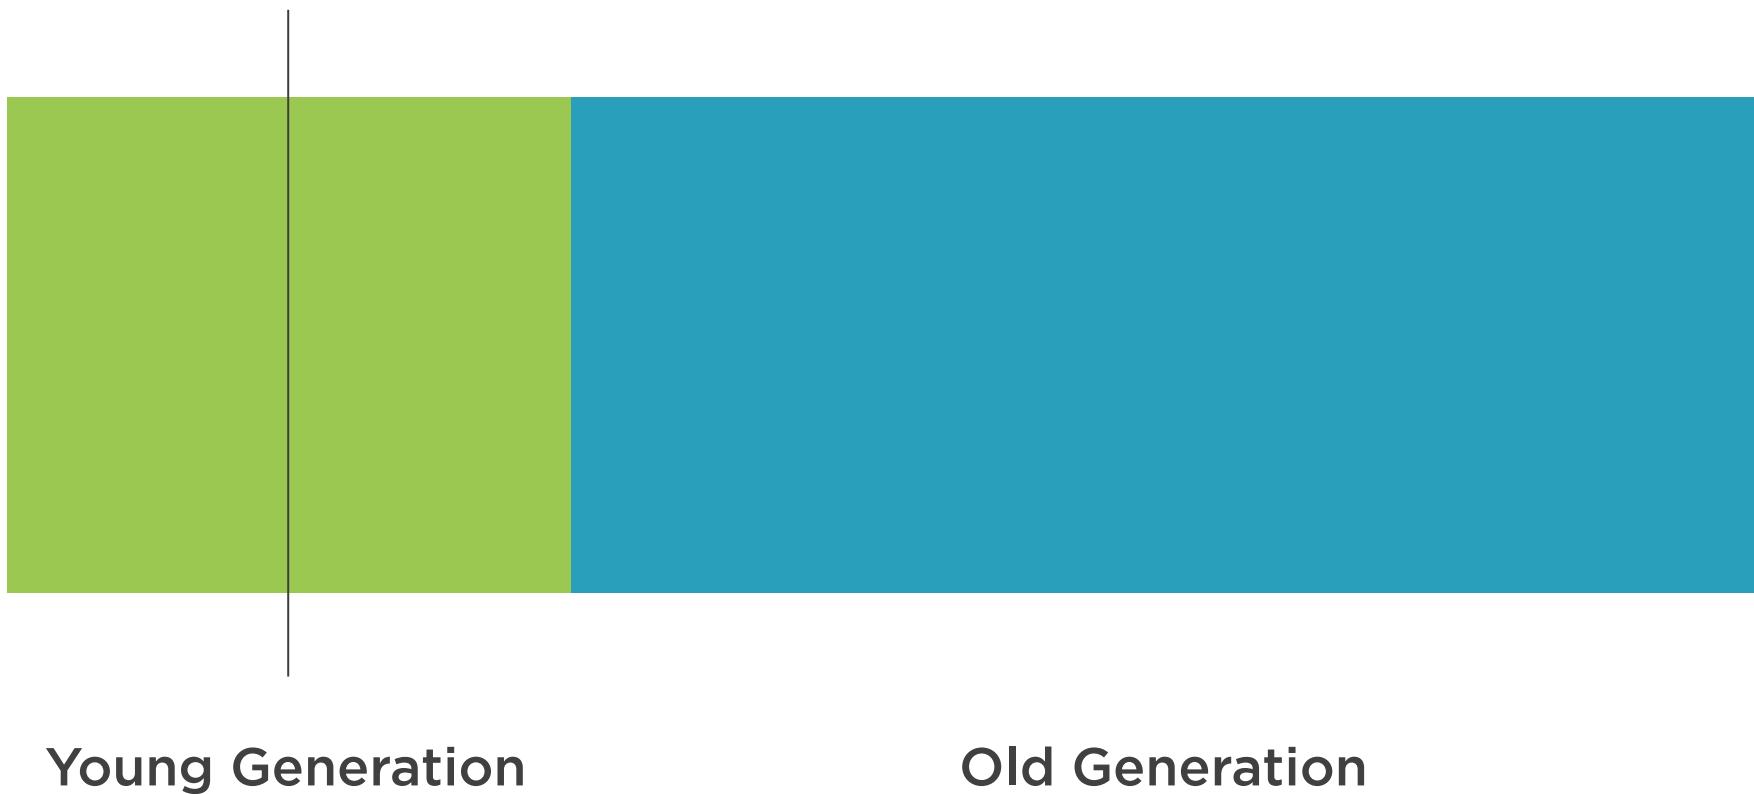


Young Generation

Old Generation



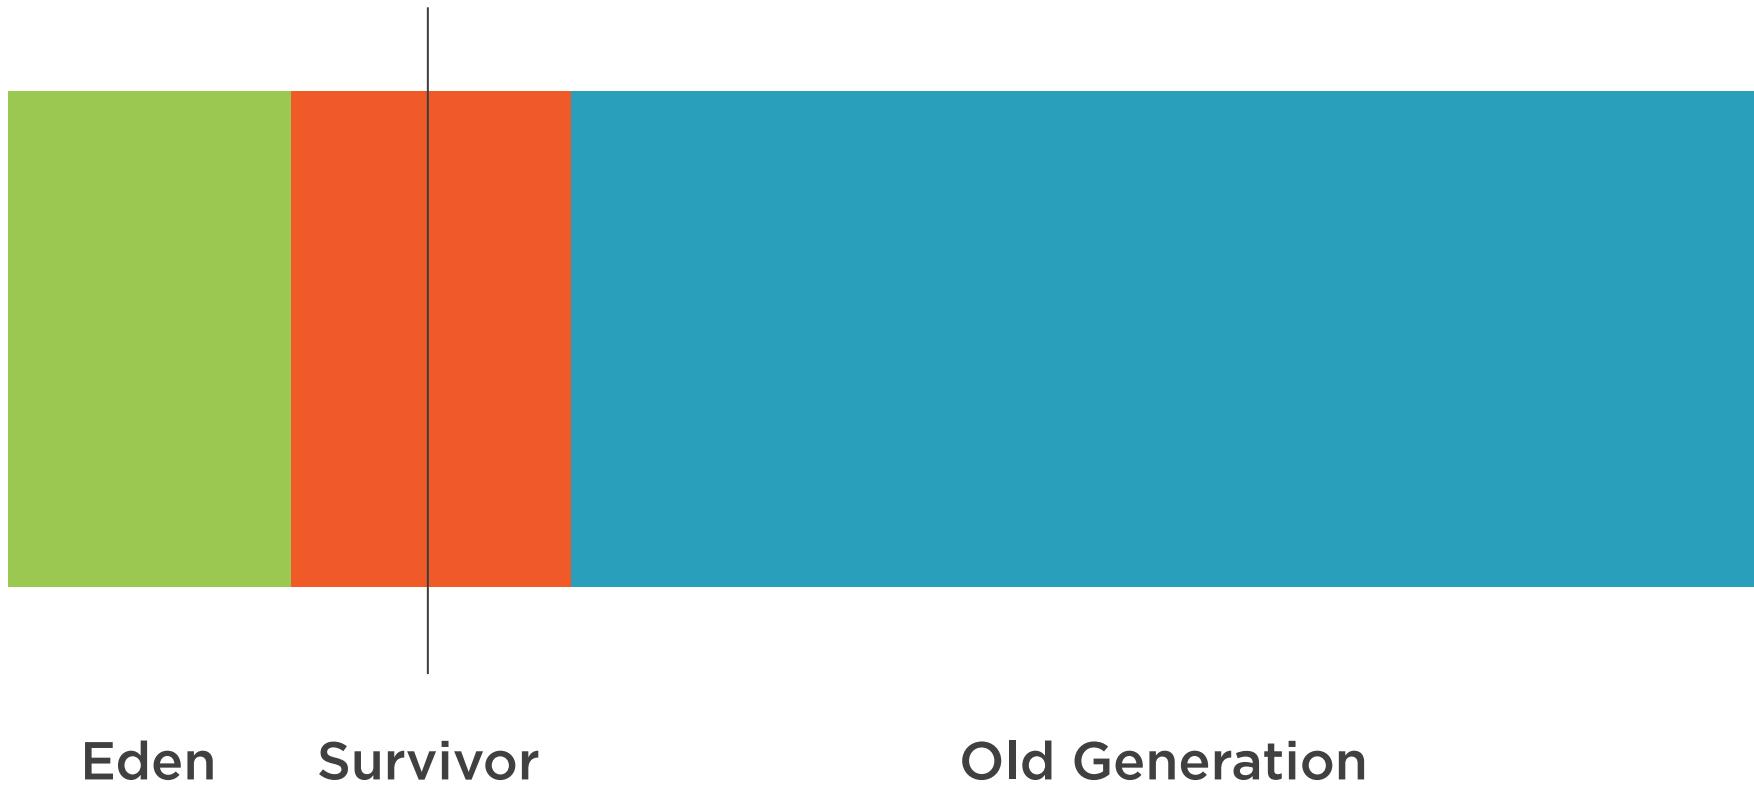
Generational Garbage Collection



Generational Garbage Collection



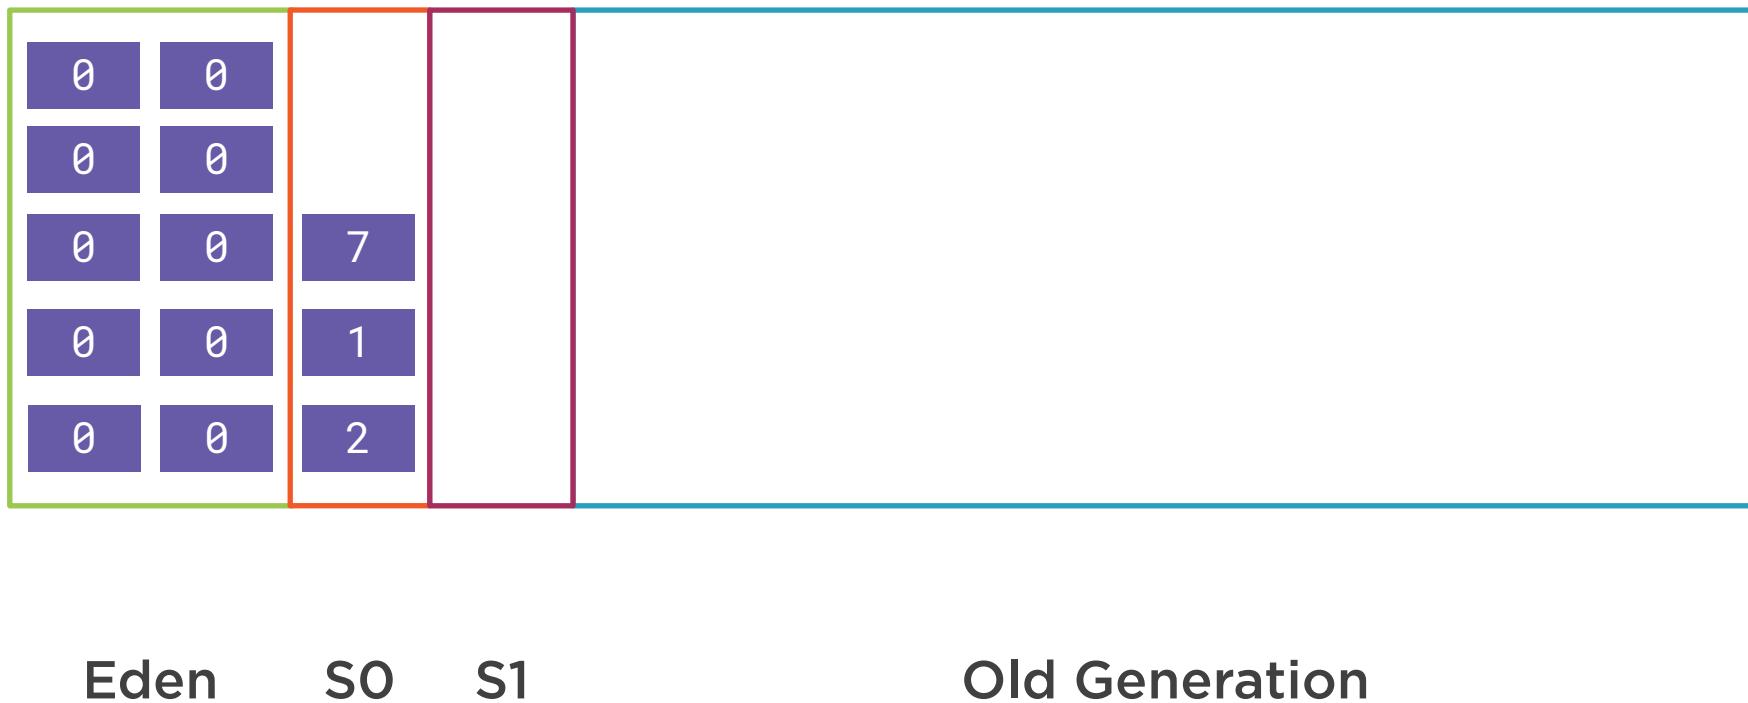
Generational Garbage Collection



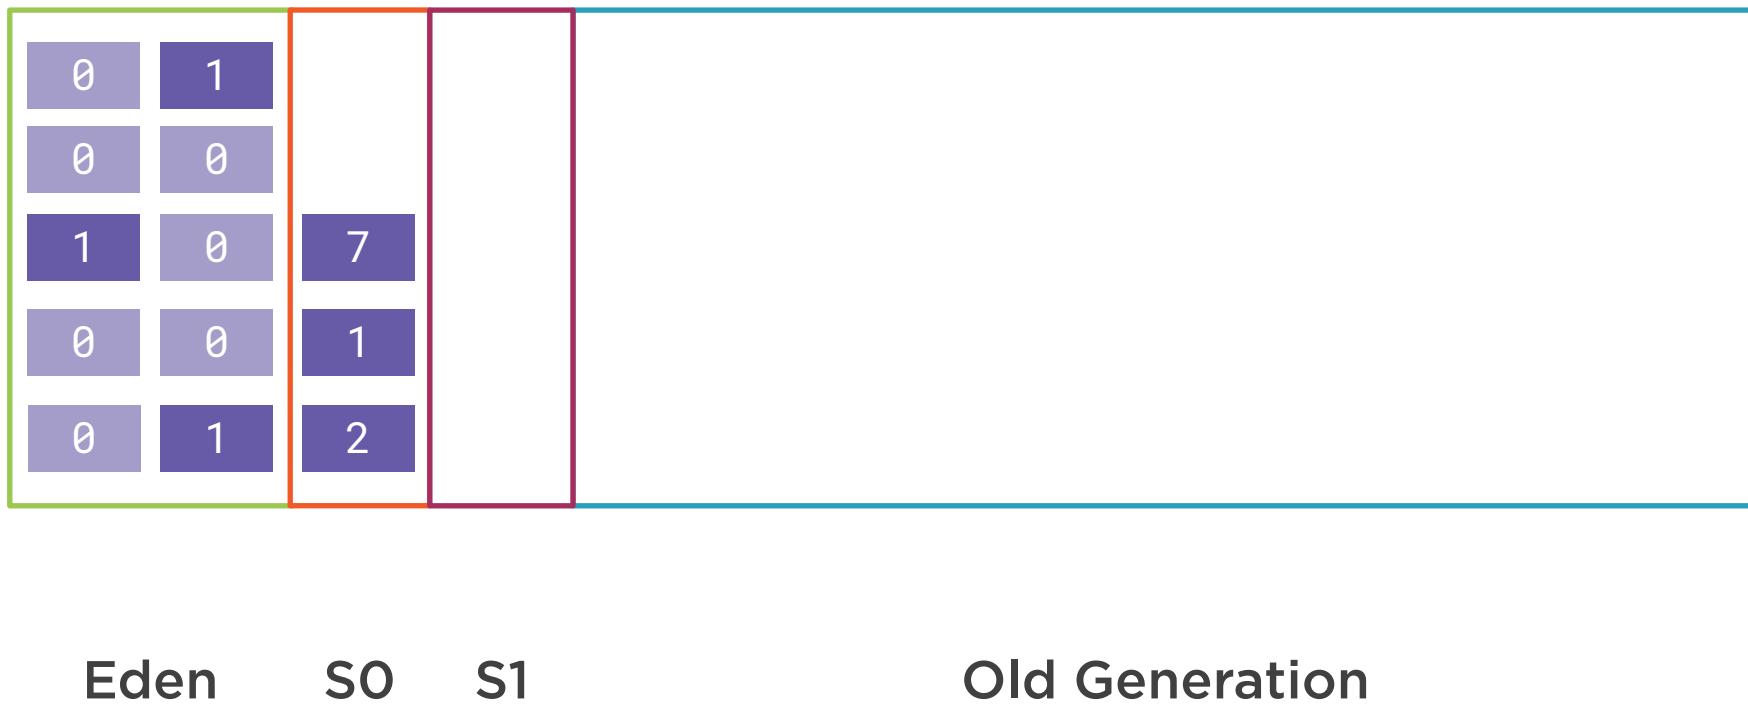
Generational Garbage Collection



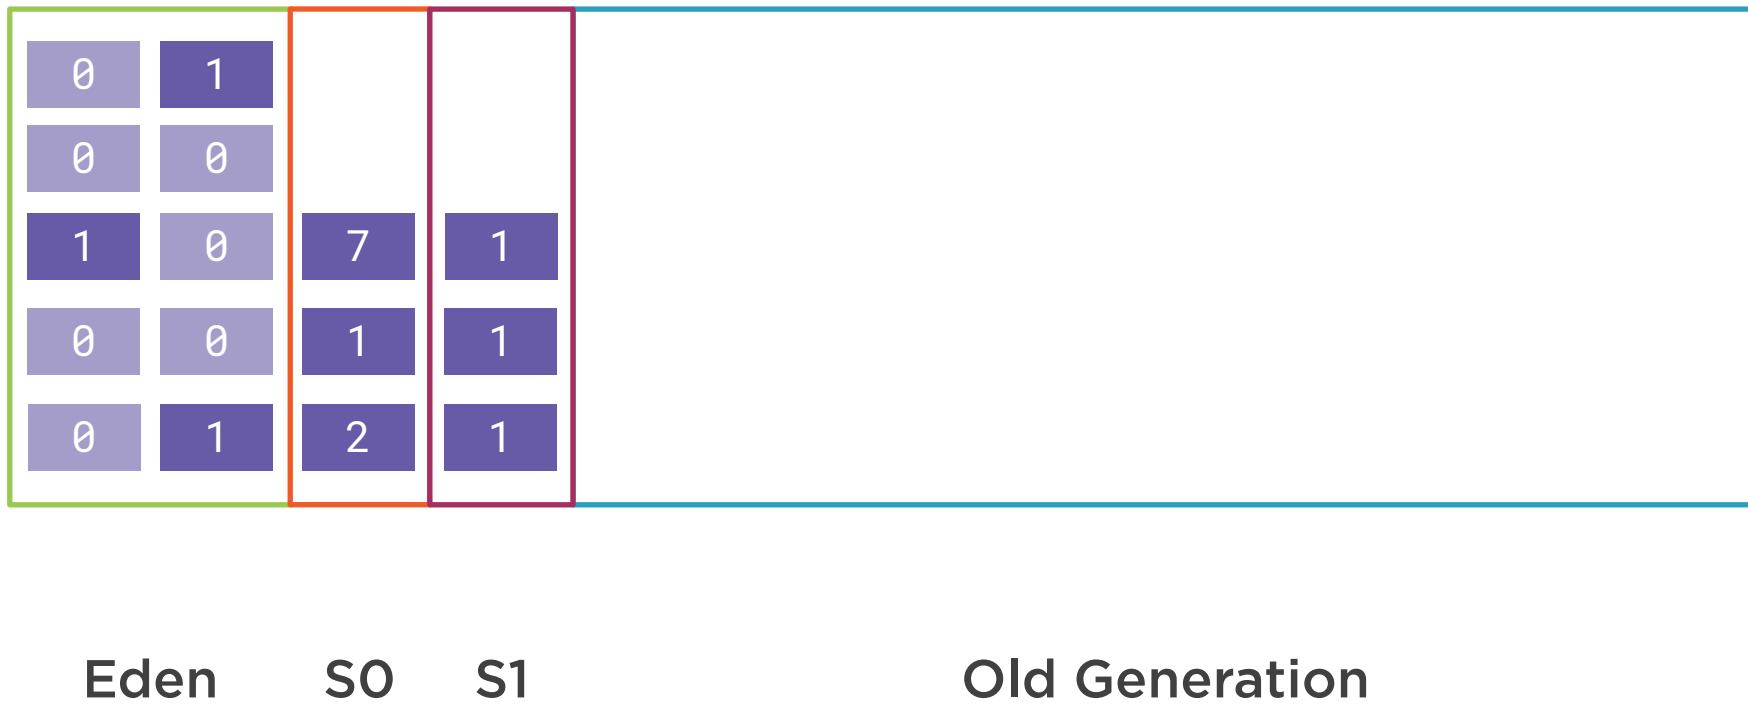
Generational Garbage Collection



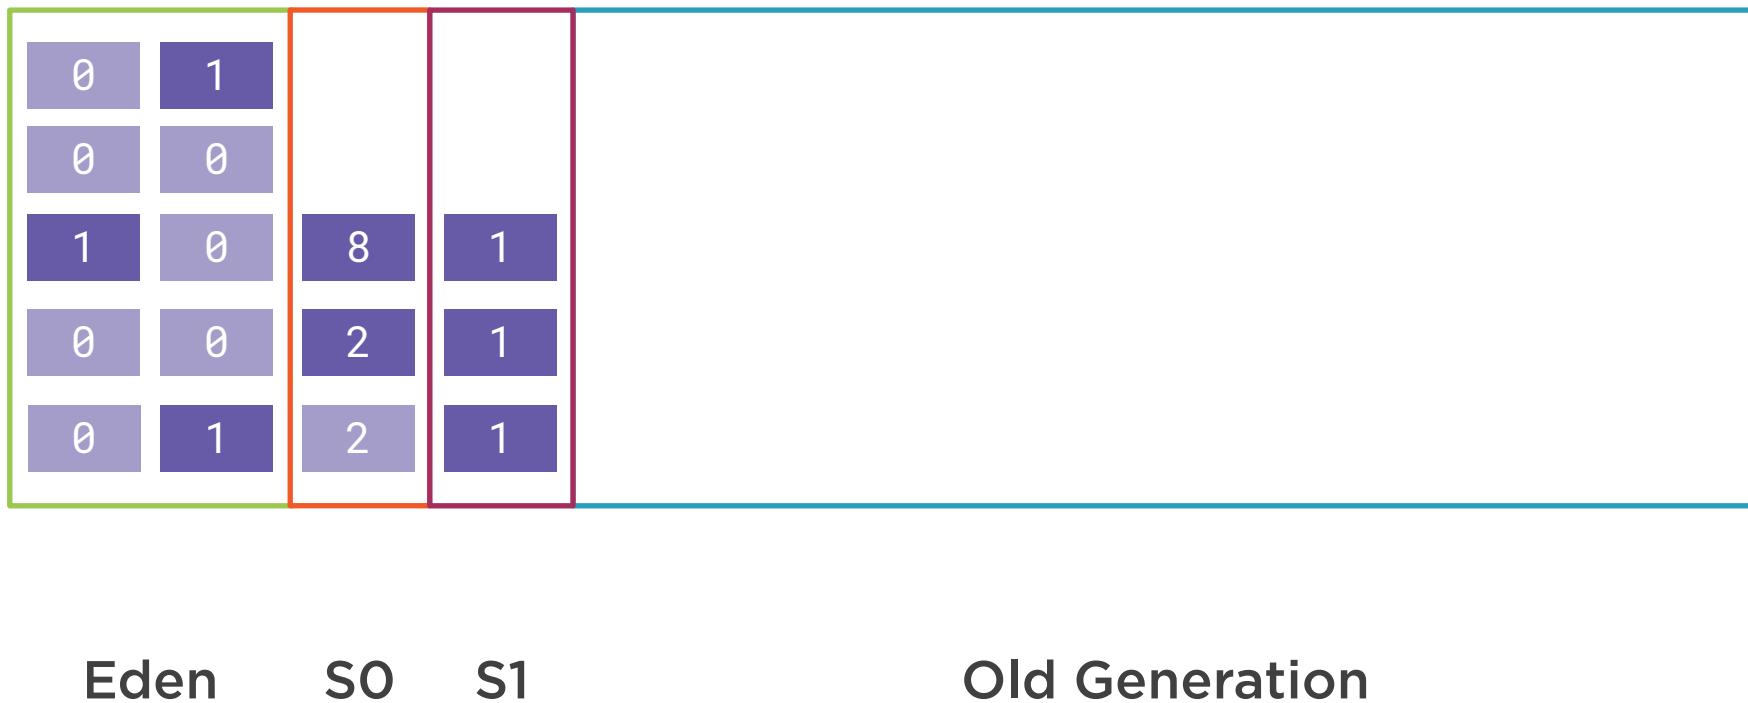
Generational Garbage Collection



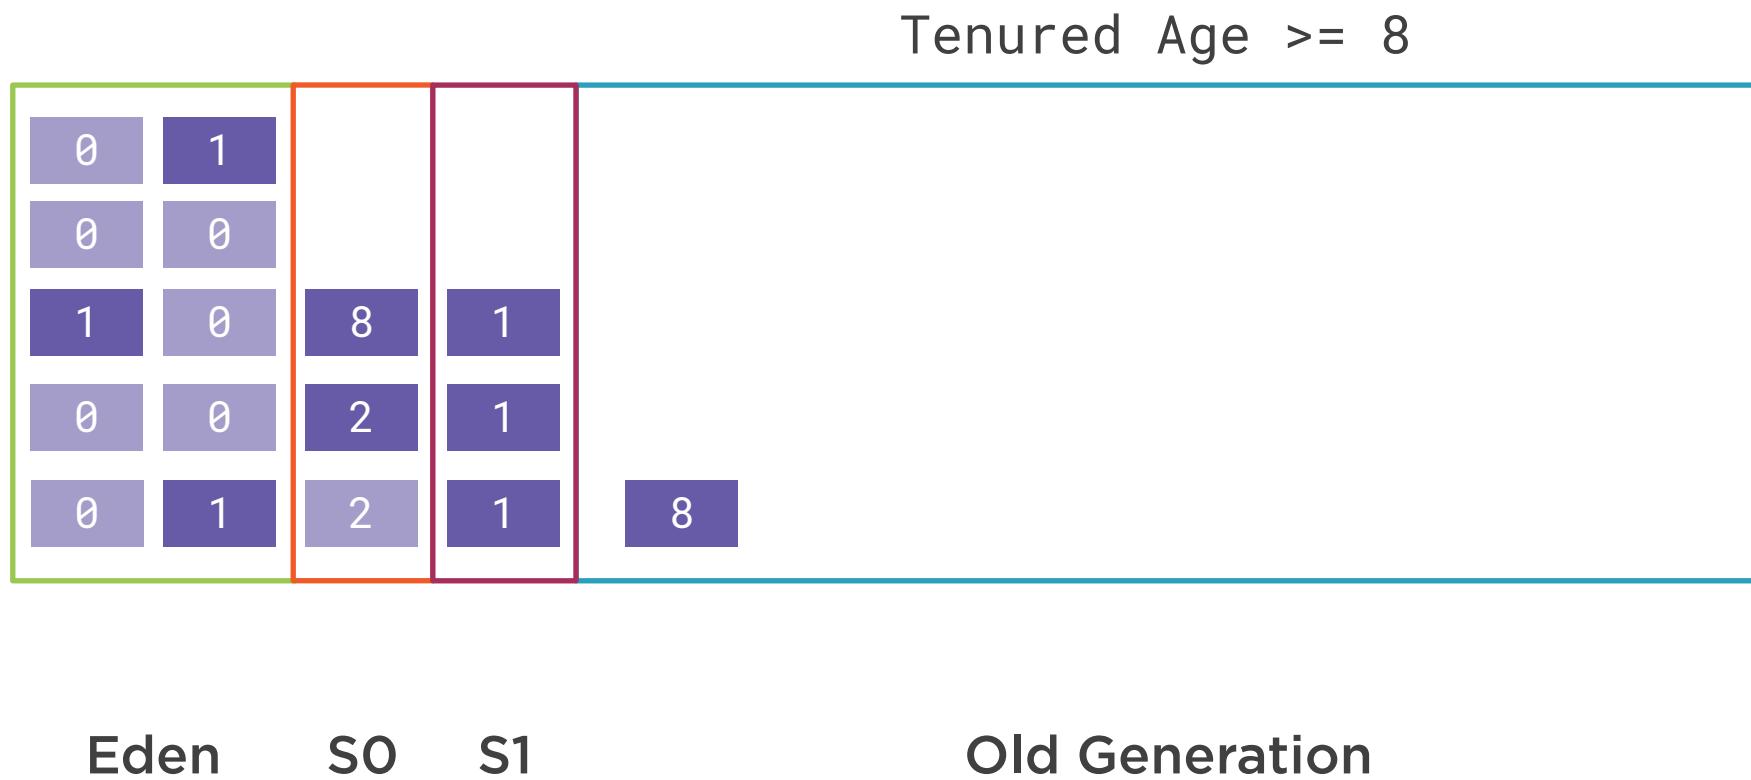
Generational Garbage Collection



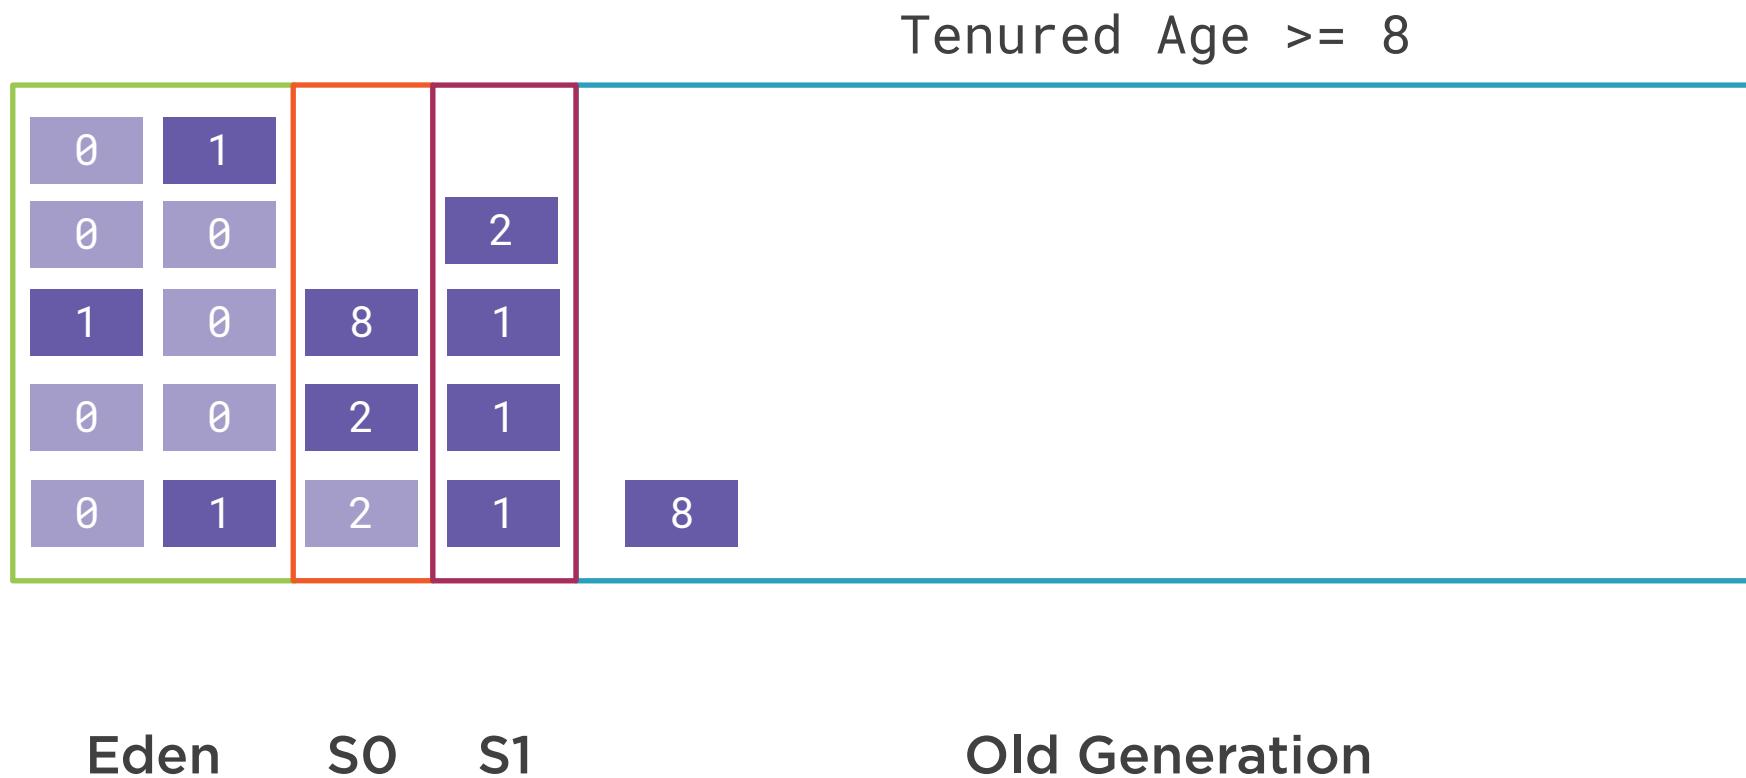
Generational Garbage Collection



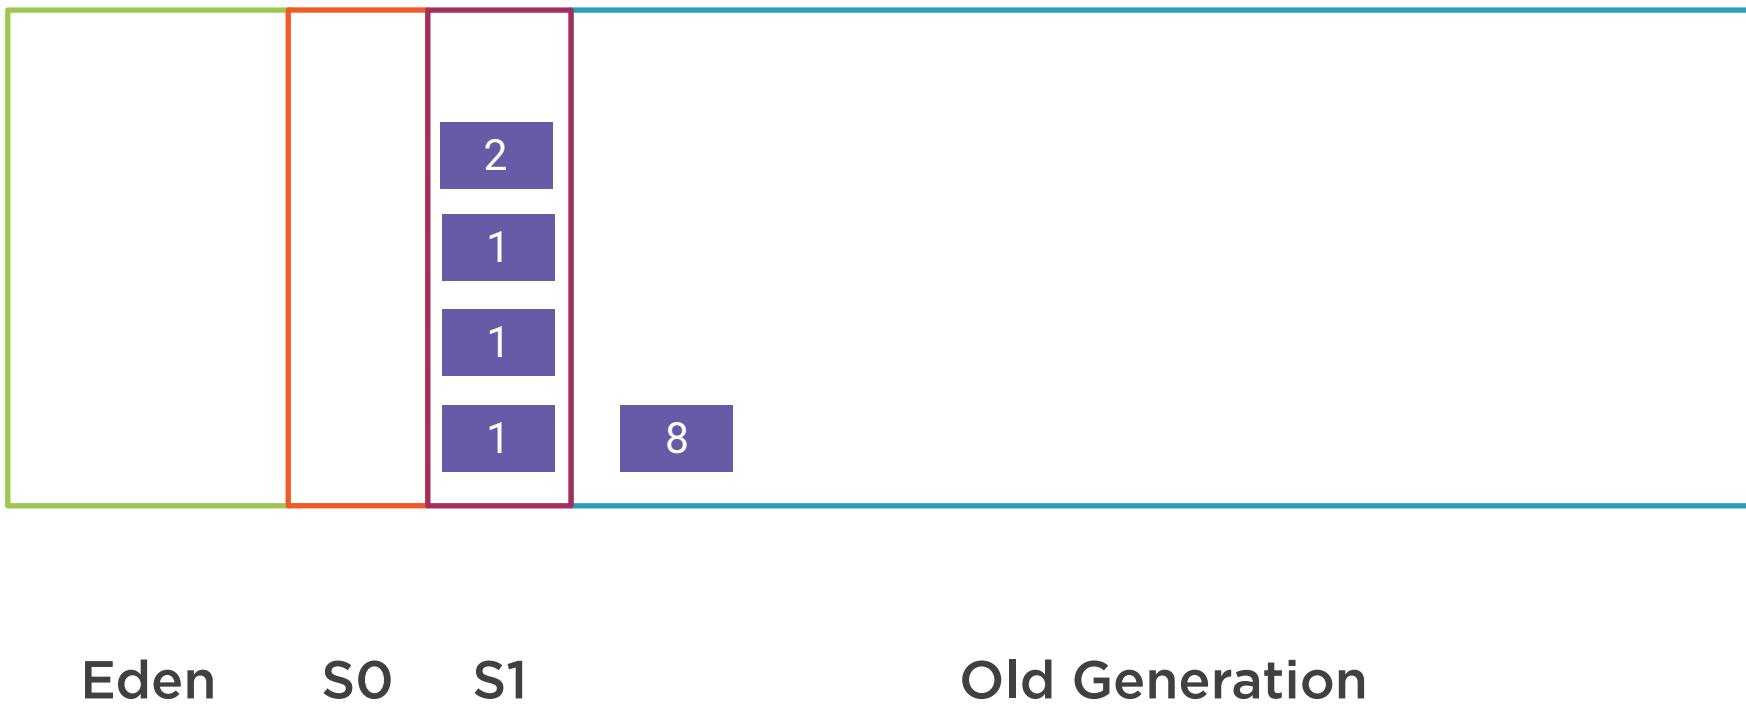
Generational Garbage Collection



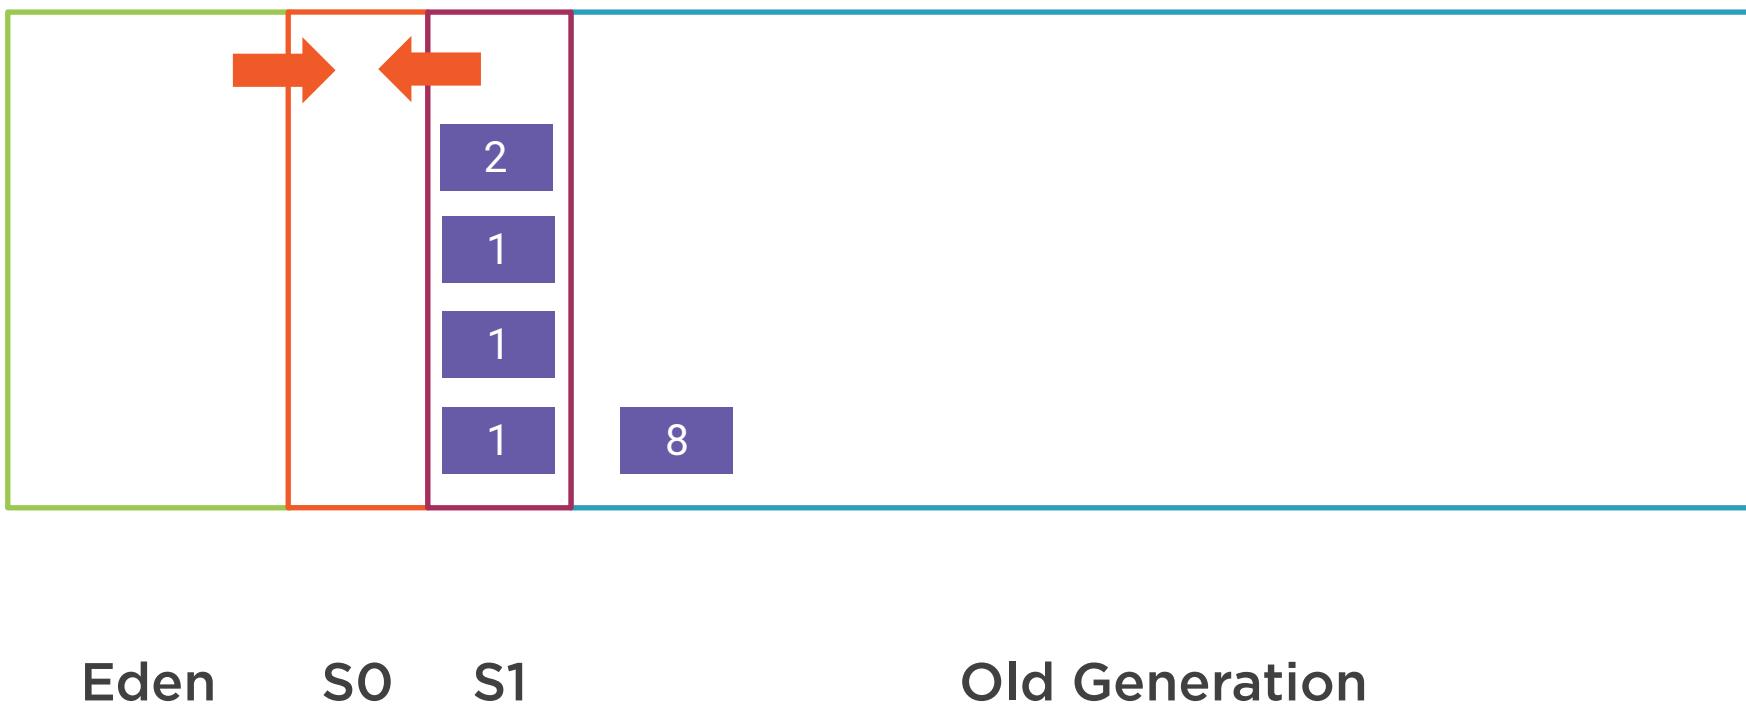
Generational Garbage Collection



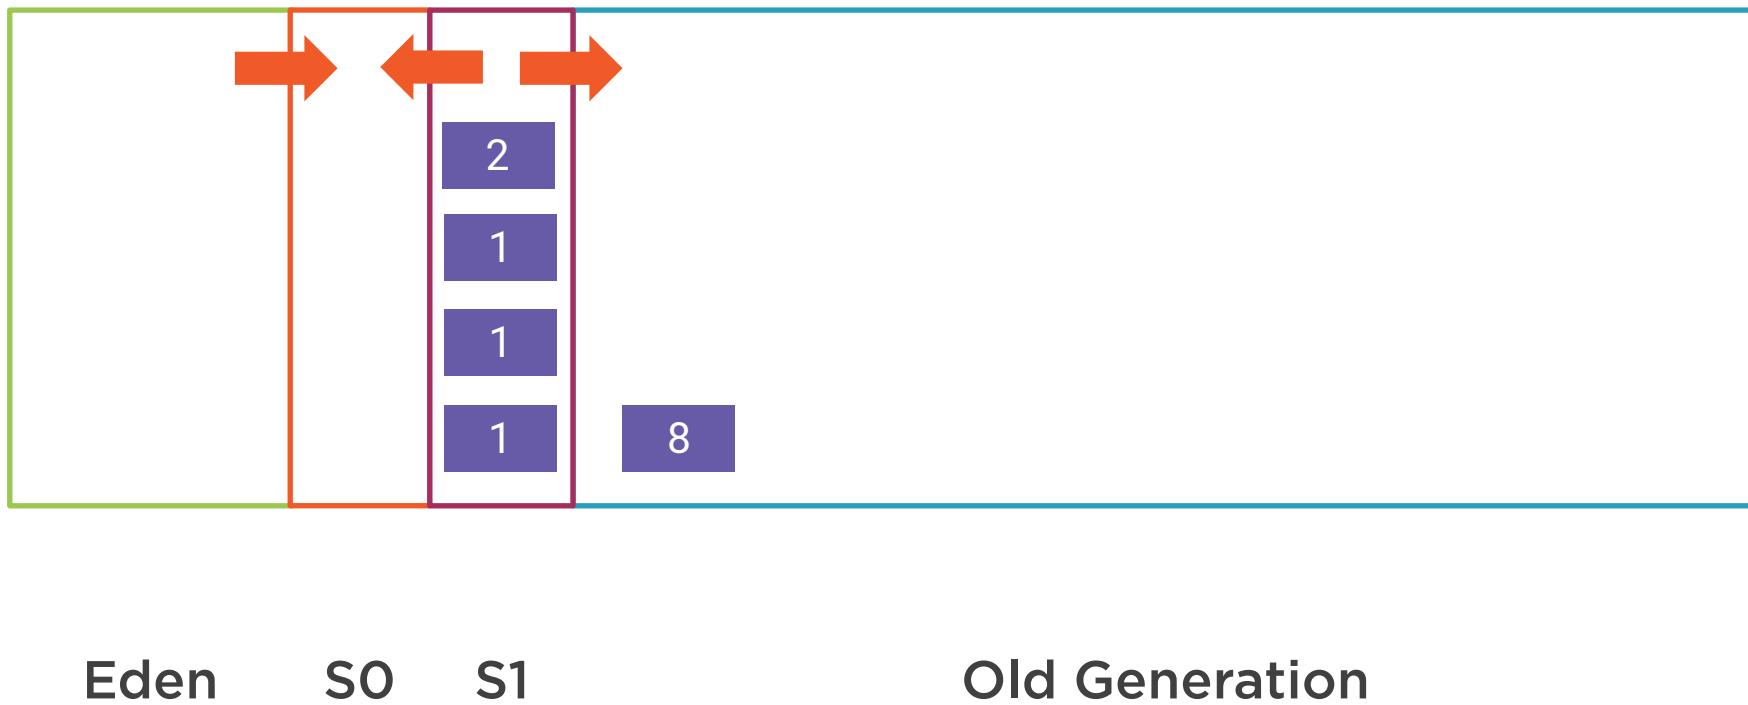
Generational Garbage Collection



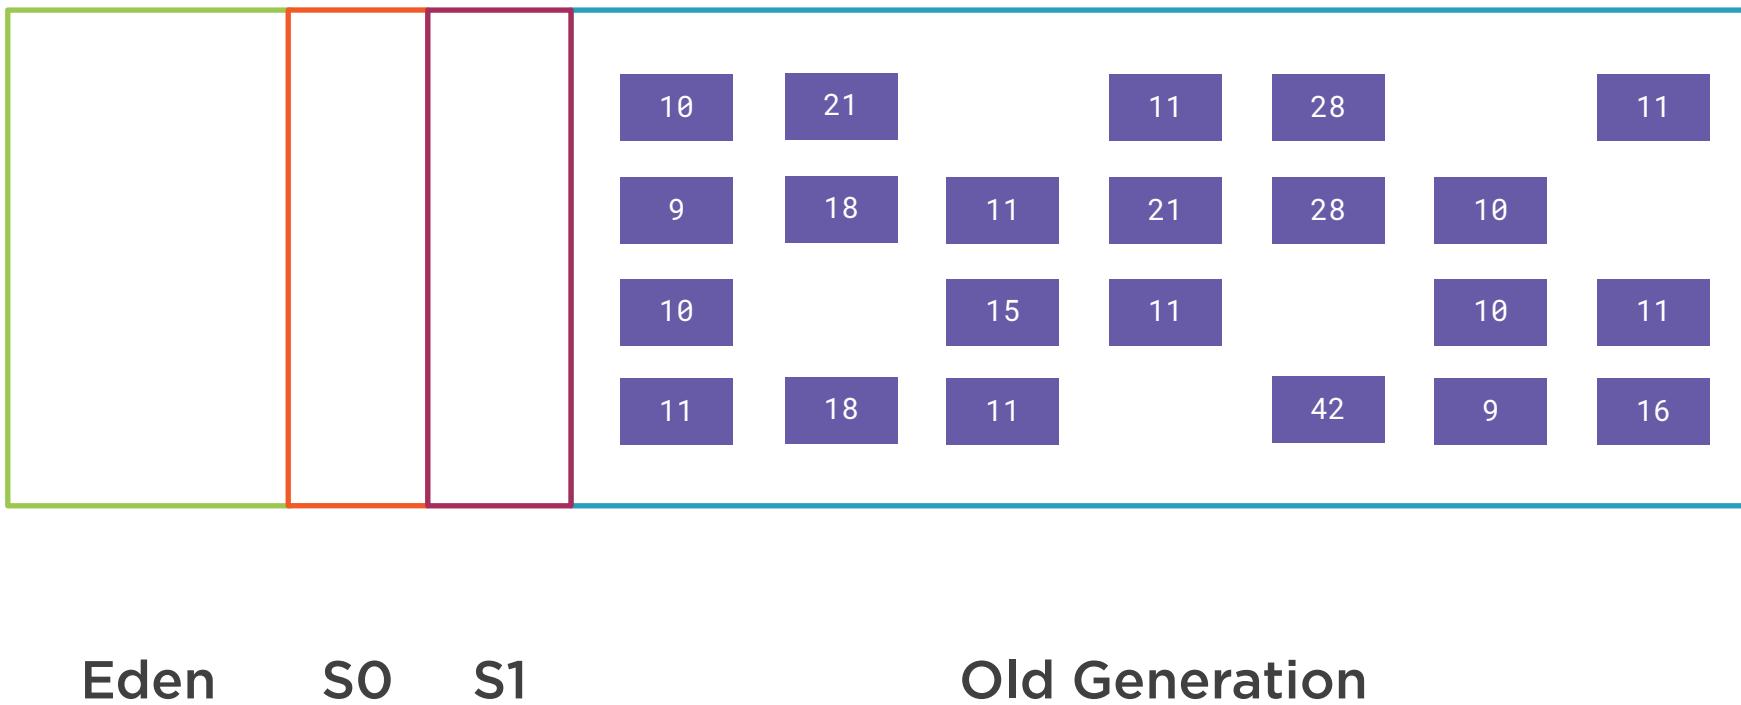
Generational Garbage Collection



Generational Garbage Collection



Generational Garbage Collection



Basic Collectors

- Stop all application threads
- Mark unreachable objects
- Free the memory
- Compact the heap
- Resume application threads



Advanced Collectors

Scan for unreachable objects while the application threads are still running

Only pause threads to free the memory and compact the heap

Known as concurrent collectors, mostly-concurrent collectors, low-pause collectors



Choosing a Garbage Collector



Factors in Choosing a Garbage Collector

**Heap size /
Amount of live
data**

**Number and
speed of available
processors**

**Pause time
requirements**



Serial Collector

Uses one thread to process the heap

To enable: `-XX:+UseSerialGC`

Use when:

- Only 1 CPU/vCPU available and no pause time requirements**
- Multiple small JVMs on a single machine (more than the number of CPUs)**
- Small live data set (up to 100MB)**



Parallel Collector

Uses multiple threads to process the heap

Fully stops all application threads which can result in long pause times but higher throughput

Use when:

- You have a batch application

To enable in Java 9↑: -XX:+UseParallelGC

Default in Java 7 and 8



CMS Collector

Can trace reachable objects and cleanup unreachable objects while application threads are running

GC threads may compete with application threads for CPU

To enable: `-XX:+UseConcMarkSweepGC`

Has been superseded by G1GC and is deprecated in Java 9↑



G1GC Collector

Default collector in Java 9 and up

Designed for multiprocessor machines with large heaps. Tries to achieve the best balance between throughput and latency

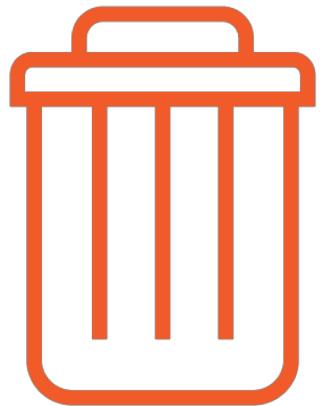
GC threads mark unreachable objects concurrently while application threads are running

Use when: you have an interactive application with low pause time requirements

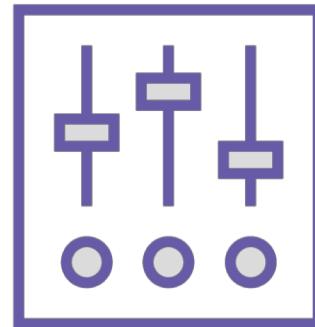
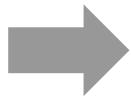
To enable in Java 7/8: -XX:+UseG1GC



GC Tuning Process



Select a collector



Tune heap and
collector settings



Try a different
collector



Shenandoah Collector

Takes the next step beyond G1GC

Can mark unreachable objects and move
reachable objects concurrently while
application threads are running

Can produce 10X reduction in pause times
with only a 10% throughput decrease



Measuring Garbage Collector Performance



GC Logging

GC Log info:

- When collections were run
- How long they lasted
- What were the pause times
- How much memory was reclaimed
- How many objects were promoted to the old generation
- Etc.



```
java -XX:+PrintGCDetails -XX:+PrintGCDateStamps -  
Xloggc:<file-path>
```

Turning on GC Logging in Java 7 / 8



```
java -XX:+PrintGCDetails -XX:+PrintGCDateStamps -  
Xloggc:<file-path>
```

Turning on GC Logging in Java 7 / 8



```
java -XX:+PrintGCDetails -XX:+PrintGCDateStamps -  
Xloggc:<file-path>
```

Turning on GC Logging in Java 7 / 8



```
java -XX:+PrintGCDetails -XX:+PrintGCDateStamps -  
Xloggc:<file-path>
```

Turning on GC Logging in Java 7 / 8



```
java -XX:+PrintGCDetails -XX:+PrintGCDateStamps -  
Xloggc:<file-path> -XX:+UseGCLogFileRotation -  
XX:NumberOfGCLogFiles=10 -XX:GCLogFileSize=10M
```

Turning on GC Logging with File Rotation in
Java 7 / 8



```
java -X:log:gc
```

Turning on GC Logging in Java 9 and Up



```
java -X:log:gc
```

Turning on GC Logging in Java 9 and Up



```
java -X:log:gc:file=<file-path>
```

Turning on GC Logging in Java 9 and Up



```
java -X:log:gc:file=<file-path>:filecount=10, filesize=10M
```

Turning on GC Logging in Java 9 and Up



```
java -X:log:gc*:file=<file-path>:filecount=10, filesize=10M
```

Turning on GC Logging in Java 9 and Up



```
jcmd <pid> VM.log what=gc output=<file-path>
```

Dynamically Turning on GC Logging in Java 9
and Up



[150.732s][info][gc,start] GC(74) Pause Young (Normal) (G1 Evacuation Pause)

[150.732s][info][gc,task] GC(74) Using 4 workers of 4 for evacuation

[150.739s][info][gc,heap] GC(74) Eden regions: 316->0(316)

[150.739s][info][gc,heap] GC(74) Survivor regions: 3->3(40)

[150.739s][info][gc,heap] GC(74) Old regions: 2007->2007

[150.739s][info][gc,heap] GC(74) Humongous regions: 4->4

[150.739s][info][gc Evacuation Pause] GC(74) Pause Young (Normal) (G1
Evacuation Pause) 2329M->2013M(2596M) 6.898ms

[150.739s][info][gc,cpu Real=0.01s] GC(74) User=0.06s Sys=0.00s



[150.732s][info][gc,start] GC(74) Pause Young (Normal) (G1 Evacuation Pause)

[150.732s][info][gc,task] GC(74) Using 4 workers of 4 for evacuation

[150.739s][info][gc,heap] GC(74) Eden regions: 316->0(316)

[150.739s][info][gc,heap] GC(74) Survivor regions: 3->3(40)

[150.739s][info][gc,heap] GC(74) Old regions: 2007->2007

[150.739s][info][gc,heap] GC(74) Humongous regions: 4->4

[150.739s][info][gc Evacuation Pause) 2329M->2013M(2596M) 6.898ms

[150.739s][info][gc,cpu Real=0.01s] GC(74) User=0.06s Sys=0.00s



[150.732s][info][gc,start] GC(74) Pause Young (Normal) (G1 Evacuation Pause)

[150.732s][info][gc,task] GC(74) Using 4 workers of 4 for evacuation

[150.739s][info][gc,heap] GC(74) Eden regions: 316->0(316)

[150.739s][info][gc,heap] GC(74) Survivor regions: 3->3(40)

[150.739s][info][gc,heap] GC(74) Old regions: 2007->2007

[150.739s][info][gc,heap] GC(74) Humongous regions: 4->4

[150.739s][info][gc Evacuation Pause) 2329M->2013M(2596M) **6.898ms**

[150.739s][info][gc,cpu Real=0.01s] GC(74) User=0.06s Sys=0.00s



[150.732s][info][gc,start] GC(74) Pause Young (Normal) (G1 Evacuation Pause)

[150.732s][info][gc,task] GC(74) Using 4 workers of 4 for evacuation

[150.739s][info][gc,heap] GC(74) Eden regions: 316->0(316)

[150.739s][info][gc,heap] GC(74) Survivor regions: 3->3(40)

[150.739s][info][gc,heap] GC(74) Old regions: 2007->2007

[150.739s][info][gc,heap] GC(74) Humongous regions: 4->4

[150.739s][info][gc Evacuation Pause) **2329M->2013M(2596M)** 6.898ms

[150.739s][info][gc,cpu Real=0.01s] GC(74) User=0.06s Sys=0.00s



JEP 158: Unified JVM Logging

Introduce a common logging system for all components
of the JVM



Analyzing the GC Log File



Manual inspection



3rd party tools



Analyzing the GC Log File



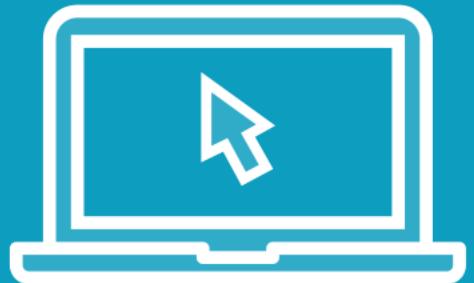
Manual inspection



3rd party tools
GCEasy.io
GCViewer



Demo



GCEasy.io



Garbage Collection Tuning



Tuning the Size of the Heap

If the heap is too small: there'll be too many GCs in order to free memory and application throughput would suffer

If the heap is too large: GC pauses would take too long and response time metrics would be affected



```
java -XmsN -XmxN
```

Setting the Heap Size



10	21	10	11	28
9	18	11	21	28
10		15	11	10
11	18	11		42

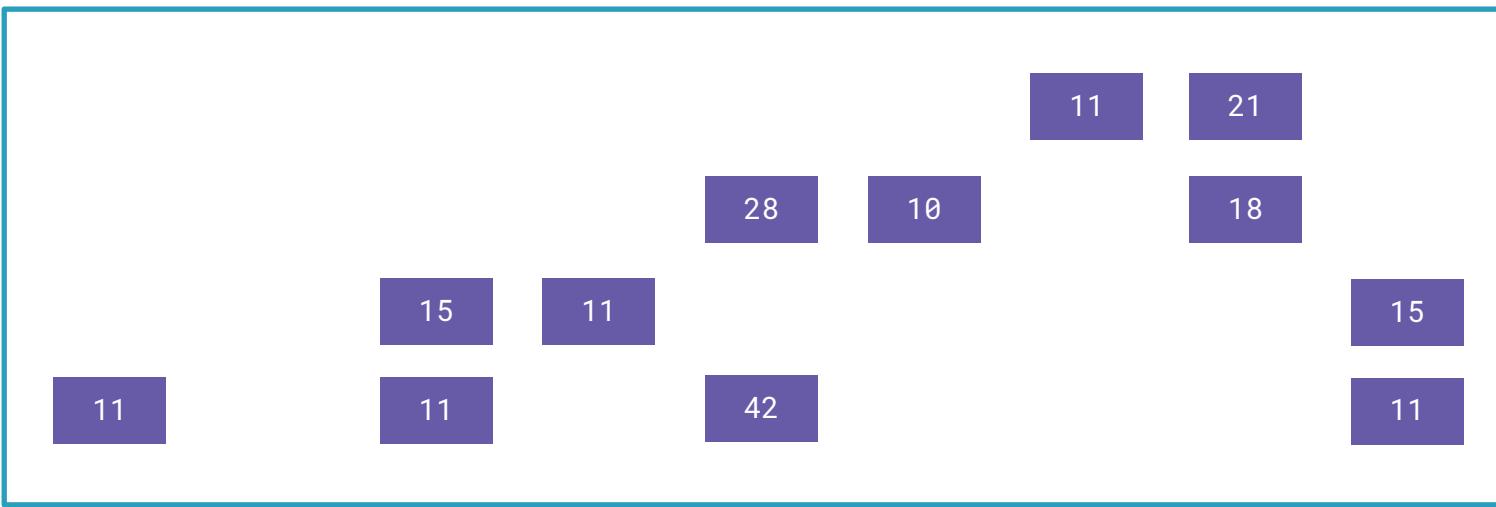


		21		11	28		11
9	18	11	21	28	10		
10		15	11		10	11	
11	18	11		42	9	16	



		21		11	28		11	21
9	18			28	10		18	
		15	11		10	11		15
11		11		42	9	16		11

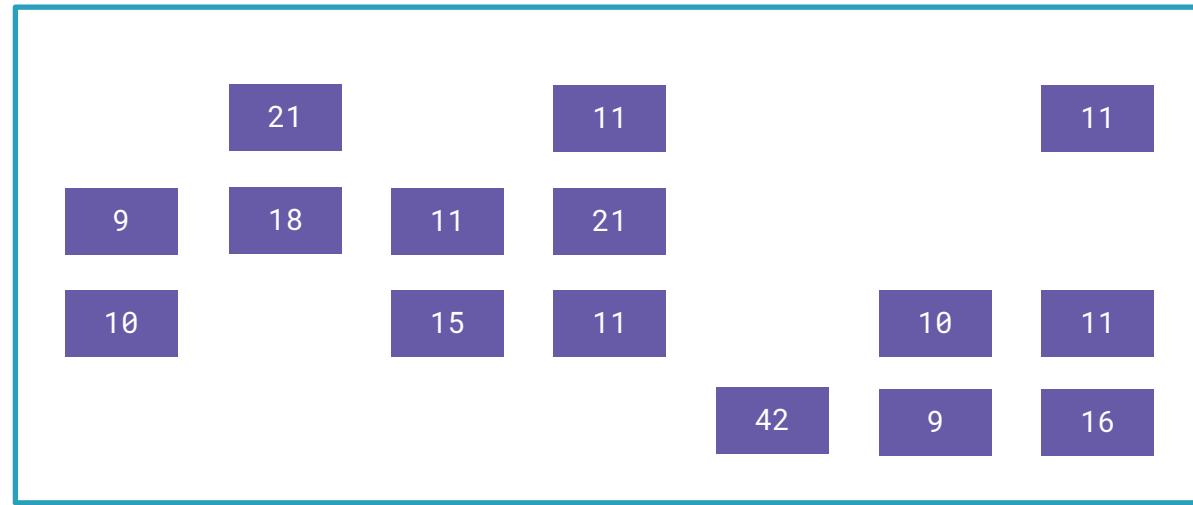




	21		11		11
9	18	11	21		
10		15	11	10	11
			42	9	16



Adaptive Sizing



Turning off Adaptive Sizing

Should do so only if you've finely tuned your application's GC behavior and sizes

Set flag: -XX:-UseAdaptiveSizePolicy

-Xms6g -Xmx6g

-XX:NewSize=1g -XX:MaxNewSize=1g



Max Heap Sizing

Max heap size should not exceed the machine's physical memory

Heap should be about 30% occupied after a full GC



MaxGCPauseMillis

Example: -XX:MaxGCPauseMillis=350

Sets a target for the maximum GC pause time

In the Parallel collector:

- **Adjusts the size of the young and old generation**
- **Adjusts the size of the heap**



MaxGCPauseMillis

In the G1GC collector:

- Adjusts the size of the young and old generation
- Adjusts the size of the heap
- Starts background processing sooner
- Changes the tenuring threshold
- Adjusts the number of old generation regions processed

G1GC default: 200 ms



Concurrent Garbage Collection Failures

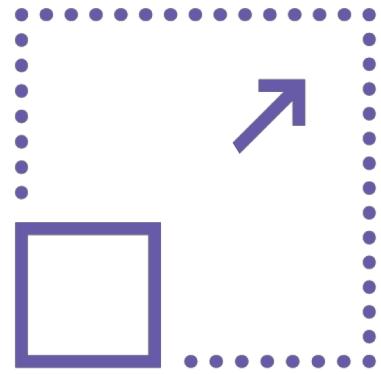
Since the application threads are still running, they may produce garbage faster than the GC threads can clean them

Known as:

- Concurrent mode failure
- Promotion failure
- Evacuation failure



Solving Concurrent GC Failures



Increase the size of
the heap

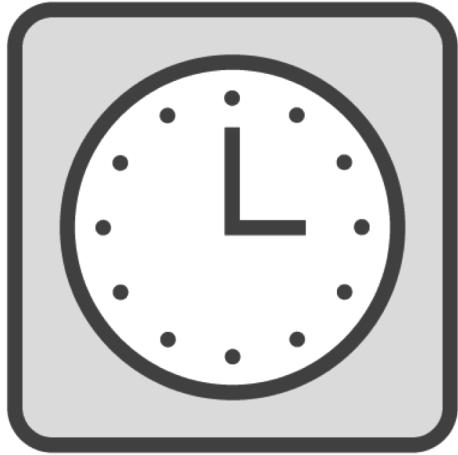


Start background
processing earlier



Use more background
threads





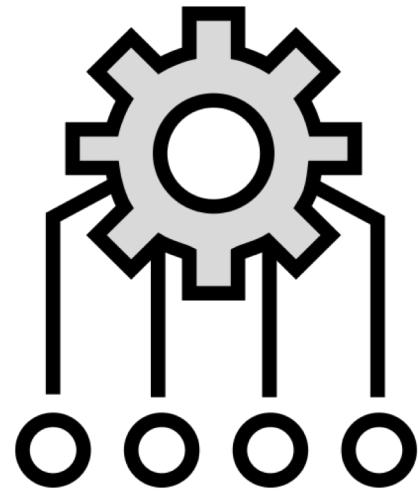
Reduce the threshold at which the G1 cycle is triggered: -

`XX:InitiatingHeapOccupancyPercent=45`

A GC cycle is triggered when the heap becomes 45% filled

If set too low, GCs will happen too frequently





To increase the number of background threads: -XX:ConcGCThreads=4

Default value: ConcGCThreads =
 $(ParallelGC + 2) / 4$

